

COMP304

Operating Systems (OS)

Operating System Structure

Didem Unat
Lecture 2

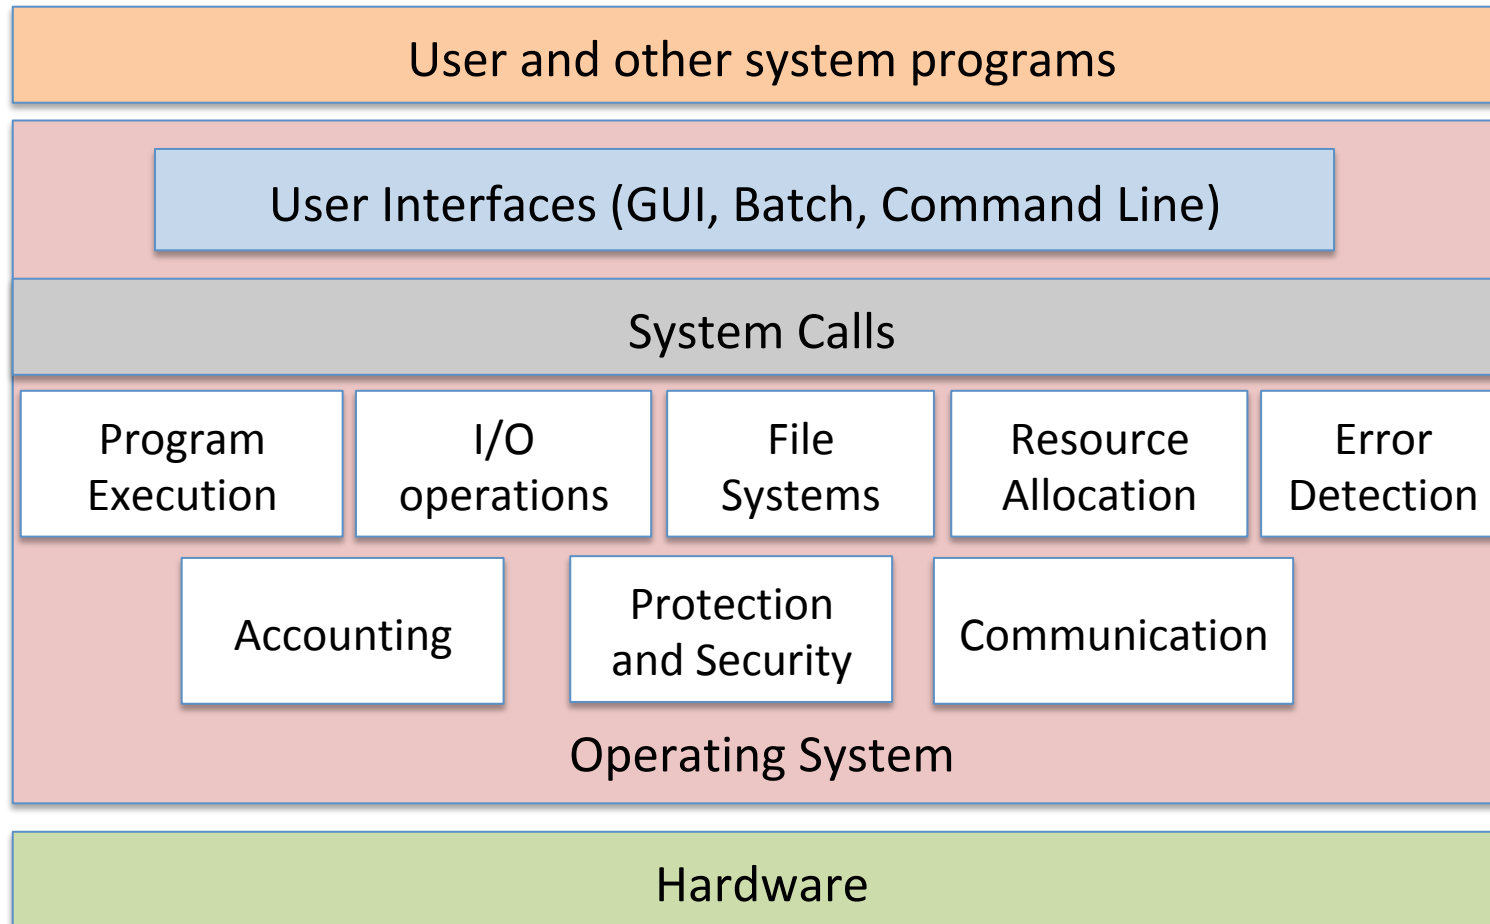
Outline

- Operating System Services
- Command Interpreter
- Dual Mode Operation
- System Calls and Types
- I/O, Memory and CPU Protection
- Operating System Design Structure

Computer Startup

- **Bootstrap program** is loaded at power-up or reboot
 - Typically stored in ROM, generally known as **firmware**
 - Initializes all aspects of a system
 - Loads operating system **kernel** into main memory and starts execution
 - The first system process is 'init' in Linux
 - When the system is fully booted, it waits for some event to occur
- **Kernel**
 - The ``one'' program running at all times (the core of OS)
 - Everything else is an application program
- **Process**
 - An executing program (active program)

Operating System Services



Operating System Services (1/3)

- **User interface** - Almost all operating systems have a user interface (**UI**).
 - Varies between **Command-Line (CLI)**, **Graphics User Interface (GUI)**, or **Batch**
- **Program execution** - The system must be able to load a program into memory and to run that program, end execution, either normally or abnormally (indicating error)
- **I/O operations** - A running program may require I/O, which may involve a file or an I/O device
- **File-system manipulation** - Programs need to read and write files and directories, create and delete them, search them, list file information, manage permissions.

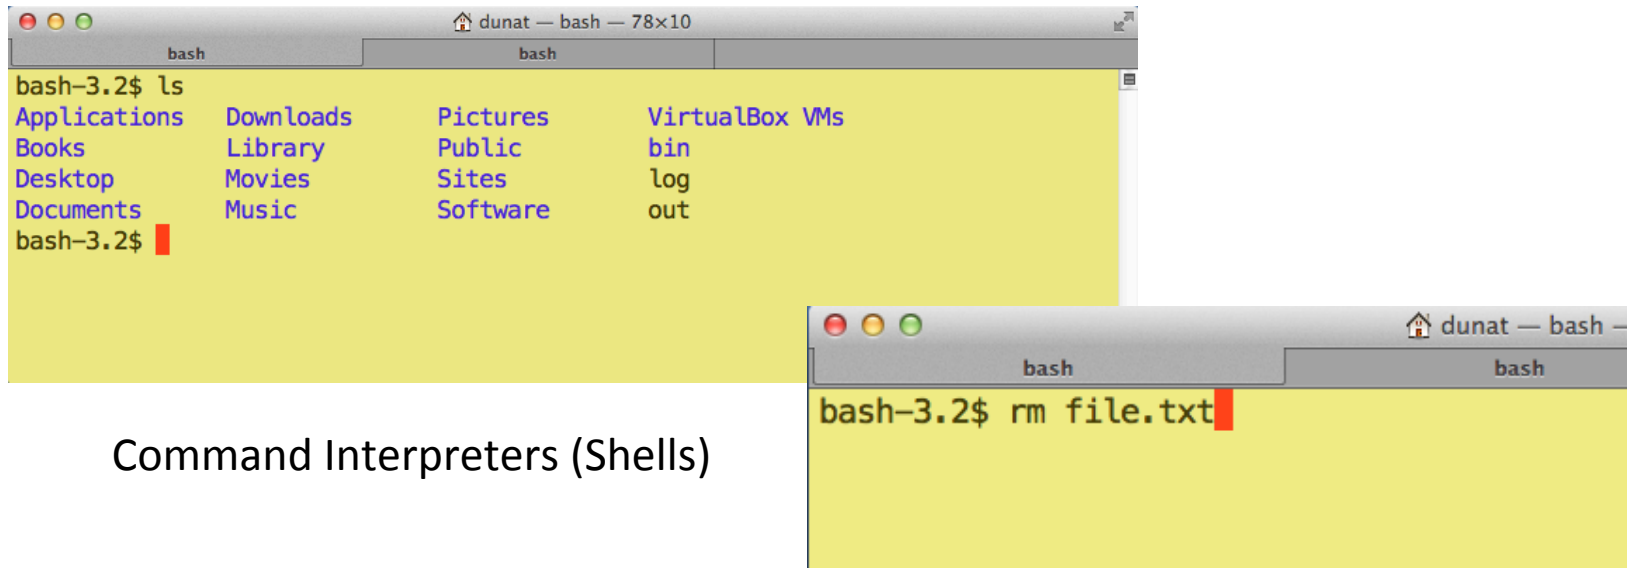
Operating System Services (2/3)

- **Communications** – Processes may exchange information, on the same computer or between computers over a network
 - Communications may be via shared memory or through message passing (packets moved by the OS)
- **Error detection** – OS needs to be constantly aware of possible errors
 - May occur in the CPU and memory hardware, in I/O devices, in user program
 - For each type of error, OS should take the appropriate action to ensure correct and consistent computing

Operating System Services (3/3)

- **Resource allocation** - When multiple users or multiple jobs running concurrently, resources must be allocated to each of them
 - Many types of resources - Some (such as CPU cycles, main memory, and file storage) may have special allocation code, others (such as I/O devices) may have general request and release code
- **Accounting** - To keep track of which users use how much and what kind of computer resources, improve response time to users
- **Protection and security** - The owners of information stored in a multiuser or networked computer system may want to control use of that information, concurrent processes should not interfere with each other

Command Interpreters



Command Interpreters (Shells)

- In UNIX everything is a file
 - Command interpreter does not understand the command (e.g. “rm”)
 - It merely uses the command to identify a file to be loaded into memory and executed.
 - For example, shell would search for a file called ‘rm’, load the file into memory and execute it with the parameter file.txt
 - Thus, programmer can add new commands to the system easily by creating new files

Src code of Linux Commands

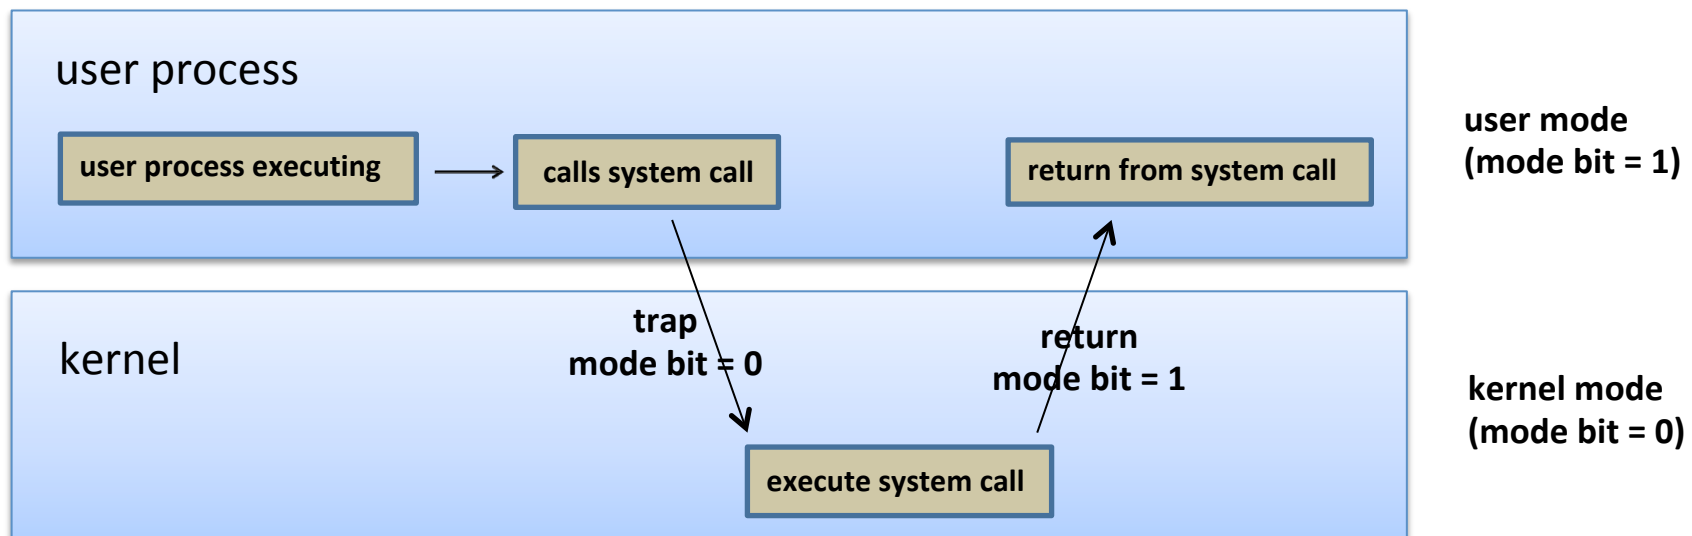
- All these basic commands are part of the **coreutils** package.
 - <http://www.gnu.org/software/coreutils/>
 - commands such as rm, ls, chmod, cp ...
 - <http://git.savannah.gnu.org/cgit/coreutils.git/tree/src>
- For example, “ls” command:
 - <http://git.savannah.gnu.org/cgit/coreutils.git/tree/src/ls.c>
 - Only 5308 code lines for a command 'easy enough'

OS Protection: Dual-Mode Operation

- **Dual-mode** operation allows OS to protect itself and other system components
 - **User mode** and **kernel mode**
 - **Mode bit** provided by hardware
 - Provides ability to distinguish when system is running user code or kernel code
 - Some instructions designated as **privileged**, only executable in kernel mode
 - For example, I/O related instructions are privileged
- Ensures that an incorrect program cannot cause other programs to execute incorrectly.

Transition from User to Kernel Mode

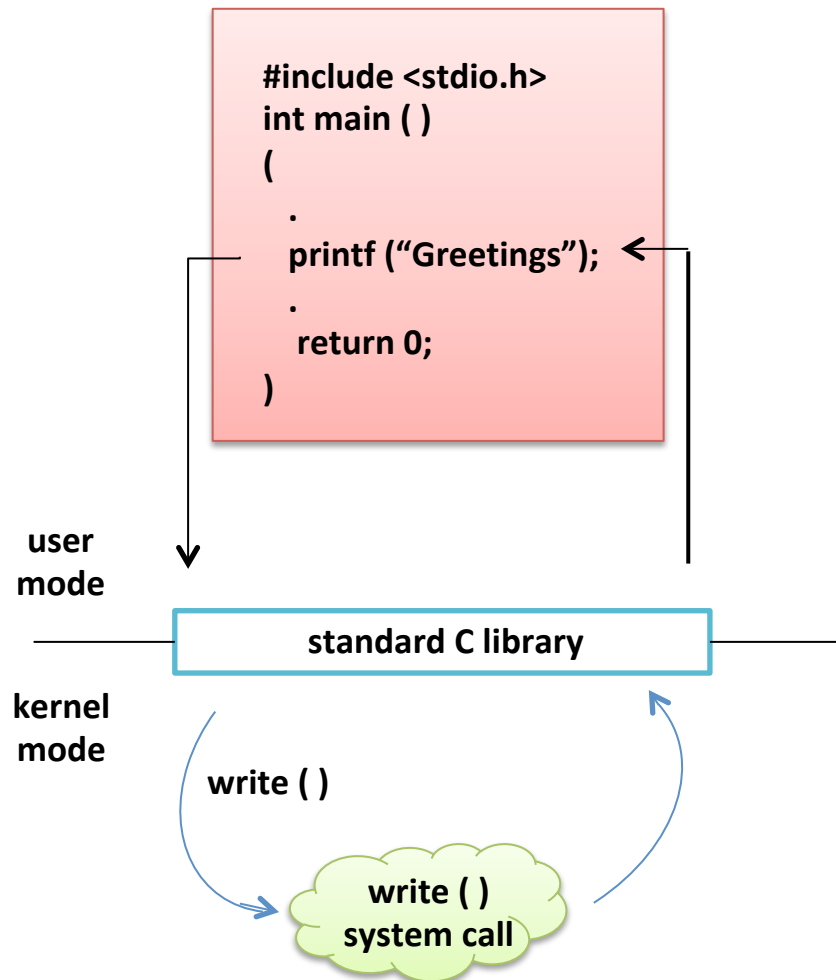
- **System Call**
 - Results in a transition from user to kernel mode
 - Return from call resets it to user mode
- Software error or a user request creates an **exception or trap**



System Calls

- Programming interface to the services provided by the OS
 - Well-defined and safe implementation for service requests
 - Typically written in a high-level language (C or C++)
- A typical OS executes 1000s of system calls per second
- Mostly accessed by programs via a high-level **Application Program Interface (API)** rather than direct system call use
 - Wrapper functions for the system calls
- Three most common APIs are
 - Windows API for Windows,
 - POSIX API for POSIX-based systems (including virtually all versions of UNIX, Linux, and Mac OS X),
 - Java API for the Java virtual machine (JVM)
- Why use APIs instead of using system calls directly?

Standard C Library Example

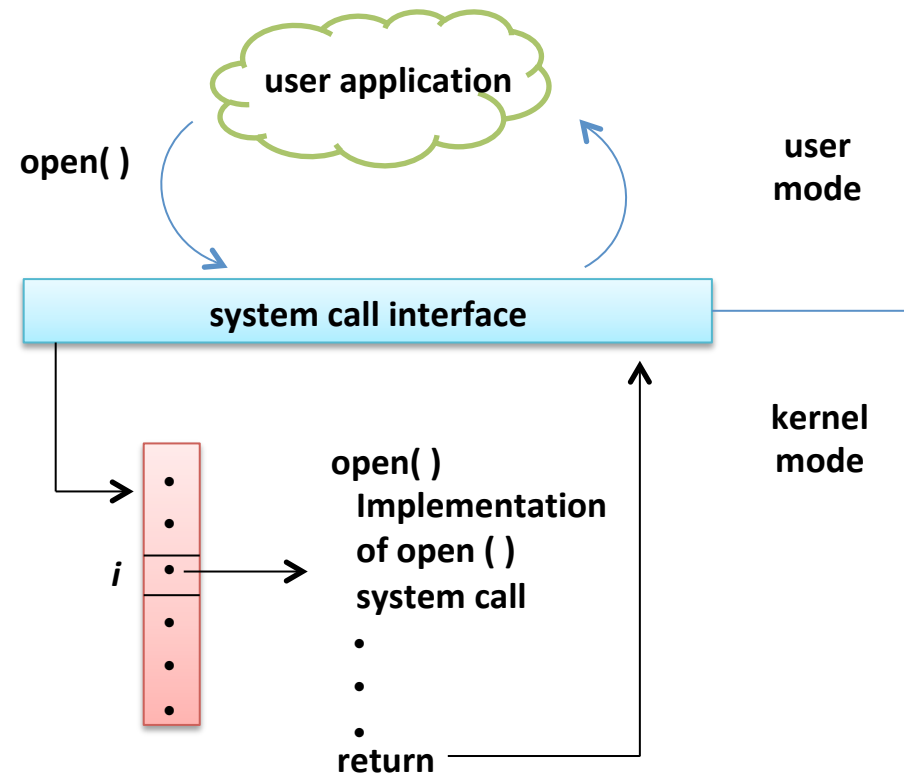


- C program invoking `printf()` library call, intercepts function call in the API and invokes the necessary system calls within the operating system
 - Calls `write()` system call
- Caller needs to know nothing about
 - how the system call is implemented
 - what it does during execution

Example: Linux System Calls

- A system call number is a unique integer in Unix-based OSs
 - There are about >300 system calls in Linux
 - A list of all registered system calls is maintained in the *system call table*
 - *Those numbers cannot be changed or recycled*
 - See the list of system calls with a command
 - (Location might differ depending on the Unix distribution)


```
cat /usr/include/asm/unistd.h | less
```



Example: Linux System Calls

```
cat /usr/include/asm/unistd.h | less
```

- One can also call a service by directly using its number
 - `syscall(system_call_number, arguments)`
- Actual Implementation of a system call is in different files in the kernel src
 - <http://syscalls.kernelgrok.com/>
 - <http://lxr.free-electrons.com/source/>



Types of System Calls

| | Windows | Unix |
|-------------------------|---|--|
| Process Control | CreateProcess() ExitProcess() WaitForSingleObject() | fork() exit() wait() |
| File Manipulation | CreateFile() ReadFile() WriteFile() CloseHandle() | open() read() write() close() |
| Device Manipulation | SetConsoleMode() ReadConsole() WriteConsole() | ioctl() read() write() |
| Information Maintenance | GetCurrentProcessID() SetTimer() Sleep() | getpid() alarm() sleep() |
| Communication | CreatePipe() CreateFileMapping() MapViewOfFile() | pipe() shmget() mmap() |
| Protection | SetFileSecurity() InitializeSecurityDescriptor() SetSecurityDescriptorGroup() | chmod() umask() chown() |

- Types of system calls classified under 6 categories. Table gives an example for Windows and Unix syscalls.

Privileged Instructions

- The dual mode of operation provides us with the means for protecting the operating system from errant users—and errant users from one another.
- We accomplish this protection by designating some of the machine instructions that may cause harm as **privileged instructions**.
 - The hardware allows privileged instructions to be executed only in kernel mode.
 - If an attempt is made to execute a privileged instruction in user mode, the hardware does not execute the instruction but rather treats it as illegal and traps it to the operating system

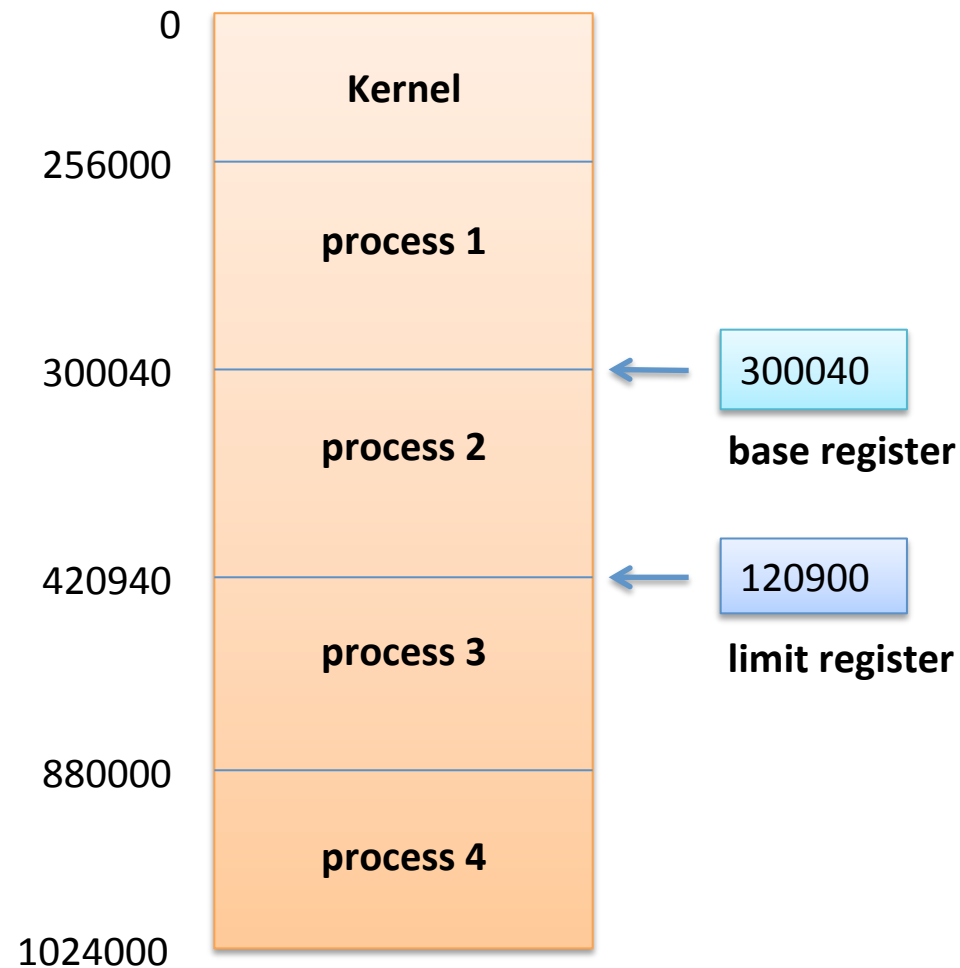
I/O Protection

- All I/O instructions are *privileged instructions*.
 - Must ensure that a user program could never gain control of the computer in **kernel** mode (i.e., a user program that, as part of its execution, stores a new address in the interrupt vector).
1. “normal” instructions, e.g., add, sub, etc.
 2. “privileged” instructions, e.g., initiate I/O switch state vectors or contexts load/save from protected memory etc.

Memory Protection

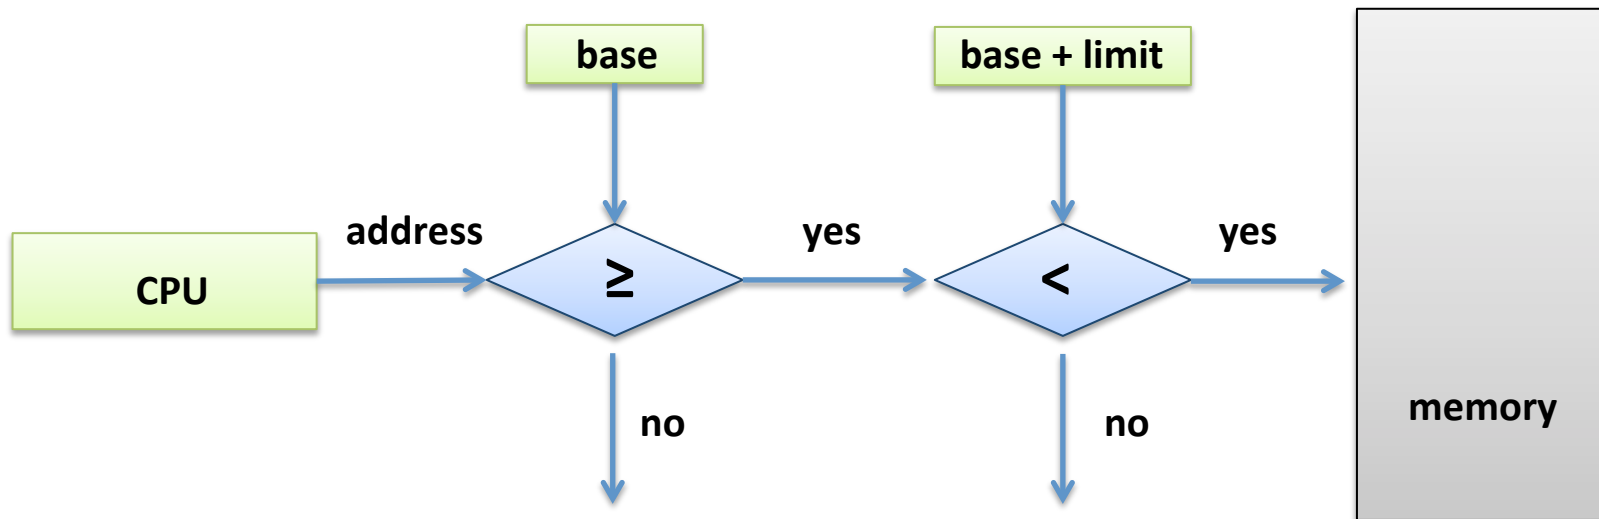
- Must provide memory protection at least for the interrupt vector and the interrupt service routines.
- In order to have memory protection, add two registers that determine the range of legal addresses a process may access:
 - **Base register** – holds the smallest legal physical memory address.
 - **Limit register** – contains the size of the range
- Memory outside the defined range is protected.

Use of a Base and Limit Registers



Hardware Protection

- When executing in kernel mode, the operating system has unrestricted access to both kernel and user's memory.
- The load instructions for the *base* and *limit* registers are **privileged instructions**.



A fault raised by hardware, notifying the operating system about an addressing error

CPU Protection

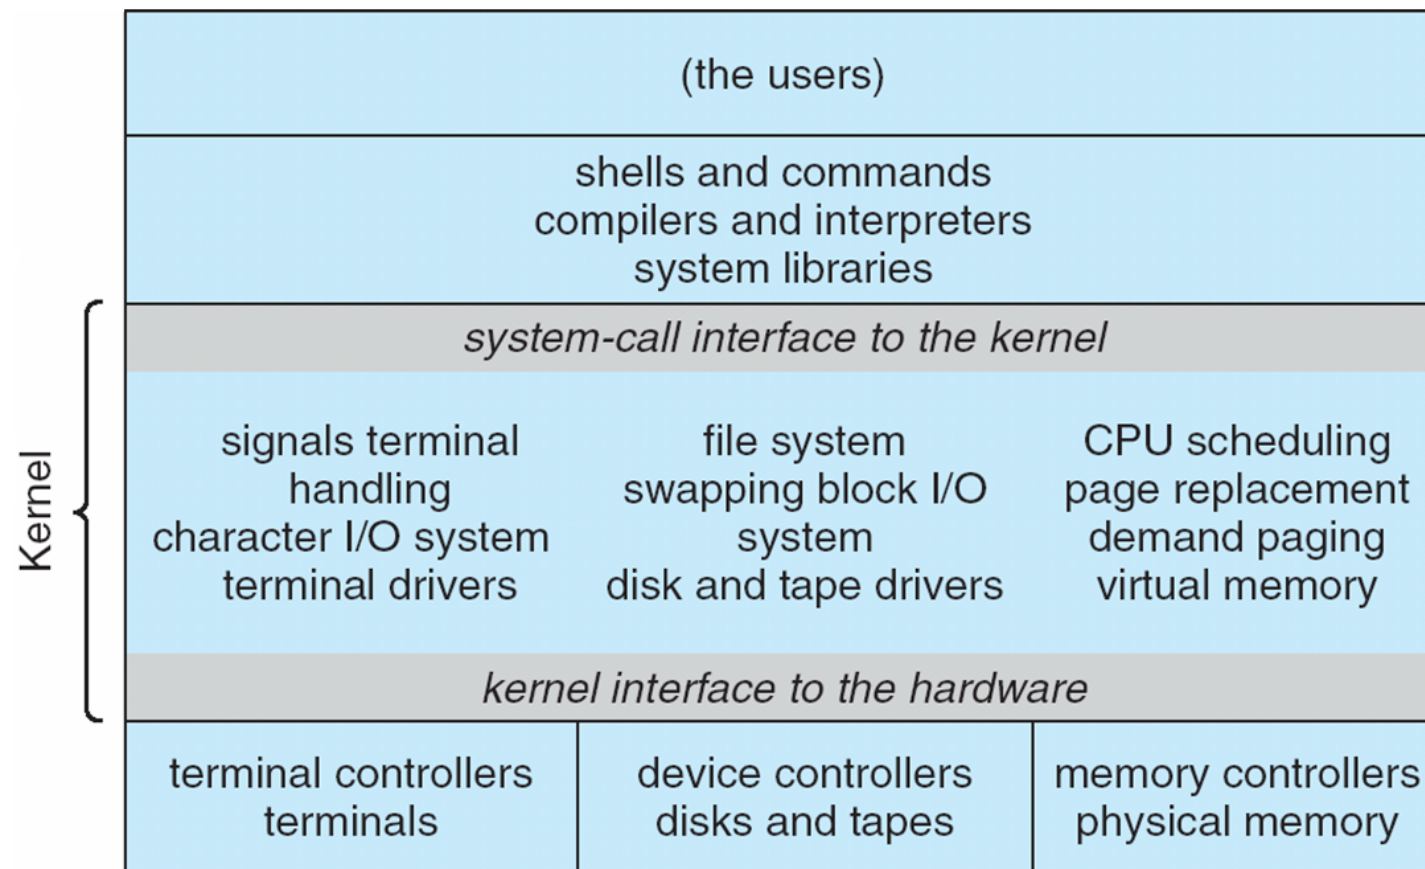
- *Timer* – interrupts computer after specified period to ensure operating system maintains control.
 - Timer is decremented every clock tick.
 - When timer reaches the value 0, an interrupt occurs.
- Timer commonly used to implement time sharing systems.
- Clearly, instructions that modify the content of the timer are privileged.

Operating System Structure

- General-purpose OS is very large program
 - Typically written in assembly, C/C++, some scripts in Perl or Python
- Various ways to structure it
 - **Monolithic Kernel:** All the OS services are implemented in the kernel. Fast OS but hard to extend
 - Ex: MS-DOS, Unix
 - **Microkernel:** Moves all the nonessential components from the kernel to user level. Smaller kernel, uses messages with system and user-level programs
 - Ex: Mach
 - **Modular Approach:** Loadable kernel modules, load additional services if needed at boot or run time
 - Ex: Solaris
- Most current OS combines all three approaches nowadays
 - Ex: Windows, Mac OS X, Linux

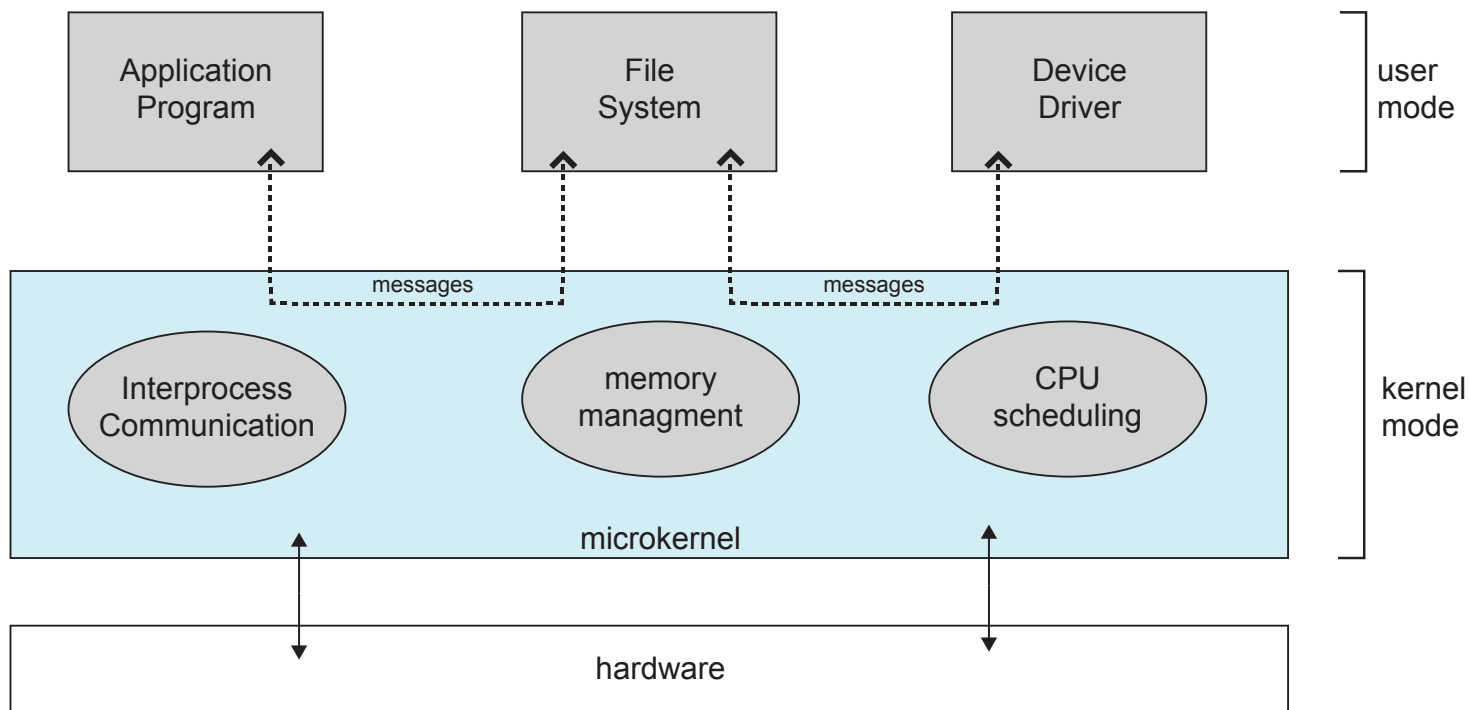
Monolithic Kernel

All the OS services are implemented in the kernel. Fast OS but hard to extend



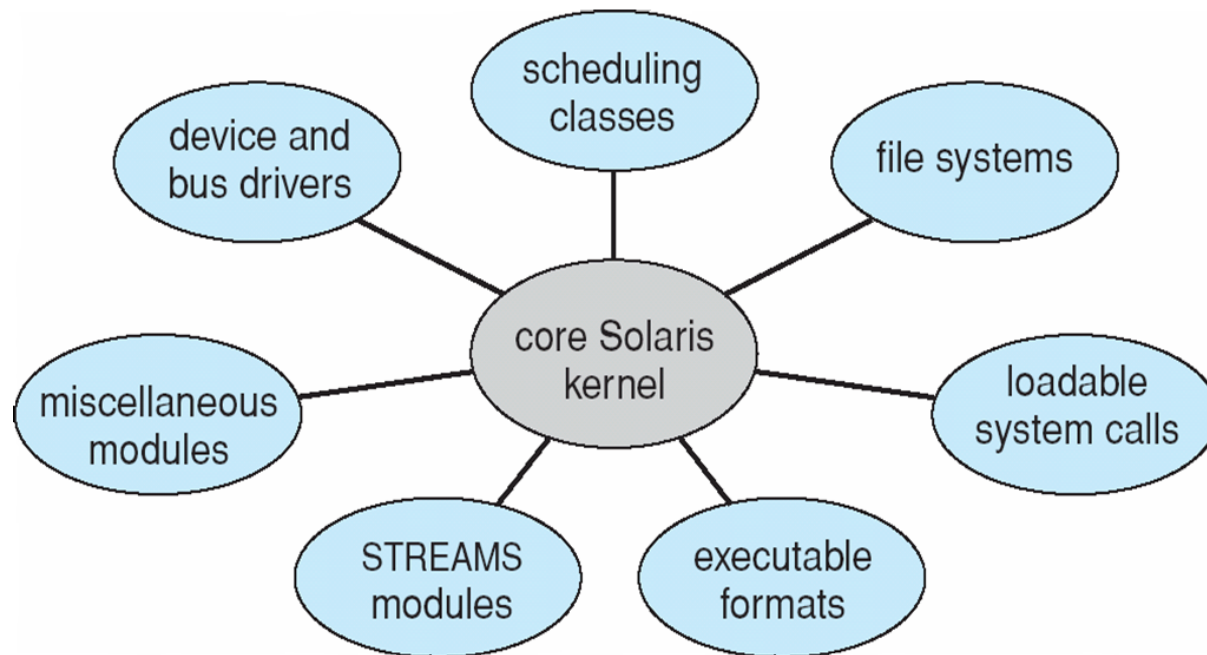
Microkernel

- Moves all the nonessential components from the kernel to user level. Smaller kernel, uses messages with system and user-level programs

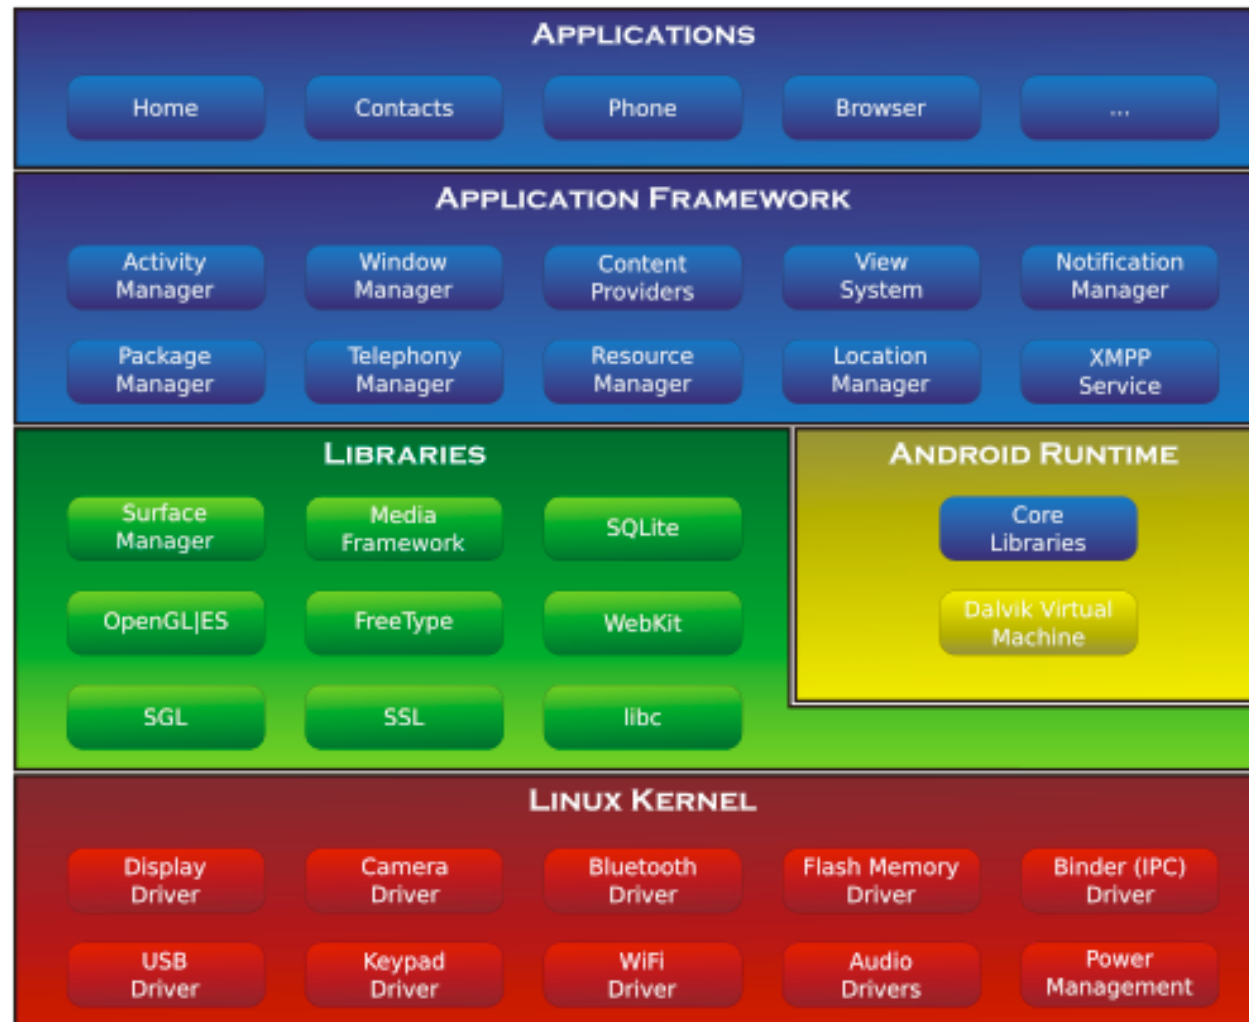


Modular Kernel

- **Modular Approach:** Loadable kernel modules, load additional services if needed at boot or run time



Android



Question

- Which of the following instructions should be privileged?
 - a. Set value of timer.
 - b. Read the clock.
 - c. Clear memory.
 - d. Issue a trap instruction.
 - e. Turn off interrupts.
 - f. Modify entries in device-status table.
 - g. Access I/O device.
- a, c, e, f, g

Question

- A ____ can be used to prevent a user program from never returning control to the operating system.

A) portal
B) program counter
C) firewall
D) Timer

D

Question

What statement concerning privileged instructions is considered false?

- A) They may cause harm to the system.
- B) They can only be executed in kernel mode.
- C) They cannot be attempted from user mode.
- D) They are used to manage interrupts.

C

Reading

- From text book
 - Read Chapter 2: Section 2.1-2.4, 2.10
- Linux System Call References
 - <http://syscalls.kernelgrok.com/>
- Acknowledgments
 - These slides are adapted from
 - Öznur Özkasap (Koç University)
 - Operating System and Concepts (9th edition) Wiley