

# Process Management

COMP304  
Operating Systems (OS)

Didem Unat  
Lecture 3

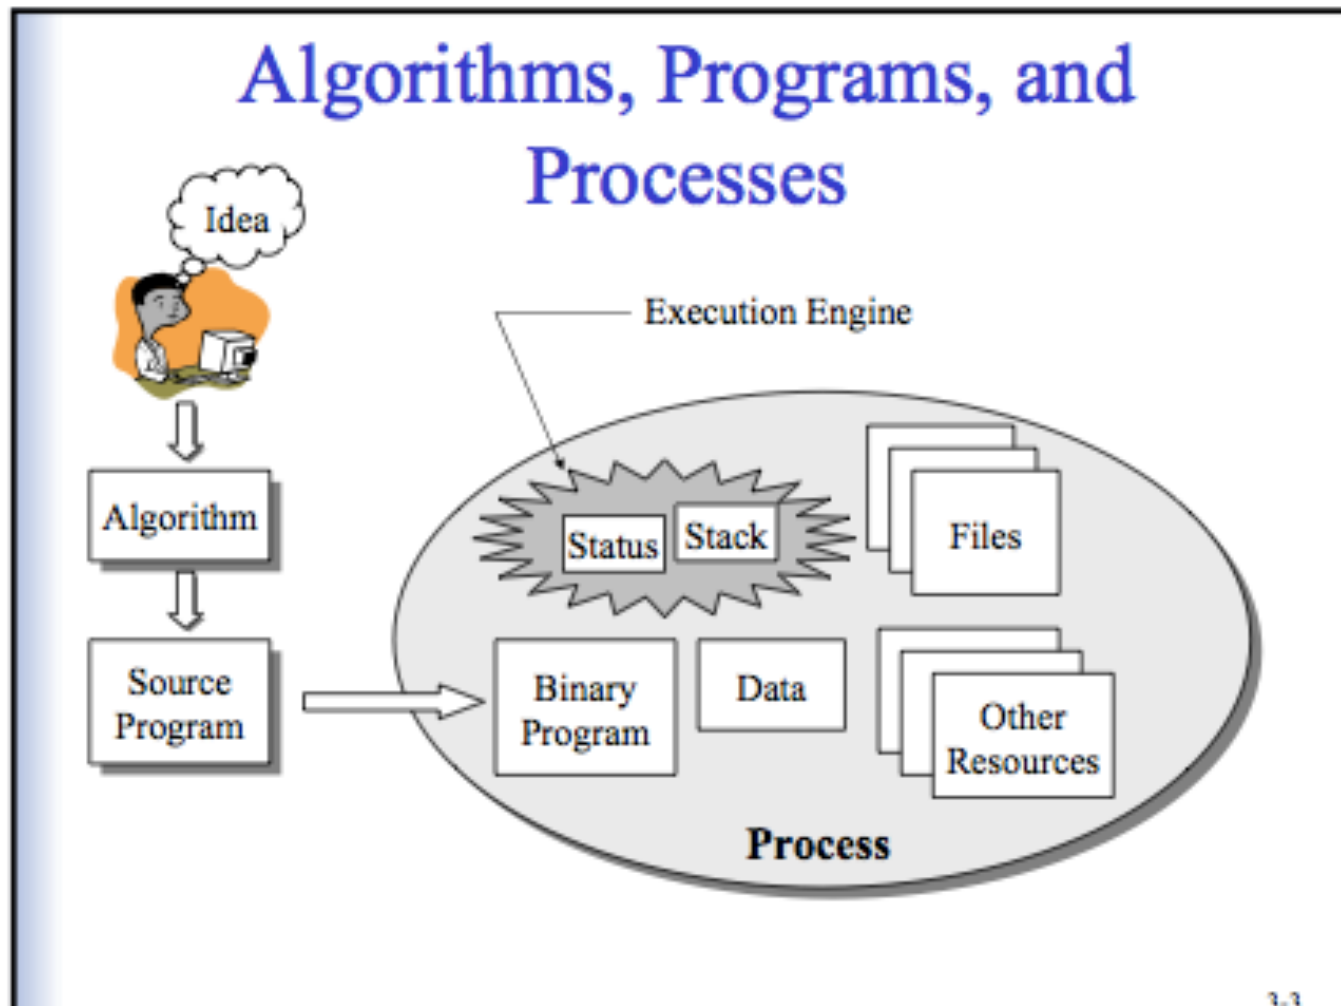
# Outline

- Process Concepts
- Process State
- Context Switch
- Schedulers
- Process Creation and Termination
- Reading
  - Chapter 3.1-3.4 from textbook – Very Good!
  - Linux Kernel Development – Chapter 3
  - HW #1 will be out this week
  - Requires Linux environment with a sudo access

# Process


- **Process** – a program in execution; process execution must progress in sequential fashion
- Program is *passive* entity stored on disk (**executable file**), process is *active*
  - Program becomes process when executable file loaded into memory
- Terms *job*, *task* and *process* are almost interchangeably used
- Execution of a program starts via GUI mouse clicks, command line entry of its name, etc
- One program can have several processes
  - Consider multiple users executing the same program
  - Ex. Multiple browsers running at the same time

# Algorithm, Program and Process



# Process

- The program code, also called **text section**
  - **Binary code**
- Current activity includes **program counter** and other processor **registers**
- **Stack** containing temporary data
  - Function parameters, return addresses, local variables
- **Data section** containing global variables
- **Heap** containing memory dynamically allocated during run time



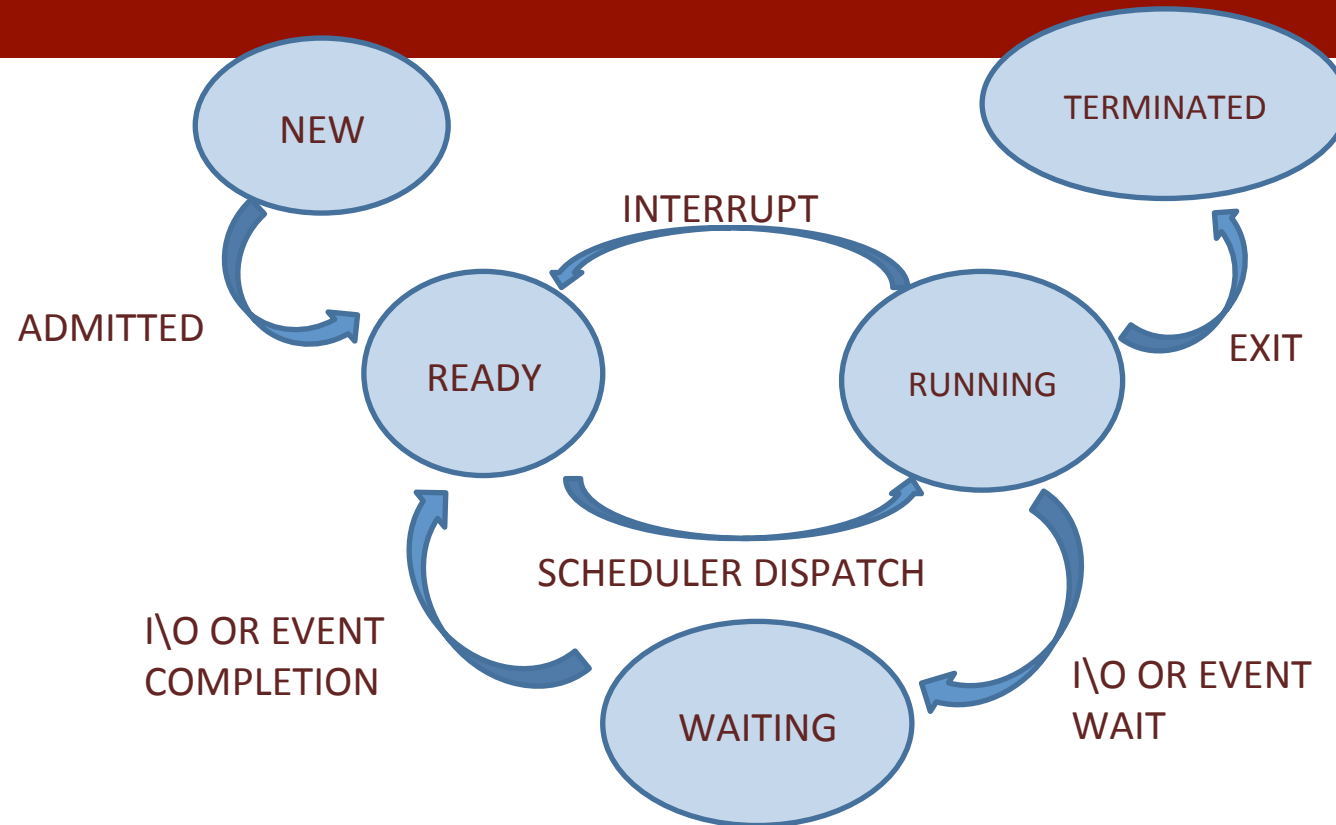
# Concurrent Execution

- OS implements an abstract machine per process
- *Multiprogramming* enables
  - $N$  programs to be *space-muxed* in executable memory, and *time-muxed* across the physical machine processor.
- Result: Have an environment in which there can be multiple programs in execution *concurrently*\*, each as a process
  - Concurrently means processes appear to execute simultaneously, they all make some progress over time

# Process State

- As a process executes, it changes its **state**
  - **new**: The process is being created
  - **running**: Instructions are being executed
  - **waiting**: The process is waiting for some event to occur
  - **ready**: The process is waiting to be assigned to a CPU
  - **terminated**: The process has finished execution

# Transition between Process States



- Process transitions from one state to another
- An animation for process states
  - <http://williamstallings.com/OS/Animation/Queensland/PROCESS.SWF>



# Process Context

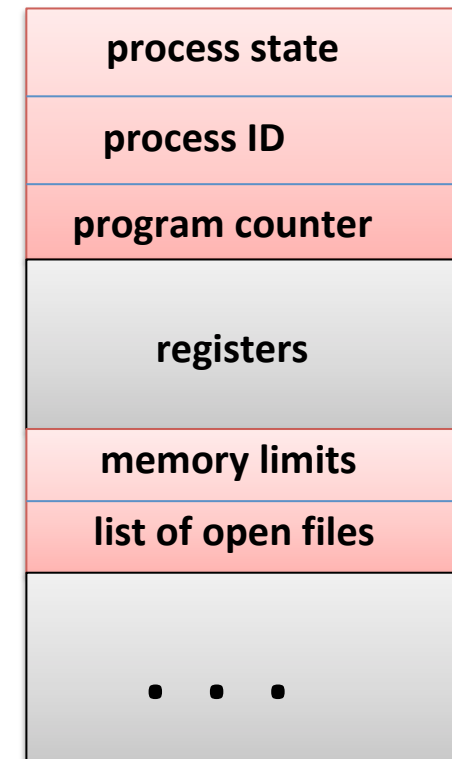
- Also called **process control block**
- When an interrupt occurs, what information OS needs to keep around so that we can reconstruct process's context as if it was never interrupted its execution?

# Process Control Block (PCB)

## Keeps the process context

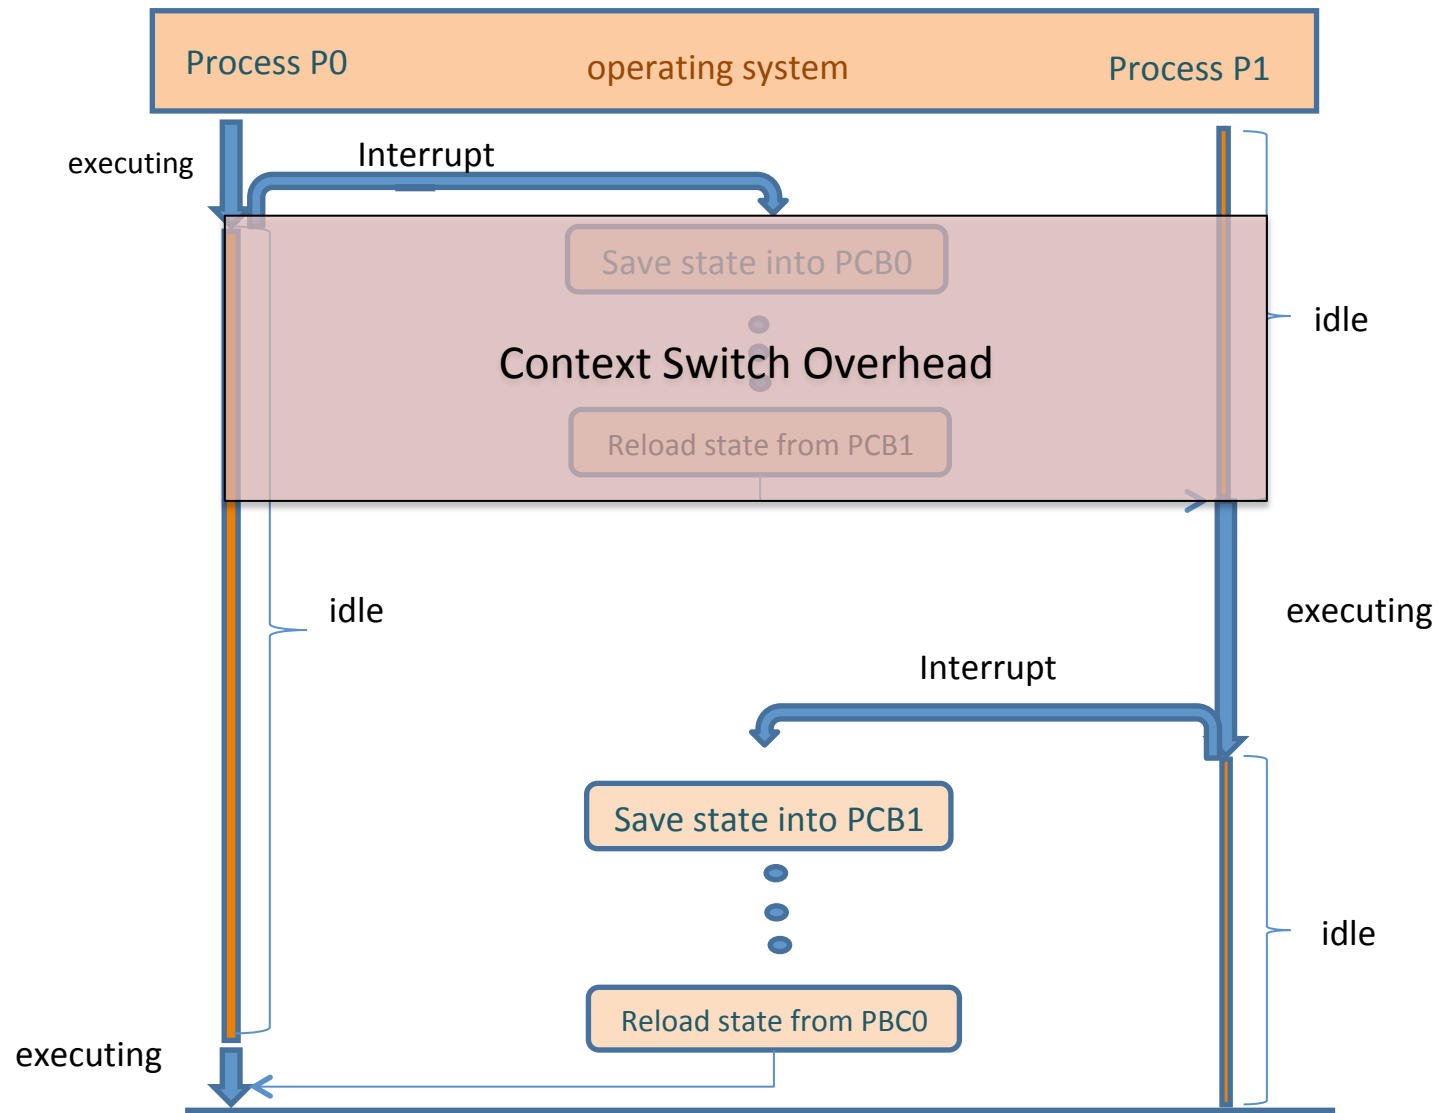
- **Process state** – running, waiting, etc
- **Program counter** – location of instruction to next execute
- **CPU registers** – contents of all process registers
- **CPU scheduling information**- priorities, scheduling queue pointers
- **Memory-management information** – memory allocated to the process
- **Accounting information** – CPU used, clock time elapsed since start, time limits
- **I/O status information** – I/O devices allocated to process, list of open files

Metadata about a process



# Context Switch

- OS needs to store and restore the context of a process
  - So that execution of the process can be resumed from the same point at a later time
  - This is called **context switch**
- Context switch is **pure overhead**
  - Why?
  - Should be very small (couple millisecond)
  - Hardware support for context switching improves the performance
- When does OS switch context?
  - In case of an interrupt
  - When process's time is up
    - Even though process has still some work to do
  - When a process terminates



- Switching between threads of a single process can be faster than between two separate processes

# Threads

- A process has at least one thread of execution
- Consider having multiple program counters per process
  - Multiple locations can execute at once
  - Word document
    - Spell checker – 1 thread
    - Typing text – 1 thread
- Must then have storage for thread details, multiple program counters in PCB or each thread has a PCB
- More on threads later (in Chapter 4)

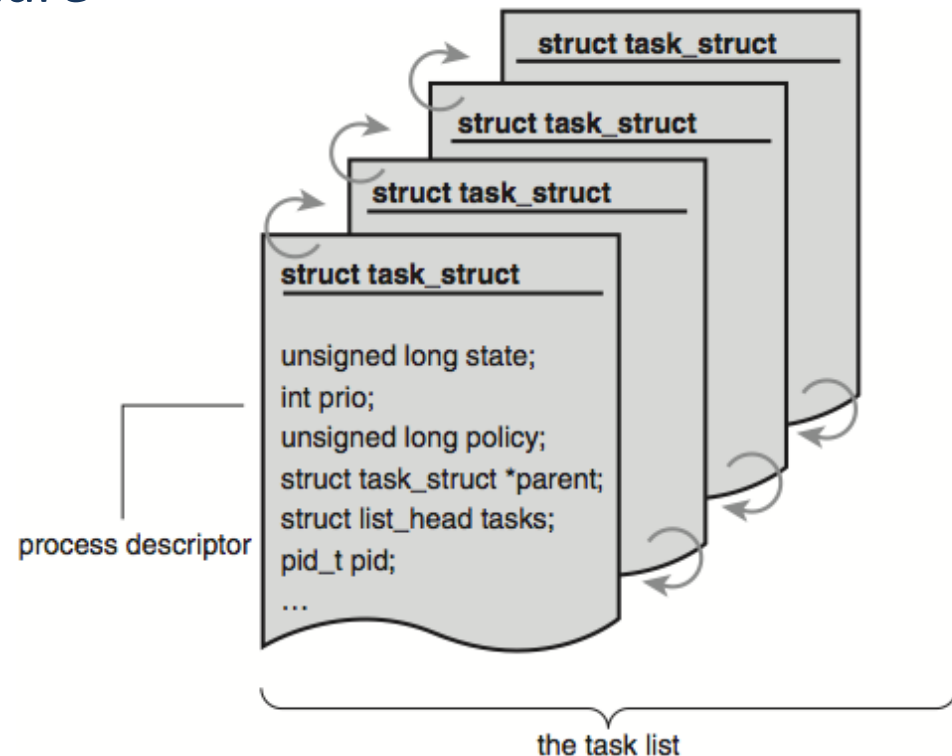
# Unix Processes

- Each process has its own address space
  - Subdivided into text, data, & stack segment – a.out file describes the address space
- OS kernel creates *descriptor (PCB)* to manage process
- *Process identifier (PID)*: User handle for the process (descriptor)

# task\_struct

- Represented by the C structure

```
task_struct {  
    pid_t pid;  
    /* process identifier */  
    long state;  
    /* state of the process */  
    unsigned int time_slice  
    /* scheduling information */  
    struct task_struct *parent;  
    /* this process's parent */  
    struct list_head children;  
    /* this process's children */  
    struct files_struct *files;  
    /* list of open files */  
    struct mm_struct *mm;  
    /* address space of this process */  
    ...  
}
```

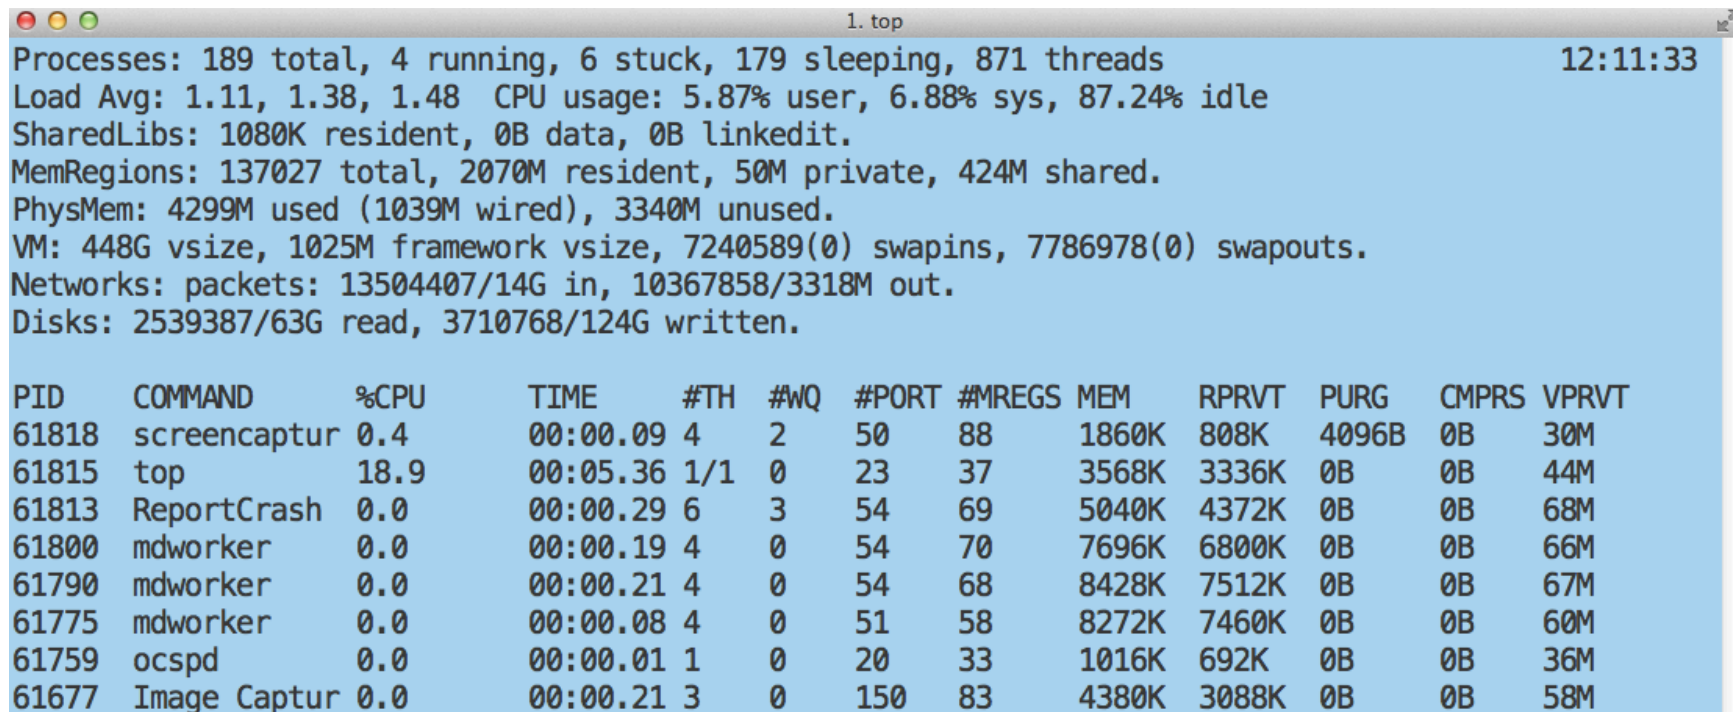


Search for task\_struct (line ~1200)

<http://lxr.free-electrons.com/source/include/linux/sched.h>

# 'top' and 'ps' commands

- top: Displays processor activity of POSIX-based OS and also displays tasks managed by kernel in real-time.
- ps: snapshot of process states



```
1. top
Processes: 189 total, 4 running, 6 stuck, 179 sleeping, 871 threads      12:11:33
Load Avg: 1.11, 1.38, 1.48  CPU usage: 5.87% user, 6.88% sys, 87.24% idle
SharedLibs: 1080K resident, 0B data, 0B linkedit.
MemRegions: 137027 total, 2070M resident, 50M private, 424M shared.
PhysMem: 4299M used (1039M wired), 3340M unused.
VM: 448G vsize, 1025M framework vsize, 7240589(0) swapins, 7786978(0) swapouts.
Networks: packets: 13504407/14G in, 10367858/3318M out.
Disks: 2539387/63G read, 3710768/124G written.

PID    COMMAND      %CPU    TIME    #TH    #WQ    #PORT  #MREGS  MEM    RPRVT  PURG    CMPRS  VPRVT
61818  screencaptur 0.4      00:00.09 4      2      50     88      1860K  808K   4096B   0B     30M
61815  top           18.9     00:05.36 1/1    0      23     37      3568K  3336K  0B      0B     44M
61813  ReportCrash  0.0      00:00.29 6       3      54     69      5040K  4372K  0B      0B     68M
61800  mdworker     0.0      00:00.19 4       0      54     70      7696K  6800K  0B      0B     66M
61790  mdworker     0.0      00:00.21 4       0      54     68      8428K  7512K  0B      0B     67M
61775  mdworker     0.0      00:00.08 4       0      51     58      8272K  7460K  0B      0B     60M
61759  ocspd        0.0      00:00.01 1       0      20     33      1016K  692K   0B      0B     36M
61677  Image Captur 0.0      00:00.21 3       0     150     83      4380K  3088K  0B      0B     58M
```

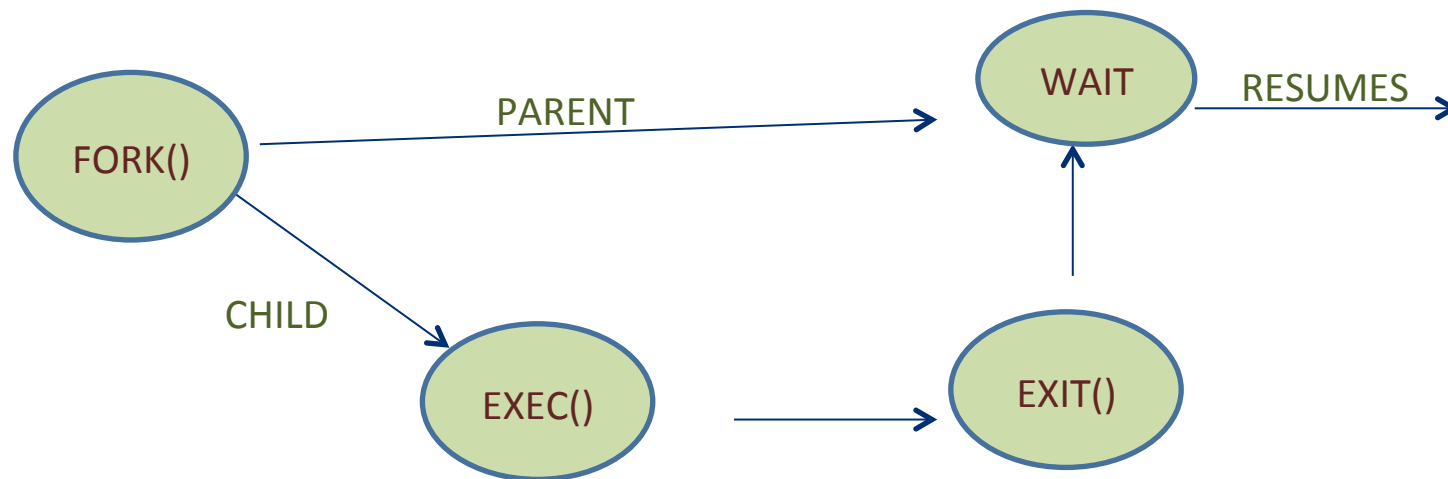


# Creating/Destroying Processes

- UNIX `fork ( )` creates a process
  - Creates a new address space
  - Copies text, data, & stack into new address space
  - Provides child with access to open files of its parent
- UNIX `wait ( )` allows a parent to wait for a child to change its state
  - This is a blocking call, parent waits until it receives a signal
  - <http://linux.die.net/man/2/wait>
- UNIX `exec ( )` system call variants (e.g. `execve()`) allow a child to run a new program

# Process Creation

- Address space
  - Child duplicates the address space of the parent
  - Child has a program loaded into it
- UNIX examples
  - **fork()** system call creates a new process
  - **exec()** system call is used after a **fork()** to replace the process' memory space with a new program



# Creating a UNIX process

```
int value, mypid=-1;
...
value = fork(); /* Creates a child process */

/* value is 0 for child, nonzero for parent */
if(value == 0) {
    /* The child executes this code concurrently with parent */
    mypid = getpid();
    printf("Child's Process ID: %d\n", mypid);
    exit(0);
}
/* The parent executes this code concurrently with child */

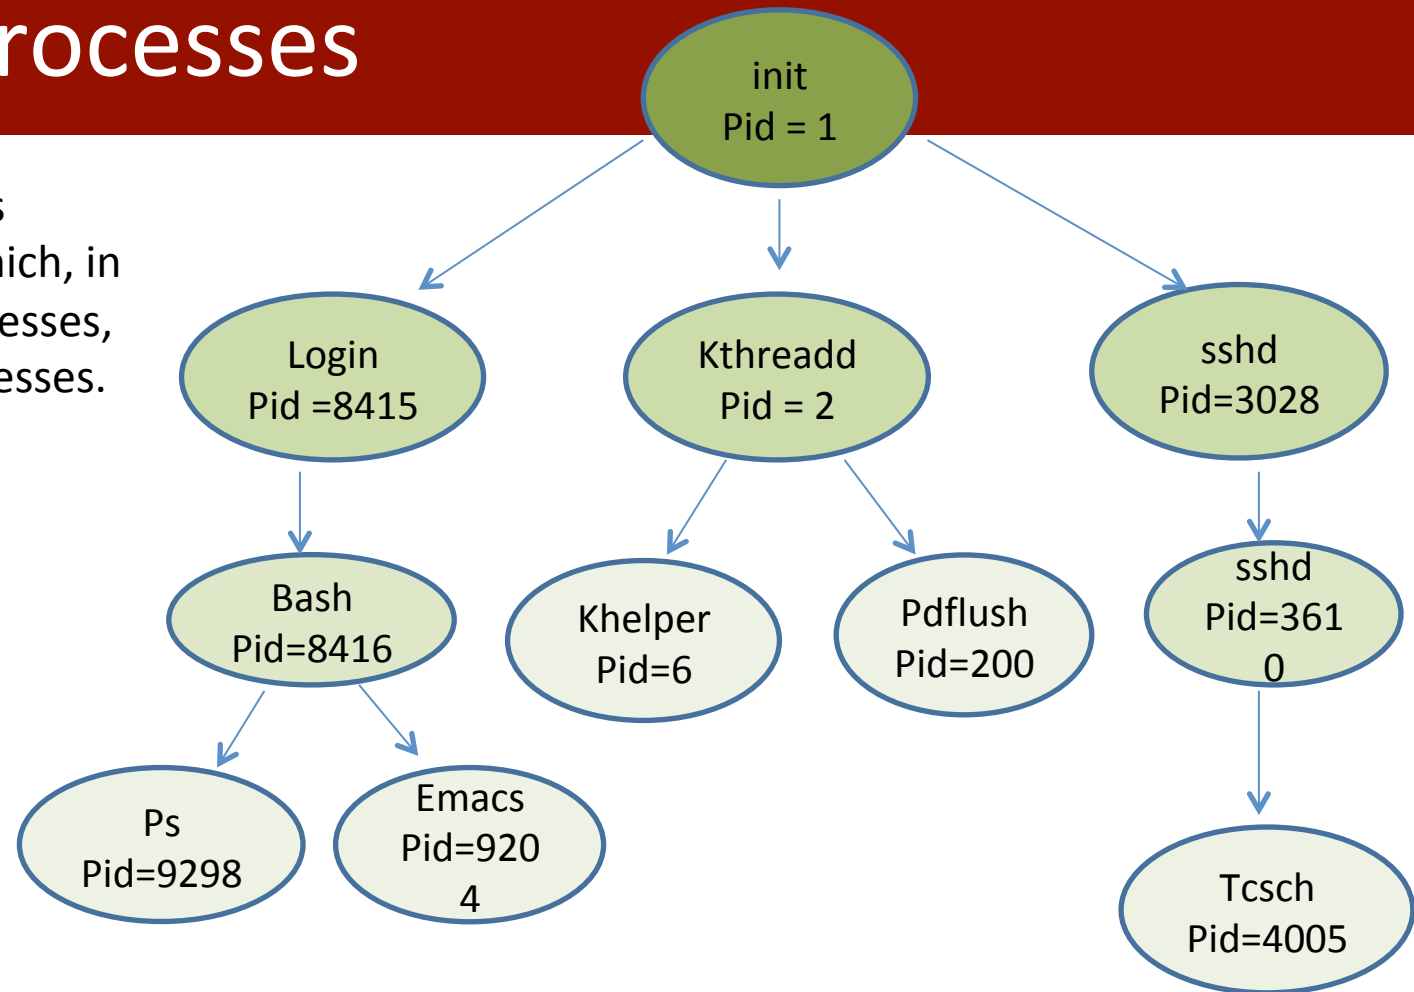
parentWorks(..);
wait(...);
...
```

# Child Executing a Different Program

```
int mypid;
...
/* Set up the argv array for the child */
...
/* Create the child */
if((mypid = fork()) == 0) {
    /* The child executes its own absolute program */
    execve("childProgram.out", argv, 0);
    /* Only return from an execve call if it fails */
    printf("Error in the exec ... terminating the child ...");
    exit(0);
}
...
wait(...);    /* Parent waits for child to terminate */
...
```

# Tree of Processes

Parent process creates children processes, which, in turn create other processes, forming a tree of processes.



- init is very first process (pid =1)
- kthread is for system processes (pid=2)
- login process is for users directly logged in to the system
- sshd process is for users remotely logged in to the system
  - Starts an openSSH SSH daemon

```

#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>
#define SIZE 5

int nums[SIZE] = {0,1,2,3,4};

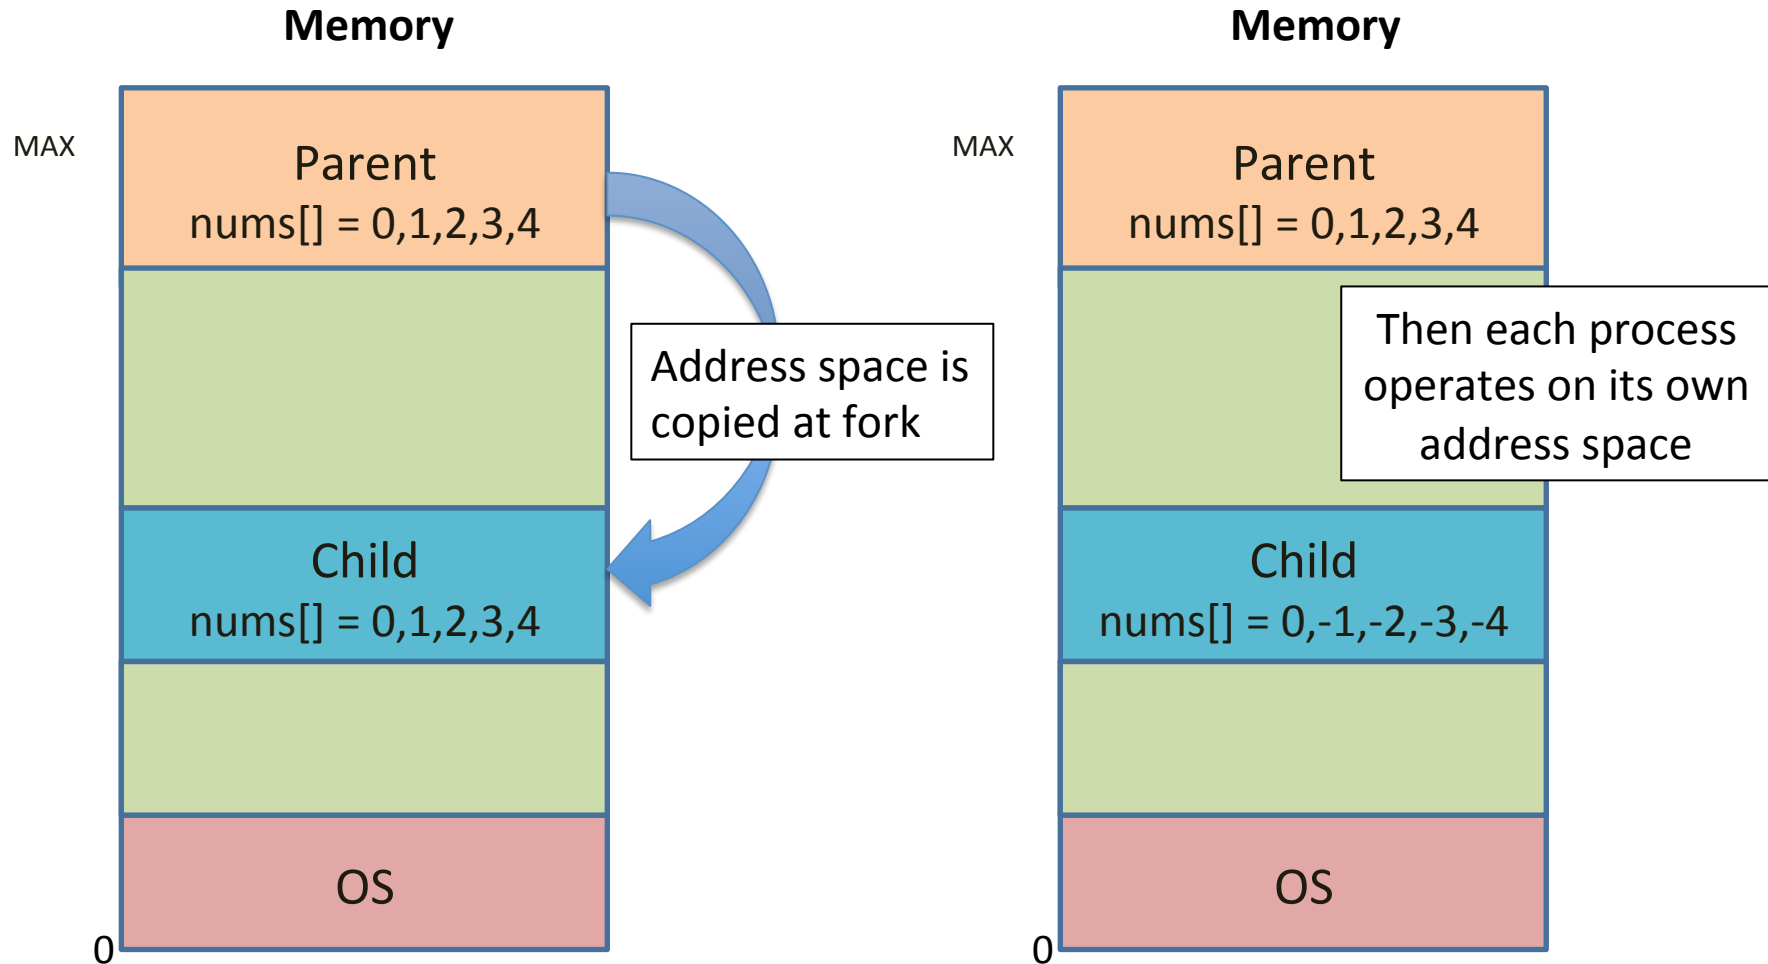
int main()
{
    int i;
    pid_t pid;
    pid = fork();

    if (pid == 0) {
        for (i = 0; i < SIZE; i++) {
            nums[i] = -i;
            printf("CHILD: %d \n",nums[i]); /* LINE X */
        }
    }
    else if (pid > 0) {
        wait(NULL);
        for (i = 0; i < SIZE; i++)
            printf("PARENT: %d \n",nums[i]); /* LINE Y */
    }
    return 0;
}

```

What output will be at  
Line X and Line Y?

# Address Spaces



# Reading

- From text book
  - Read Chapter 3.1-3.3
  - Linux Kernel Development (Chapter 3)
- Acknowledgments
  - These slides are adapted from
    - Öznur Özkasap (Koç University)
    - Operating System and Concepts (9<sup>th</sup> edition) Wiley