

# Real-Time CPU Scheduling and Linux Scheduler

Didem Unat  
Lecture 7

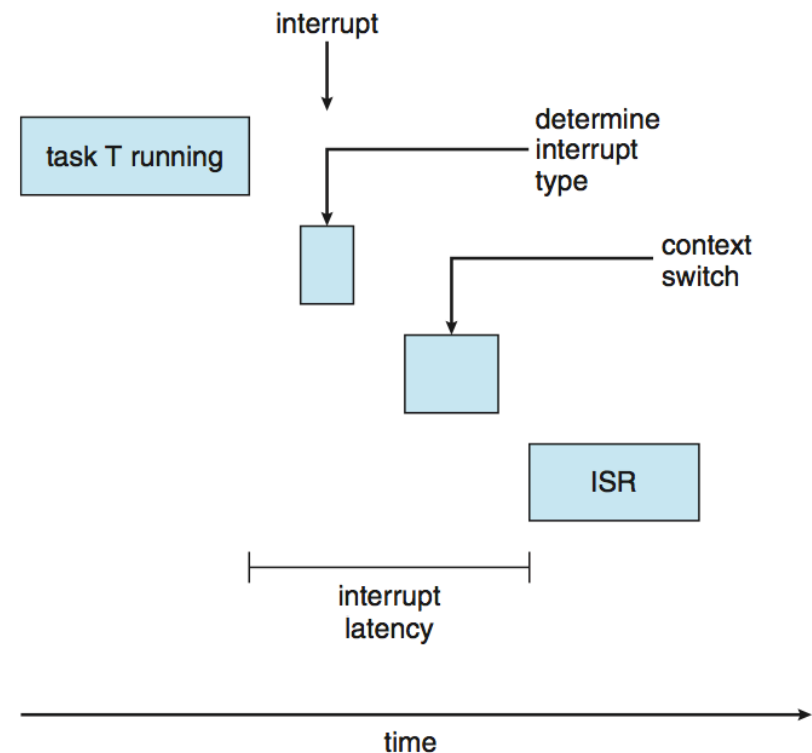
COMP304 - Operating Systems (OS)

# Real-Time CPU Scheduler

- **Real-time programs** must guarantee response within strict time constraints, often referred to as deadlines
- **Soft real-time systems** – no guarantee as to when critical real-time process will be scheduled, degrades the system's quality of service
  - Ex: updates the flight plans for an airline, live broadcasting
- **Hard real-time systems** – missing a deadline is a total system failure
  - Mission critical: a real-time deadline must be met, regardless of system load
  - Ex: Anti-lock brakes on a car, heart pacemakers and many medical devices
- Not all the Operating Systems are real-time operating systems

# Real-Time CPU Scheduling

- **Event latency:** time that elapses from when an event occurs to when it is serviced
- Two types of latencies affect performance
  1. **Interrupt latency** – time from arrival of interrupt to start of routine that services the interrupt
  2. **Dispatch latency** – time for schedule to take current process off CPU and switch to another



# Real-Time OSs

- Event driven systems switch between tasks based on their priorities while time sharing systems switch the task based on clock interrupts.
- Design goal is not high throughput, but rather a guarantee of service for a high priority job
- Real-time OS is more frequently dedicated to a narrow set of applications.
  - Targeted usages is typically embedded systems, robots
- Some open source real-time OSs:
  - uKOS
  - [Apache Mynewt OS](#)
  - Atomthreads ...
- [http://www.wikiwand.com/en/Comparison\\_of\\_real-time\\_operating\\_systems](http://www.wikiwand.com/en/Comparison_of_real-time_operating_systems)

# Linux Scheduler

## History

Linux Version	Scheduler
Pre 2.5	Multi-level Feedback Queue
Pre 2.6.23	O(1) Scheduler
Post 2.6.23	Completely fair scheduler

# Basic Philosophies in Linux

- Priority is the primary scheduling mechanism
- Priority is *dynamically adjusted* at run time
- Try to distinguish **interactive** processes from **non-interactive**
- Use large quanta for important processes
  - Modify quanta based on CPU usage for the next run
- Associate processes to CPUs – process affinity

# Priority

- Each task has a **static priority** that is set based upon the nice value specified by the task.
  - *static\_prio* in *task\_struct*
  - *Default is 120*
- For normal tasks, the **static priority** is  $100 + \text{nice}$ .
- Each task has a **dynamic priority** that is set based upon a number of factors
  - *prio* in *task\_struct*

# Niceness

- Niceness
  - a process is nicer to others if it has a higher nice value
  - Default is inherited from its parent (usually 0)
  - Ranges from -20 to +19
- Nice value of -20 maps to global priority 100
- Nice value of +19 maps to priority 139

Value can be set via **nice()** system call or **nice** command

```
bash$ nice -n 19 tar cvzf archive.tgz largefile
```



# Prior to 2.5

In the 2.4 kernel, this was the scheduling algorithm:


- Each task got a number of CPU ticks (*jiffies*) made available to the task each scheduling interval, or *epoch*.
- The number of new ticks given was determined from the nice value for the task. It was roughly:

$$((20 - \text{nice}) * \text{HZ} / 800) + 1.$$

- Each task had a *counter*, which was the number of CPU ticks still left for the task to use in the current epoch.
- Unused ticks in a particular epoch decayed by 50% for use in the next interval.

# Linux O(1) Scheduler

- Version 2.5 moved to constant order  $O(1)$  scheduling time
  - Preemptive, priority based
  - Two priority ranges: time-sharing and real-time
  - **Real-time** range from 0 to 99 and **normal** (time-sharing) range from 100 to 140
  - Higher priority gets larger time quantum
  - Scales well with the number of processes



# Real-Time Scheduling

- Linux has a soft real-time scheduler
  - No hard real-time guarantees
  - All real-time processes are higher priority than any normal processes
- Processes with priorities [0, 99] are real-time
  - saved in *rt\_priority* in the *task\_struct*
  - scheduling priority of a real time task is:  $99 - \text{rt\_priority}$
- A process can be converted to real-time via *sched\_setscheduler* system call

# Scheduling Policies

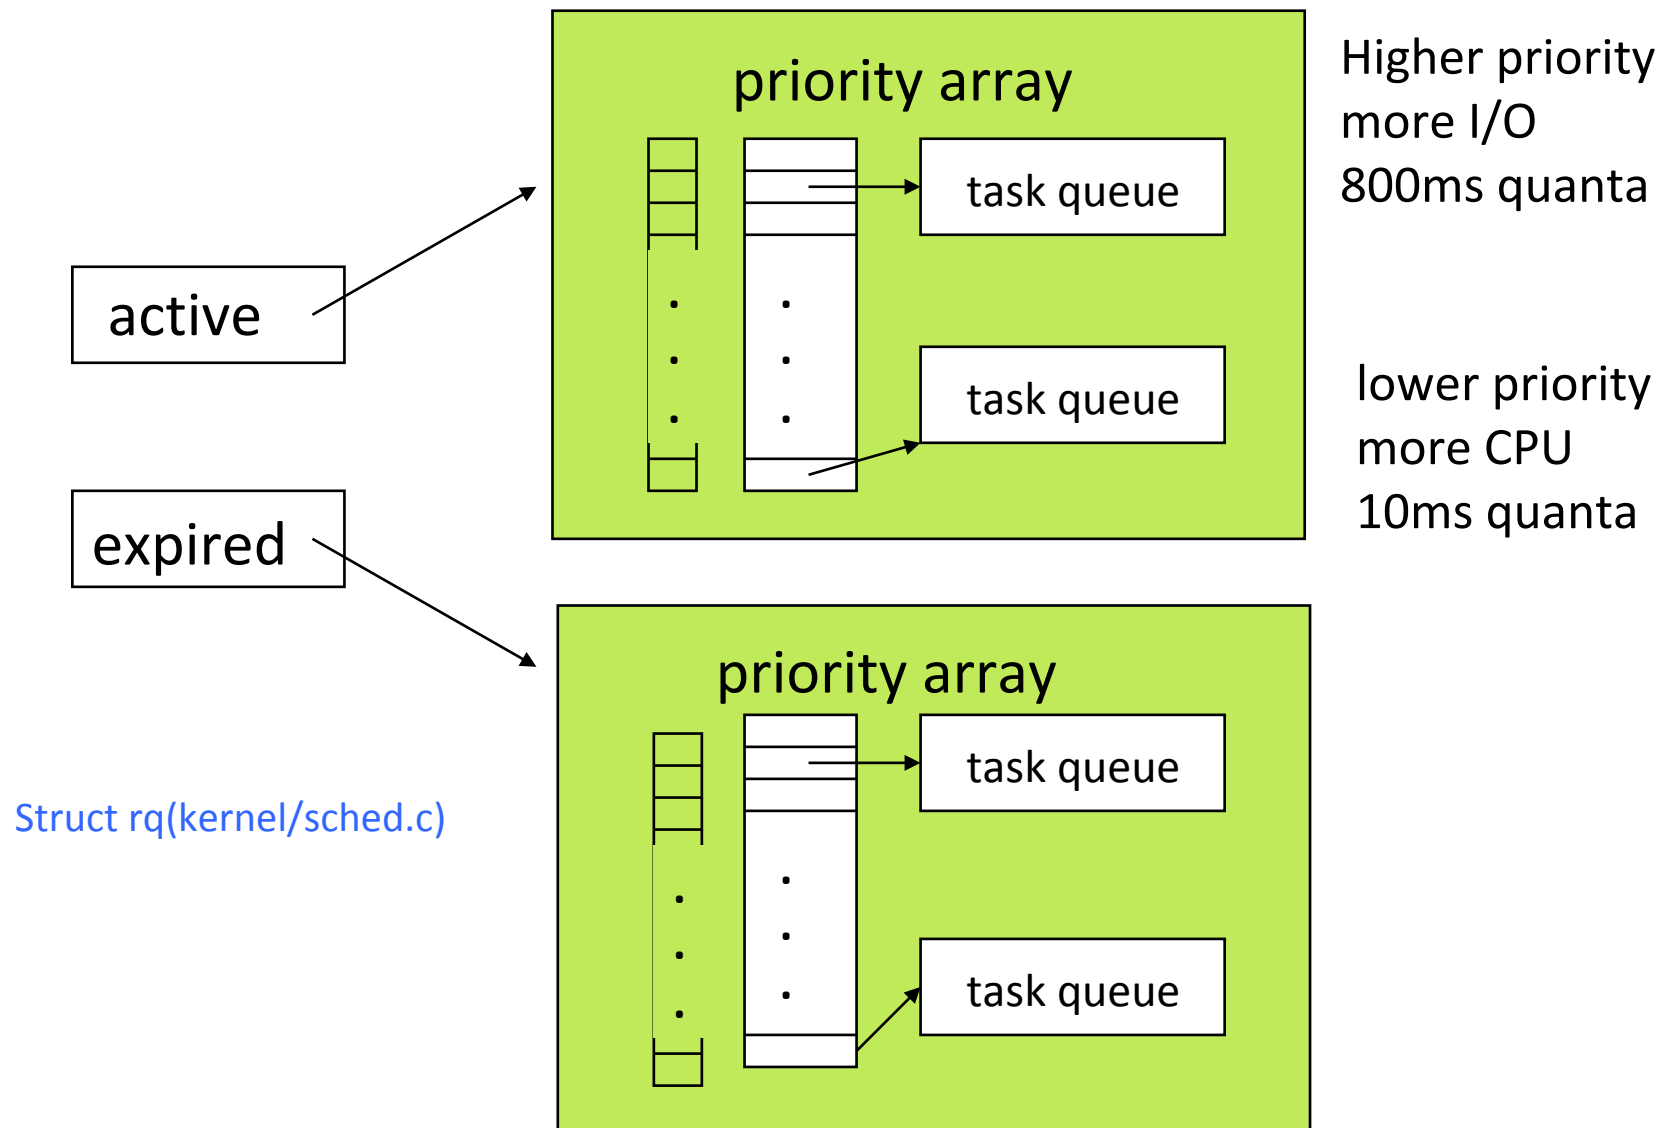
- Real-time processes
  - First-in, first-out: **SCHED\_FIFO**
    - Static priority
    - Process is only preempted for a higher-priority process
    - No time quanta; it runs until it blocks or yields voluntarily
  - Round-robin: **SCHED\_RR**
    - RR within the same priority level
    - A time quanta (800 ms)
- Normal processes have
  - **SCHED\_OTHER**: standard processes
  - **SCHED\_BATCH**: batch style processes
  - **SCHED\_IDLE**: low priority tasks

# O(1) Scheduler

- Task runnable as long as time left in time slice (**active**)
- If no time left (**expired**), not runnable until all other tasks use their slices
- All runnable tasks tracked in per-CPU **runqueue** data structure
  - Two priority arrays (active, expired)
  - When no more active, arrays are swapped

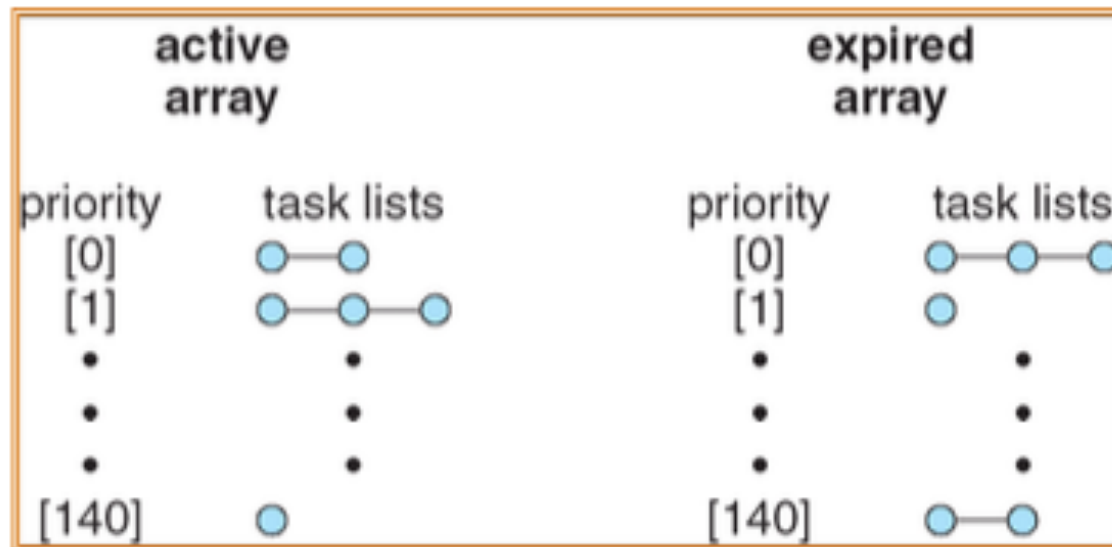
# Runqueues

140 separate queues, one for each priority level in two sets: active and expired



# Runqueues

- Two arrays of priority queues
  - **active** and **expired**
  - Total 140 priorities [0, 140)
  - Smaller integer = higher priority



# Scheduling Algorithm for Normal Processes

- Find the highest-priority non-empty queue in **rq->active**; if none, simulate aging by swapping active with expired
- **Next** = Find the first process on that queue
- Calculate **next's** quantum size and its **next's** priority
- Context switch to **next**
- Let it run
- When its time is up, put it on the **expired list**
- Repeat



# Simulate Aging

- After running all of the active queues, the active and expired queues are swapped
- There are pointers to the current arrays; at the end of a cycle, the pointers are switched
- Swapping active and expired gives low priority processes a chance to run
- Advantage:  $O(1)$ 
  - Processes are touched only when they start or stop running

# Find highest priority non-empty queue

- Time complexity  $O(1)$ 
  - Depends on the number of priority levels, not the number of processes
- Implementation: a bitmap for fast look up
  - 140 queues
  - A few comparisons to find the first non-zero bit

# Calculating Time Slices

- *time\_slice* in the *task\_struct*
- Calculate Quantum where
  - If ( $SP < 120$ ): Quantum =  $(140 - SP) \times 20$
  - if ( $SP \geq 120$ ): Quantum =  $(140 - SP) \times 5$where SP is the *static priority*
- Higher priority process gets longer quanta
- Basic idea: important processes should run longer

# Typical Quanta

Priority:	Static Pri	Niceness	Quantum
Highest	100	-20	800 ms
High	110	-10	600 ms
Normal	120	0	100 ms
Low	130	10	50 ms
Lowest	139	20	5 ms

# Issues with $O(1)$ RR Scheduler

- Not easy to distinguish between CPU and I/O bound
  - I/O bound typically needs better interactivity
- Finding right time slice isn't easy
  - Too small: good for I/O bound but high overhead
  - Too large: good for CPU bound but poor interactivity
- Priority is relative but time slice is absolute
  - Nice 0, 1: time slice 100 and 95 msec: 5% difference
  - Nice 19,20: time slice 10 and 5: 100 % difference

# Completely Fair Scheduler (CFS)

- Starting from Linux kernel version 2.6
- Not based on run queues as in  $O(1)$  scheduler
- Not based on time slices
- Note that CFS is used only for normal processes, for real-time processes, Linux still use priority based FCFS and RR schedulers

# Completely Fair Scheduler (CFS)

- Core ideas: dynamic time slice and order
- Don't use fixed time slice per task
  - Instead, fixed time slice across all tasks
  - Scheduling Latency
- Don't use round robin to pick next task
  - Pick task which has received least CPU so far
  - Equivalent to dynamic priority

# CFS

- CFS calculates how long a process should run as a function of the total number of runnable processes.
  - If there are N runnable processes, then each should be afforded  $1/N$  of the processor's time.
  - CFS adjusts the allotment by weighting each process's allotment by its nice value.
  - Small nice value => higher weight
  - Large nice value => lower weight
  - Then process's time slice is proportional to its weight divided by the total weight of all runnable processes.

$\text{Timeslice}(\text{task}) = \text{Timeslice}(t) * \text{prio}(t) / \text{Sum\_all\_t'}(\text{prio}(t'))$

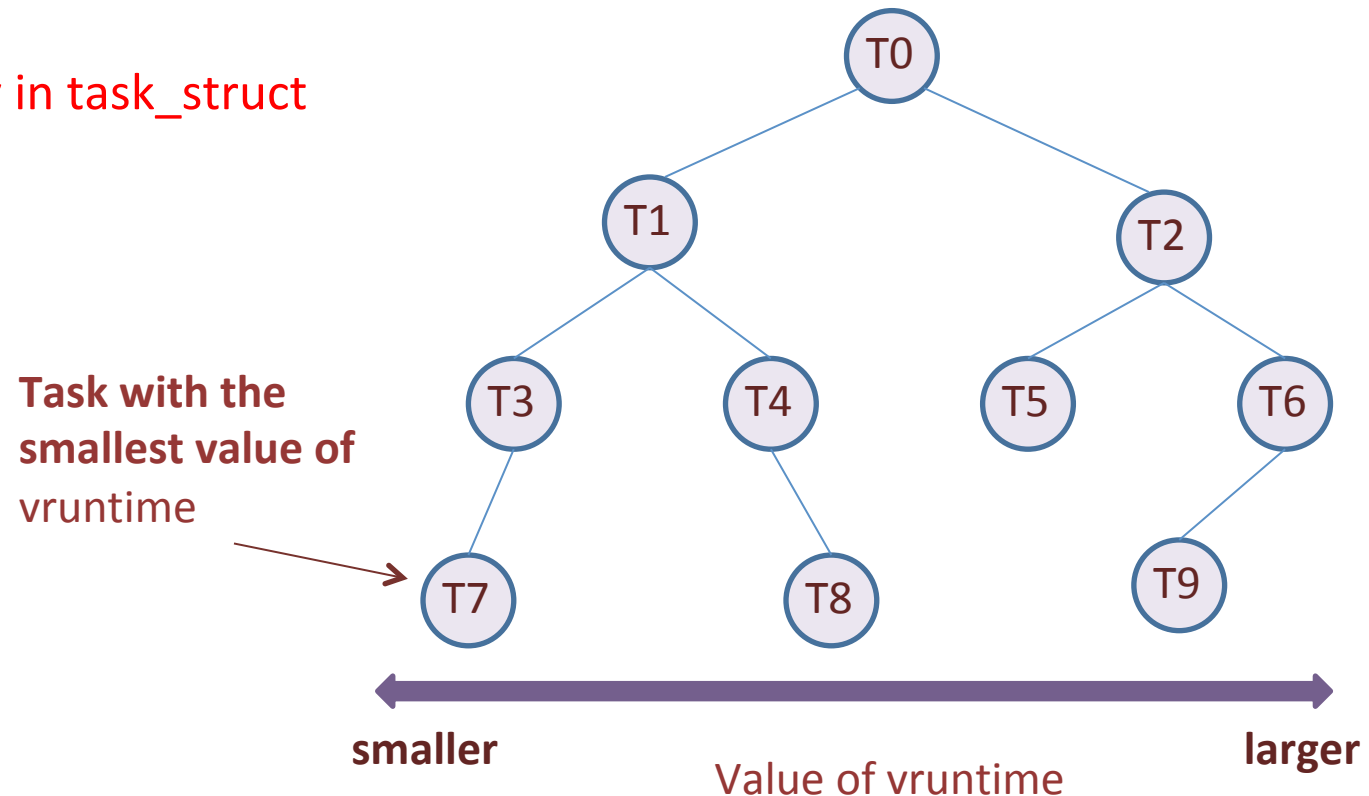
$\text{Timeslice}(t) = \text{latency} / \text{nr\_tasks}$



# CFS Tree

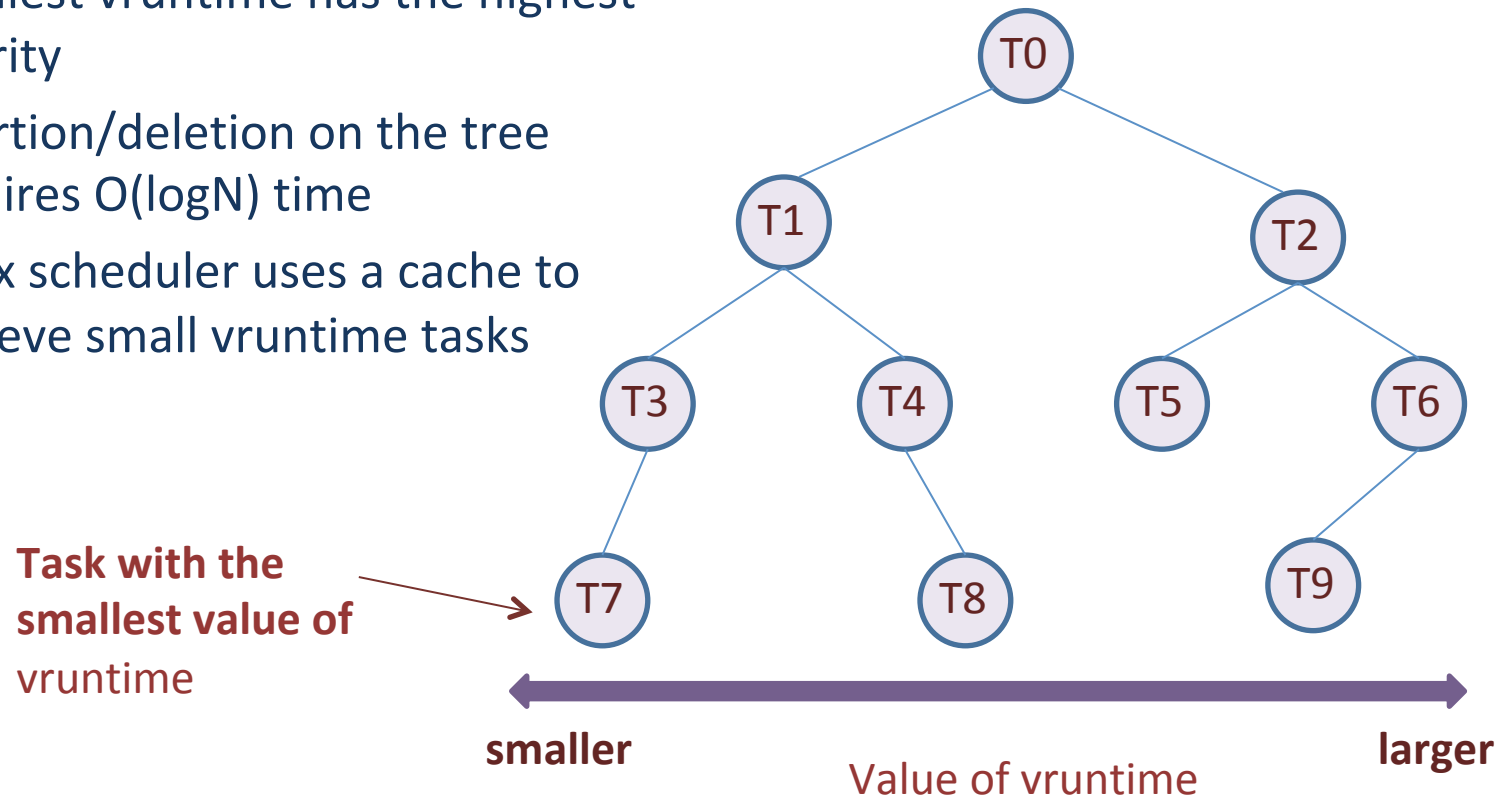
- Each runnable task is placed in a red-black tree
  - A balanced binary search tree whose key is based on the value of vruntime (task with min runtime so far)

sched\_entity in task\_struct



# CFS Tree

- When a task becomes runnable, it is added to the tree (red/black tree)
- Not runnable tasks (e.g. waiting for I/O) are removed from the tree
- Smallest vruntime has the highest priority
- Insertion/deletion on the tree requires  $O(\log N)$  time
- Linux scheduler uses a cache to retrieve small vruntime tasks



# CFS (con.t)

- Two tasks have the same nice values
- One task is I/O bound, other is CPU-bound
  - I/O bound normally runs for a short period before it is interrupted for an I/O operation
  - CPU-bound normally exhausts all its quantum
- Vruntime will eventually be lower for the I/O bound task than for the CPU-bound task
  - Vruntime is weighted by process priority

# Picking the next process

- Pick task with minimum runtime so far
- Every time process runs for  $t$  ns
  - $Vruntime += t$
- How does this impact I/O vs CPU bound tasks?
  - Task A needs 1 msec every 100 sec (I/O bound)
  - Task B, C need 80 msec every 100 msec (CPU bound)
  - After 10 times that A, B and C have been scheduled.
    - $Vruntime(A) = 10$
    - $Vruntime(B,C) = 800$
    - A gets priority but it quickly releases CPU.

# CFS Algorithm

- The leftmost node of the scheduling tree is chosen (as it will have the lowest spent *execution time*), and sent for execution.
- If the process simply completes execution, it is removed from the system and scheduling tree.
- If the process reaches its *maximum execution time* or is otherwise stopped (voluntarily or via interrupt) it is reinserted into the scheduling tree based on its new spent *execution time*.
- The new leftmost node will then be selected from the tree, repeating the iteration.

Choosing a task can be done in constant time, but reinserting a task after it has run requires  $O(\log N)$  operations

# Multiprocessor Scheduling

- Each processor maintains a red/black tree
- Each processor only selects processes from its own tree to run
- It's possible for one processor to be idle while others have jobs waiting in their run queues
- Periodically, rebalance
  - void load\_balance()!
    - Attempts to move tasks from one CPU to another

# Processor Affinity

- Each process has a bitmask saying what CPUs it can run on
- Normally, of course, all CPUs are listed
- Processes can change the mask
- The mask is inherited by child processes (and threads), thus tending to keep them on the same CPU
- not allowed to run on the current CPU (as indicated by the *cpus\_allowed* bitmask in the *task\_struct*)

# Adding a new Scheduler Class to Linux

- The Scheduler is modular and extensible
- Each scheduler class has priority within hierarchical scheduling hierarchy
  - Priorities defined in sched.h, e.g. `#define SCHED_RR 2`
  - Linked list of sched\_class sched\_class.next reflects priority
- Core functions:
  - `kernel/sched.c`, `include/linux/sched.h`
  - Additional classes: `kernel/sched_fair.c`, `sched_rt.c`
- Process changes class via
  - `sched_setscheduler` syscall
- Each class needs
  - New sched\_class structure implementing scheduling functions
  - New sched\_entity in the task\_struct



# OS Schedulers

Operating System	Preemption	Algorithm
<a href="#">Amiga OS</a>	Yes	Prioritized <a href="#">round-robin scheduling</a>
<a href="#">FreeBSD</a>	Yes	<a href="#">Multilevel feedback queue</a>
<a href="#">Linux kernel</a> before 2.6.0	Yes	<a href="#">Multilevel feedback queue</a>
Linux kernel 2.6.0–2.6.23	Yes	<a href="#">O(1) scheduler</a>
Linux kernel after 2.6.23	Yes	<a href="#">Completely Fair Scheduler</a>
<a href="#">classic Mac OS</a> pre-9	None	<a href="#">Cooperative scheduler</a>
<a href="#">Mac OS 9</a>	Some	Preemptive scheduler for MP tasks, and cooperative for processes and threads
<a href="#">macOS</a>	Yes	<a href="#">Multilevel feedback queue</a>
<a href="#">NetBSD</a>	Yes	<a href="#">Multilevel feedback queue</a>
<a href="#">Solaris</a>	Yes	<a href="#">Multilevel feedback queue</a>
<a href="#">Windows 3.1x</a>	None	<a href="#">Cooperative scheduler</a>
<a href="#">Windows 95</a> , <a href="#">98</a> , <a href="#">Me</a>	Half	Preemptive scheduler for 32-bit processes, and cooperative for 16-bit processes
<a href="#">Windows NT</a> (including 2000, XP, Vista, 7, and Server)	Yes	<a href="#">Multilevel feedback queue</a>

# Reading

- Read Chapter 16.5 Linux Scheduling
- Read Chapter 5
- Read Chapter 4 (Linux Kernel Development)
- Acknowledgments
  - These slides are adapted from
    - Öznur Özkasap (Koç University)
    - Operating System and Concepts (9<sup>th</sup> edition) Wiley
  - Linux Scheduling
    - Linux Overview. COMS W4118 Spring 2008 slideserve.com
    - Prof. Kaustubh R. Joshi from Columbia
  - <http://www.algorithmsandme.com/2014/03/scheduling-o1-and-completely-fair.html#.VPgpbMZLOWc>