

# Paging

Didem Unat

Lecture 17

COMP304 - Operating Systems (OS)

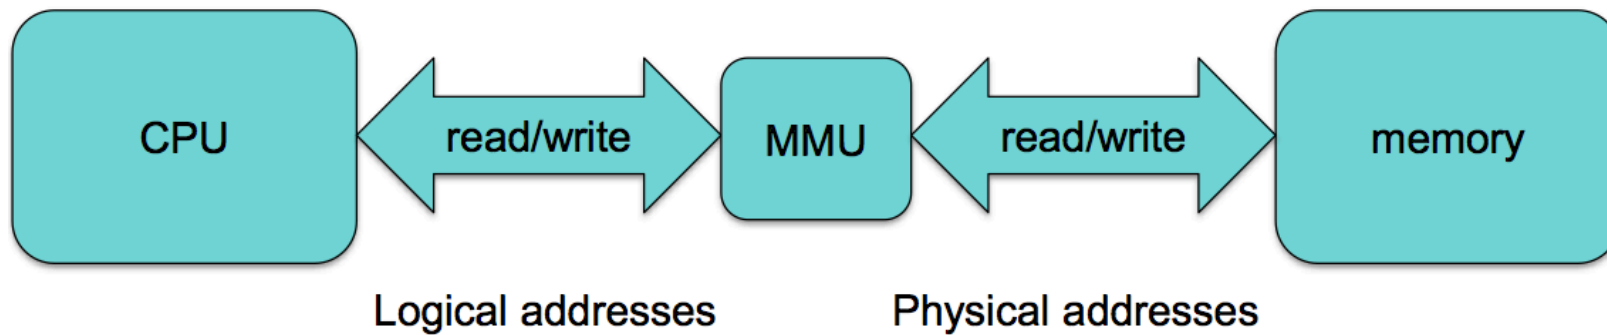
# Logical vs. Physical Address Space

- The concept of a **logical address space** that is bound to a separate **physical address space** is central to proper memory management.
  - **Logical address** – generated by the CPU; also referred to as **virtual address**.
  - **Physical address** – address seen by the memory unit.
- The user program deals with **logical** addresses; it never sees the **real physical** addresses.

# Logical Addressing

Memory management unit (MMU):

- Real-time, on-demand translation between *logical* (virtual) and *physical* addresses

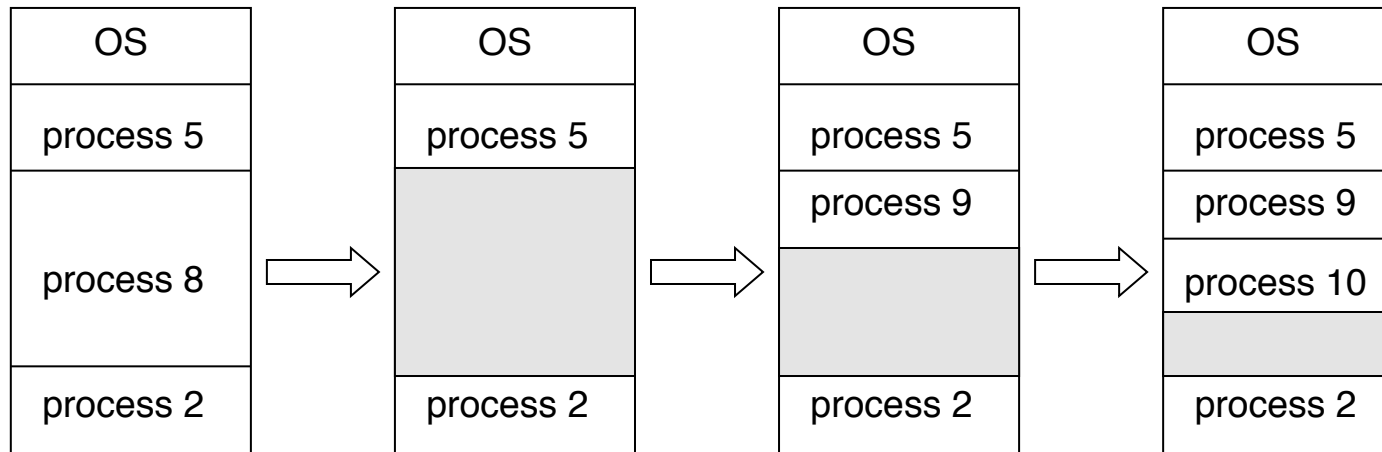


# Memory Allocation

- Main memory must support both OS and user processes
- Limited resource, must allocate efficiently
- Three methods:
  - Contiguous memory allocation
  - Segmentation
  - Paging

# 1. Contiguous Allocation

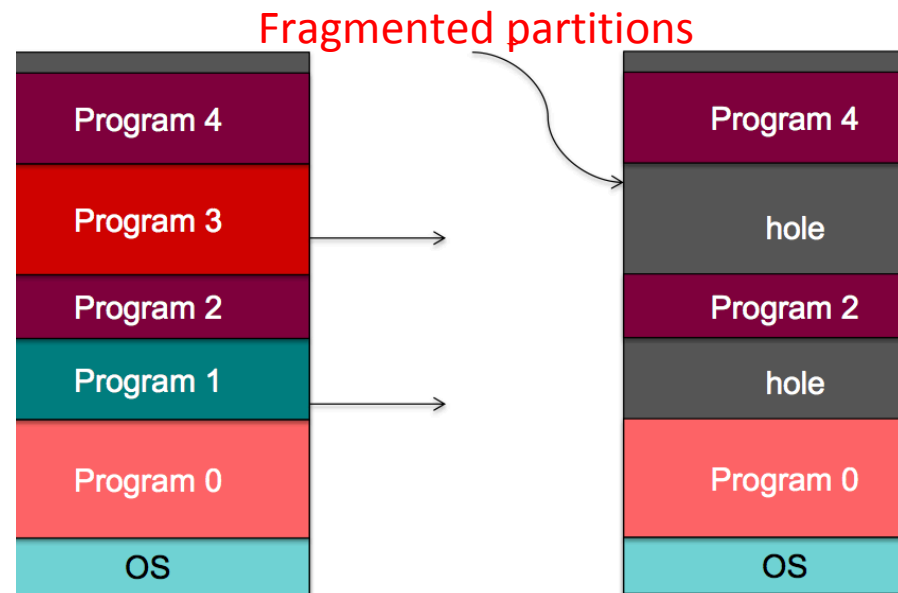
- Multiple-partition allocation
  - **Hole** – block of available memory; holes of various sizes are scattered throughout memory.
  - When a process arrives, it is allocated memory from a hole large enough to accommodate it.
  - Operating system maintains information about:  
allocated partitions and free partitions (holes)



# Fragmentation

- **External Fragmentation**

- total memory space exists to satisfy a request, but it is not contiguous
- Also a common problem in disk as well




- **Internal Fragmentation**

- Allocated memory may be slightly larger than requested memory; this size difference is memory internal to a partition, but not being used.

## 2. Segmentation

- Memory allocation mechanism that supports user view of memory.
- Users prefer to view memory as a collection of **variable-sized segments** – similar to programmer's view of memory
- A program is a collection of segments. A segment is a logical unit such as:
  - main program,
  - function,
  - method,
  - object,
  - local variables, global variables,
  - common block,
  - stack,
  - symbol table, arrays

# User's View of a Program




logical address space


Logical address space is a collection of segments



# Logical View of Segmentation



user space



physical memory space

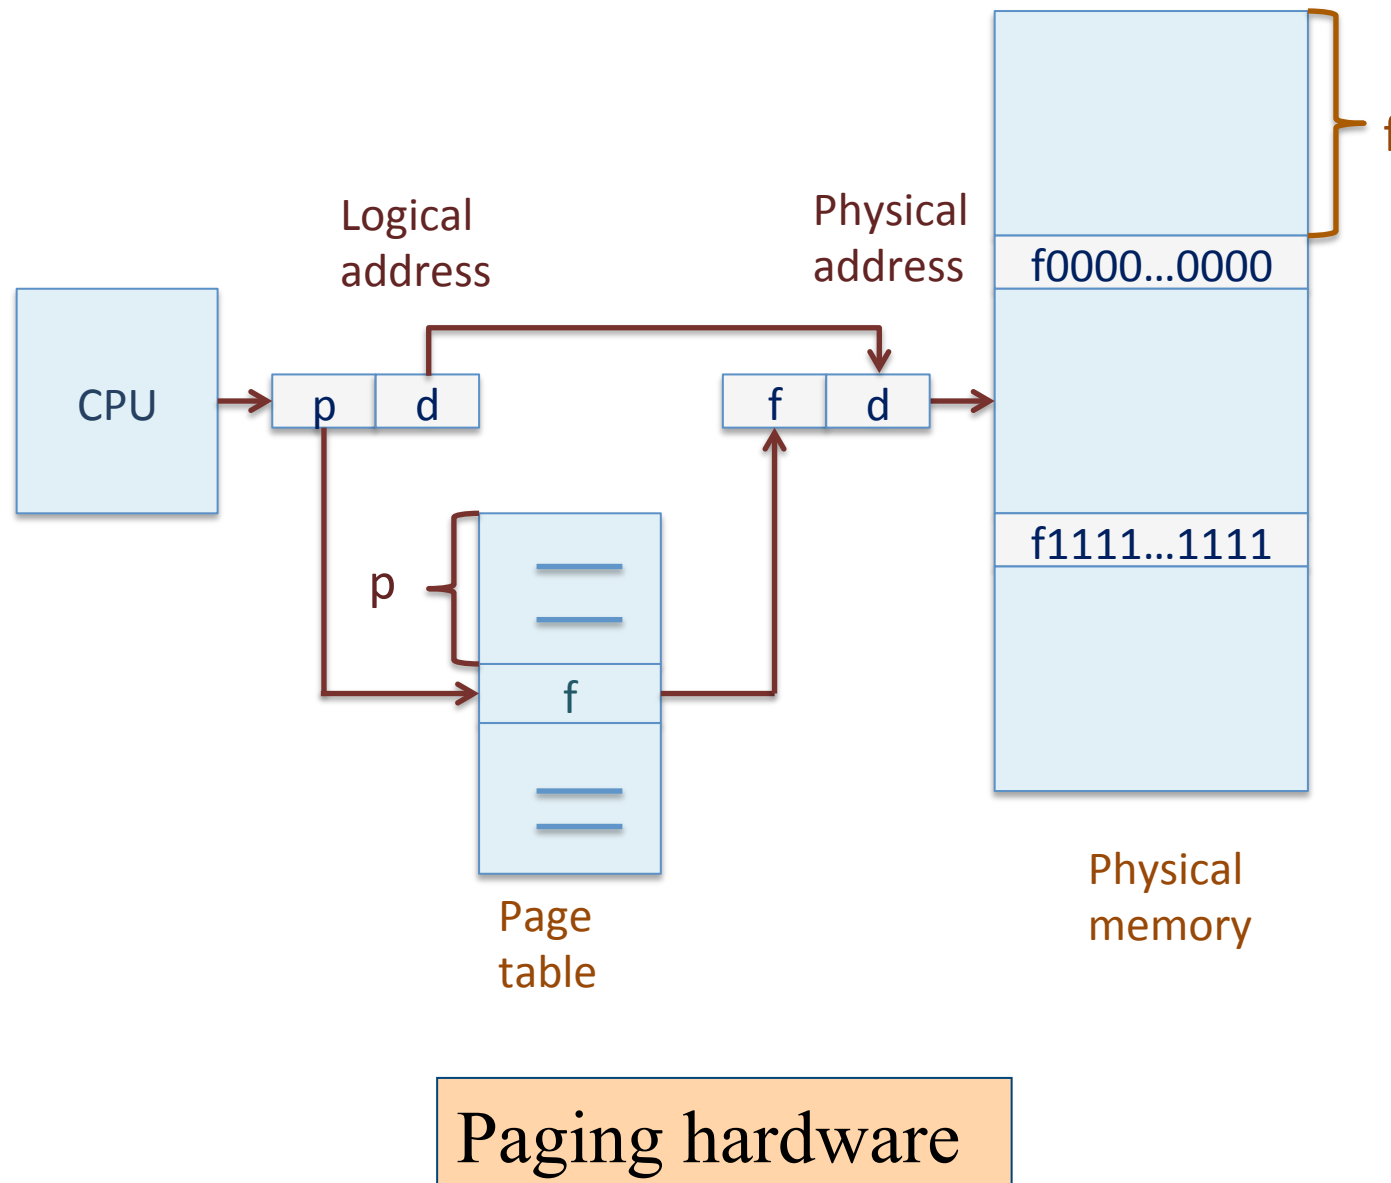
# 3. Paging

- Logical address space of a process can be noncontiguous; process is allocated physical memory.
- Divide physical memory into fixed-sized blocks called frames (size is power of 2, between 512 bytes and 16MB).
- Divide logical memory into blocks of the same size regions called pages.
- Keep track of all free frames.
  - Set up a page table to translate logical to physical addresses.

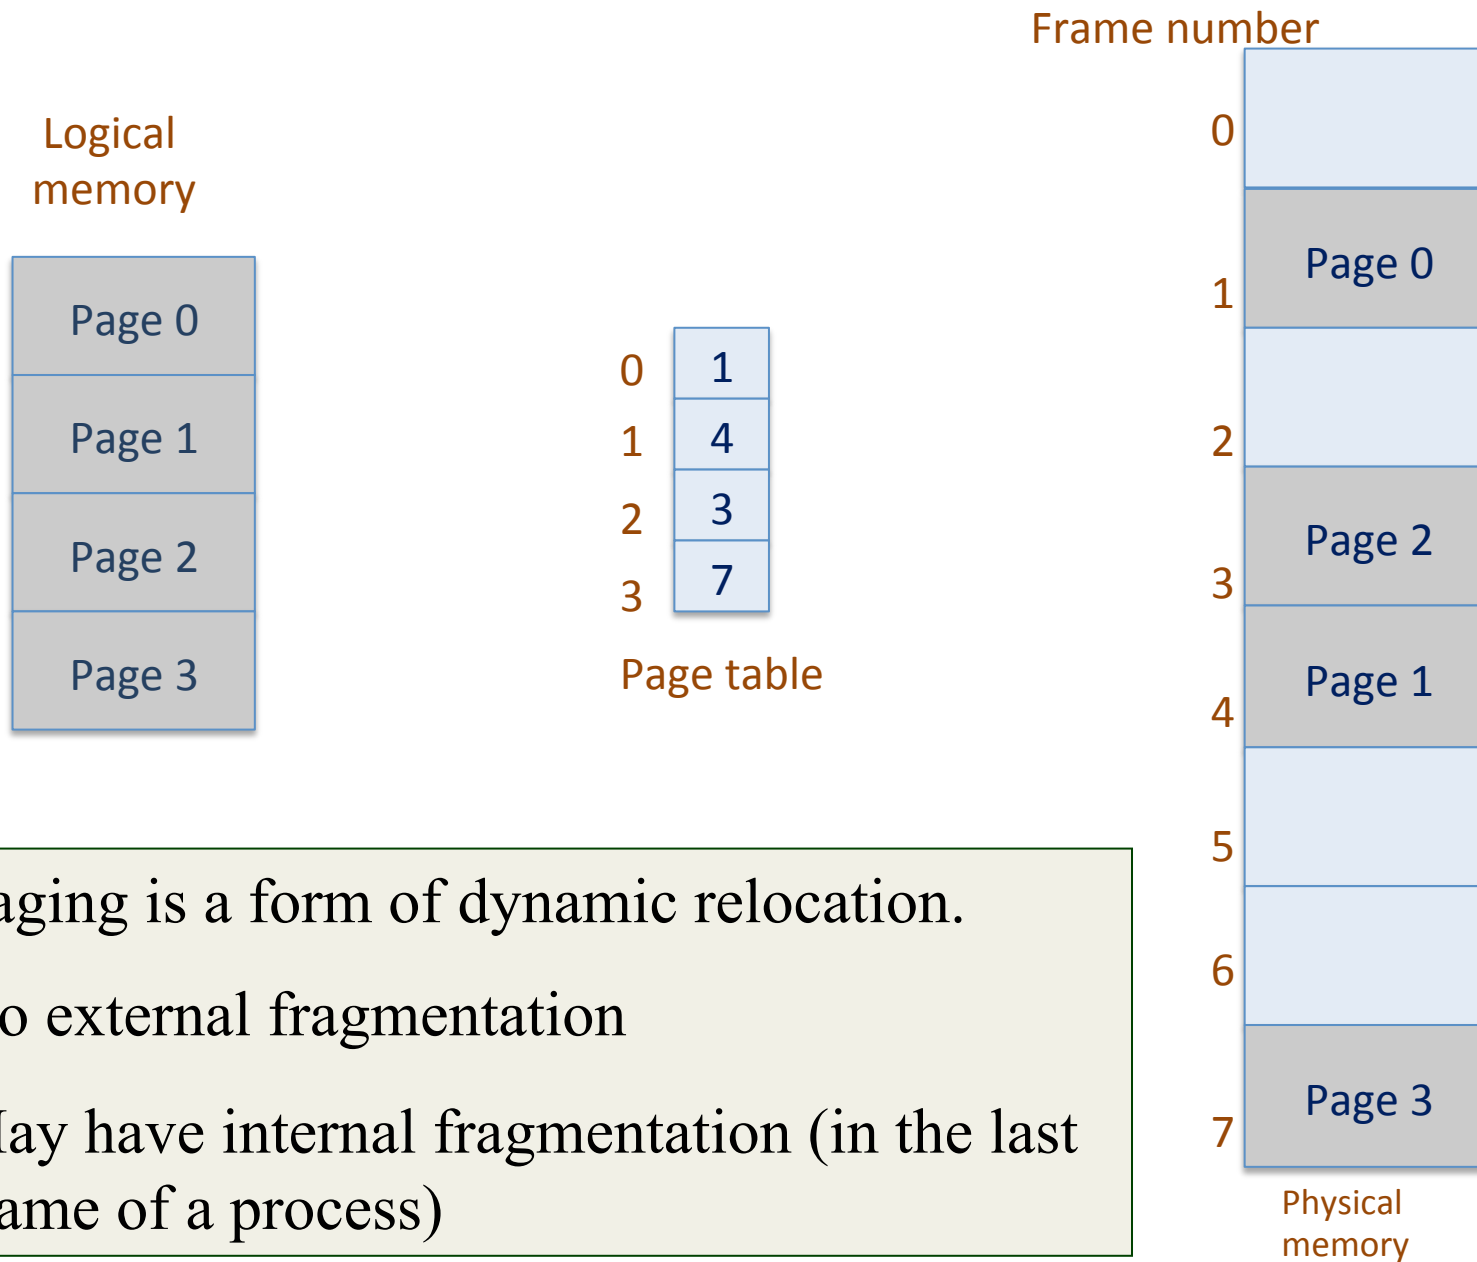
# Address Translation Scheme

- Address generated by CPU (logical address) is divided into:
  - **Page number (p)** – used as an index into a page table which contains base address of each page in physical memory.
  - **Page offset (d)** – combined with base address to define the physical memory address that is sent to the memory unit.

# Address Translation Architecture



# Paging Model of Logical and Physical Memory



# Paging Example

0	a
1	b
2	c
3	d
4	e
5	f
6	g
7	h
8	i
9	j
10	k
11	l
12	m
13	n
14	o
15	p

Logical memory

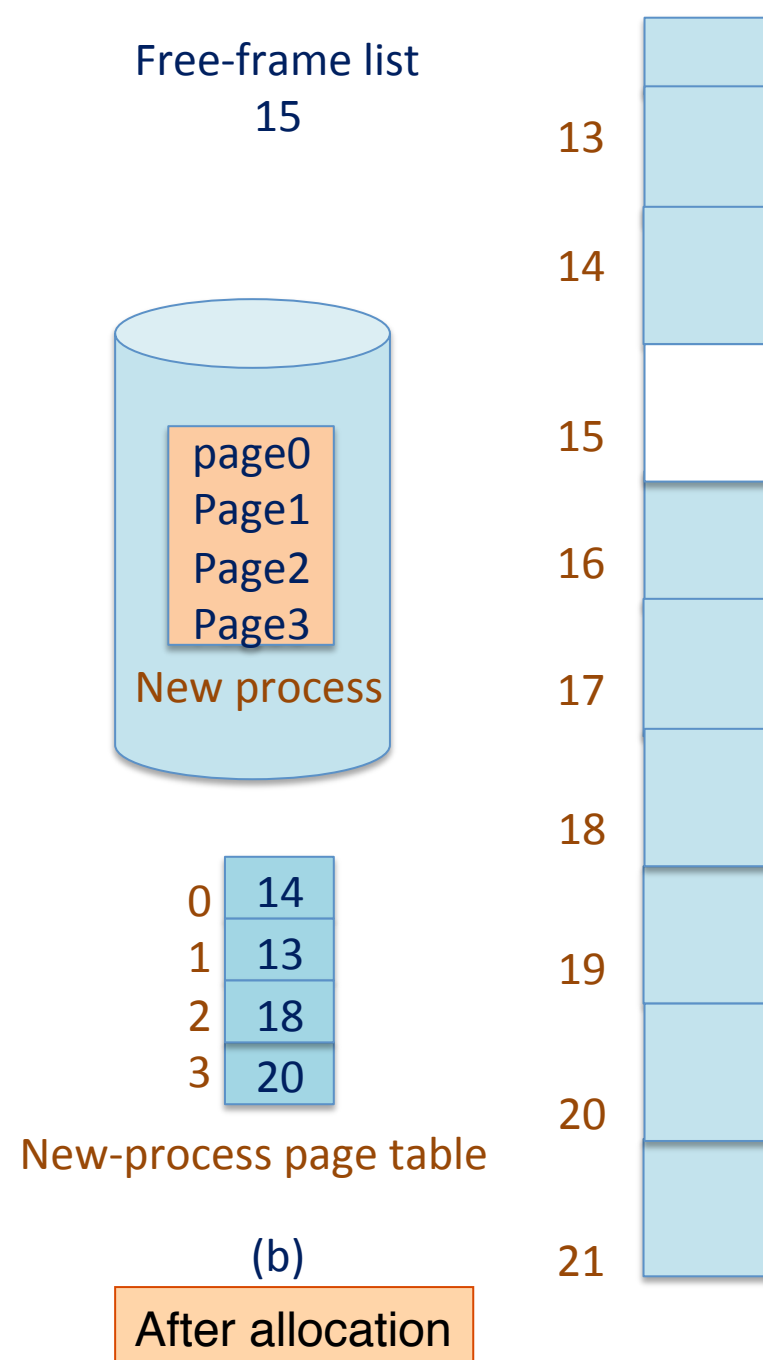
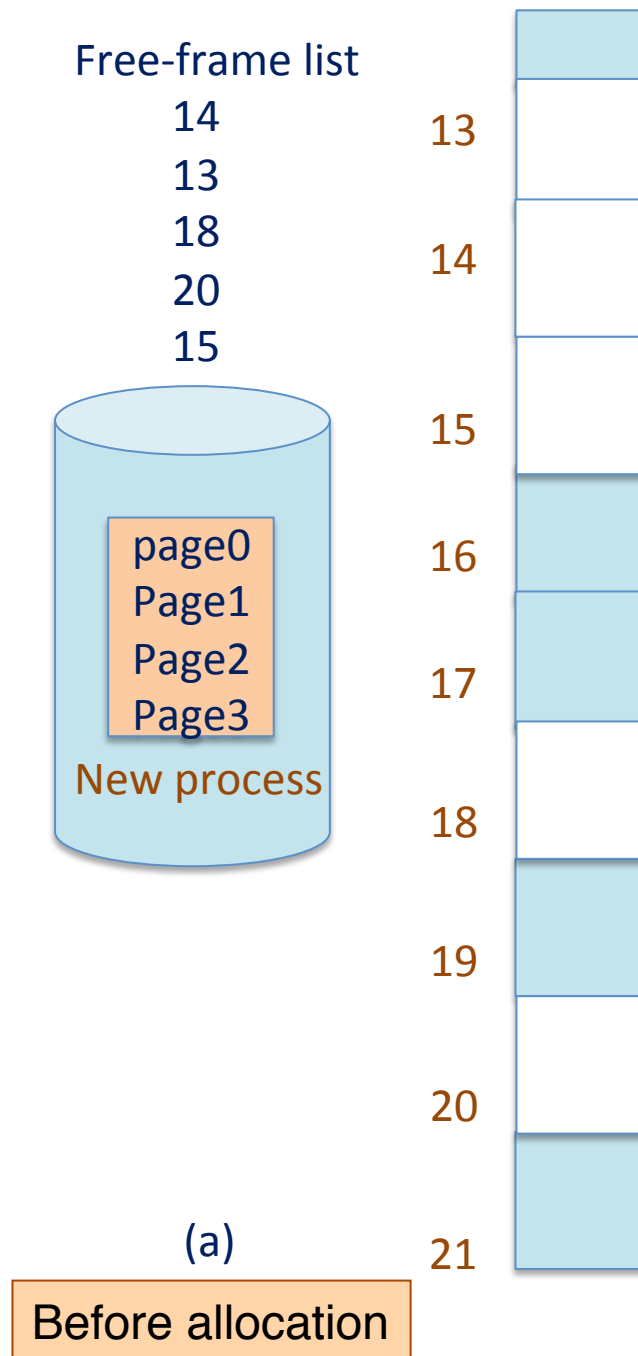
0	5
1	6
2	1
3	2

Page table

32-byte memory and 4-byte frames

0	
4	i j k l
8	m n o p
12	
16	
20	a b c d
24	e f g h
28	

Physical memory



# Implementation of Page Table

- Page table is kept in main memory.
- Page-table base register (PTBR) points to the page table.
- Page-table length register (PTLR) indicates size of the page table.
- **Problem:** Every data/instruction access requires two memory accesses:
  - one for the page table and
  - one for the data/instruction.
- **Solution:** The two memory access problem can be solved by the use of a special fast-lookup hardware cache called **associative memory** or **translation look-aside buffers (TLBs)**



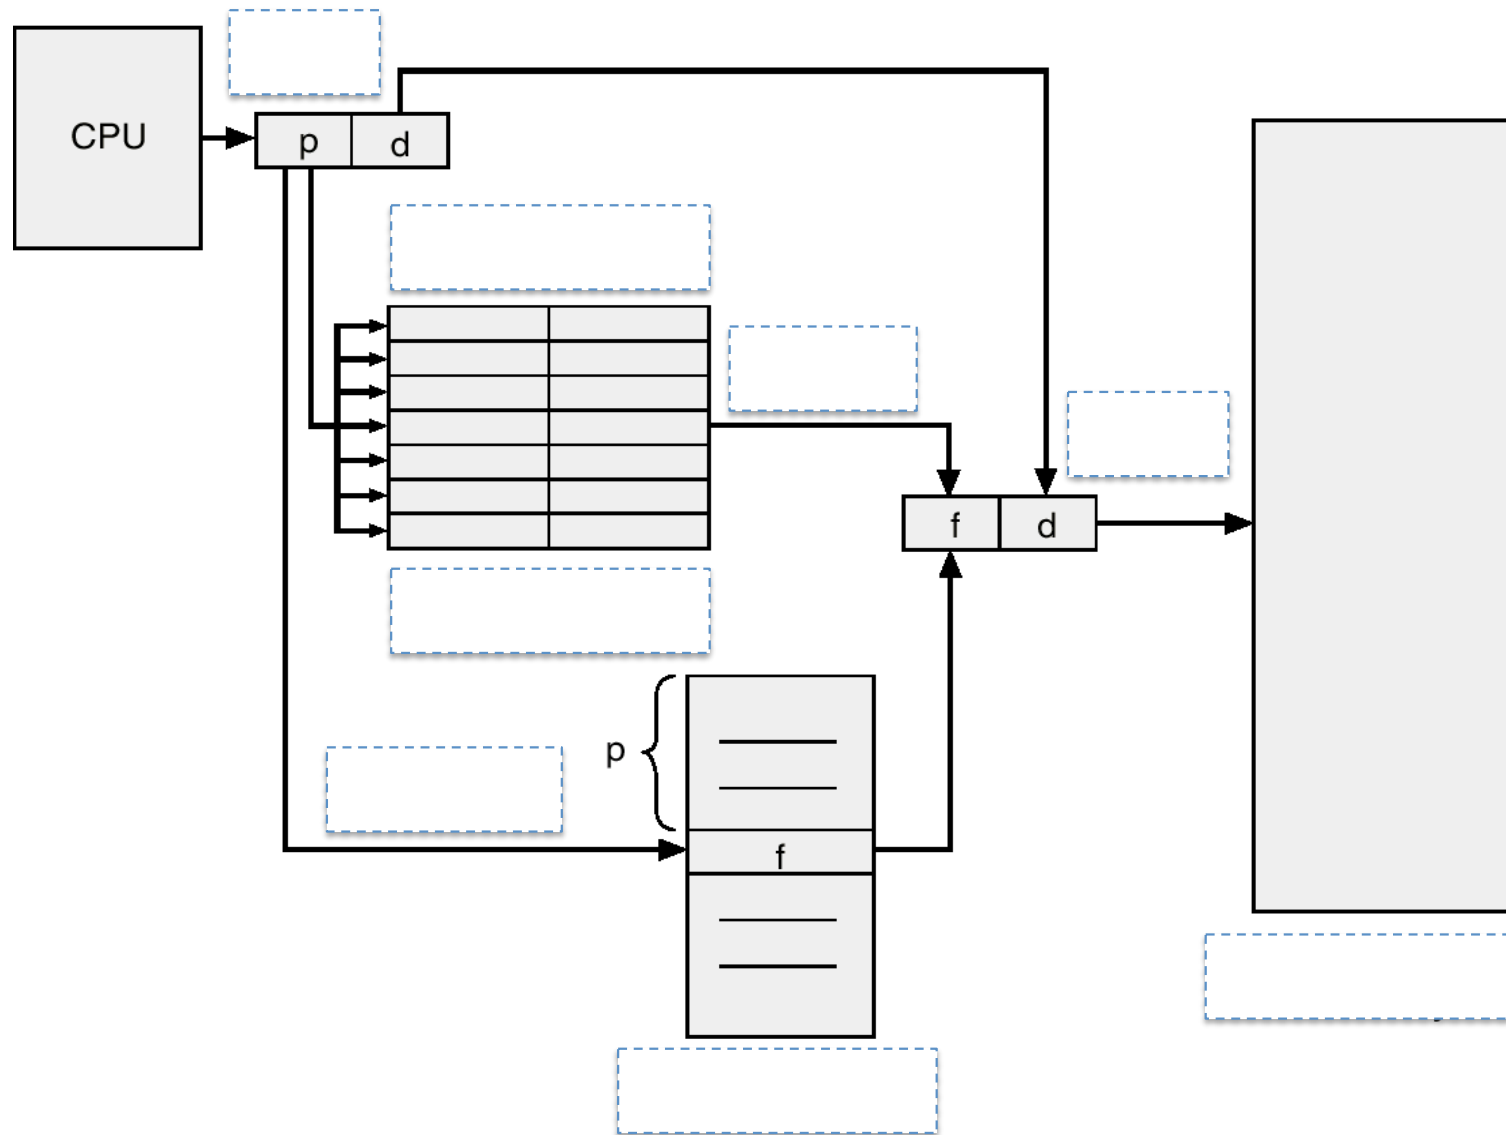
# Associative Memory

- Associative memory – parallel search

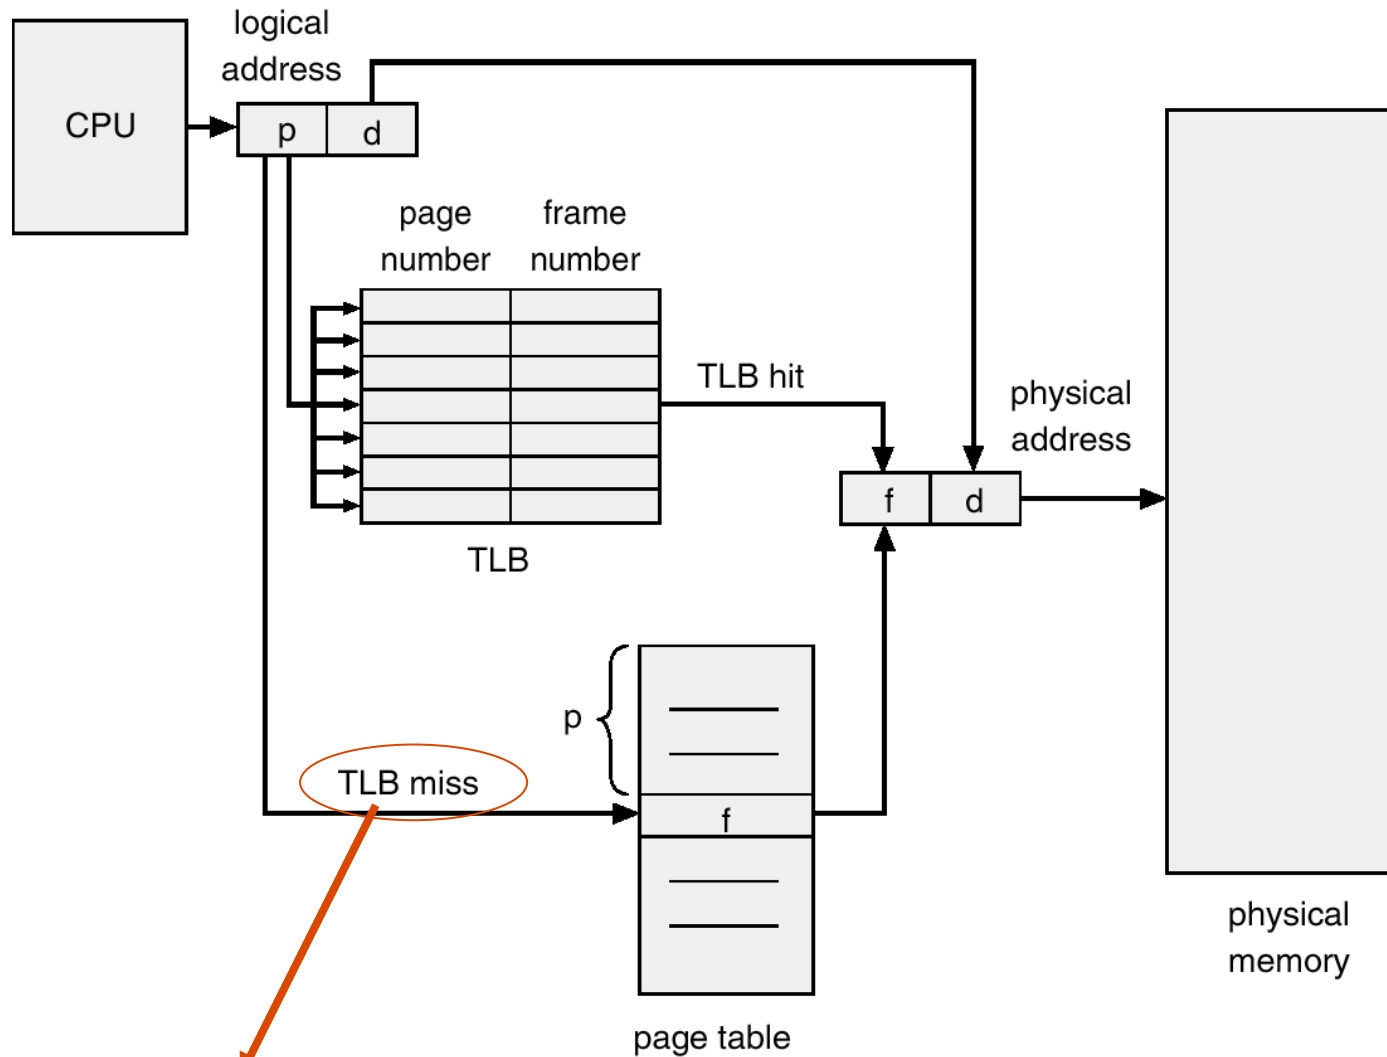
Page #	Frame #

- Address translation (p, d)
  - If p is in associative register, get frame # out.
  - Otherwise get frame # from page table in memory
- Search is fast
- TLB contains some of the page table entries (64 – 1024)

# Paging Hardware with TLB



# Paging Hardware with TLB



Page # and frame # is added to TLB

# Effective Access Time

- **Associative Lookup in TLB** =  $\varepsilon$  time unit
- **Hit ratio** – percentage of times that a page number is found in the associative registers (TLB)
  - ratio related to number of associative registers.
  - assume a hit ratio  $\alpha$
- **Effective Access Time (EAT)**




Example: **memory cycle time** is 1 time unit

$$\begin{aligned} \text{EAT} &= (1 + \varepsilon) \alpha + (2 + \varepsilon)(1 - \alpha) \\ &= 2 + \varepsilon - \alpha \end{aligned}$$

# Memory Protection

- Memory protection implemented by associating **protection bit** with each page/frame.
- **valid-invalid bit** attached to each entry in the page table:
  - **valid** indicates that the associated page is in the process' logical address space, and is thus a legal page.
  - **invalid** indicates that the page is not in the process' logical address space.

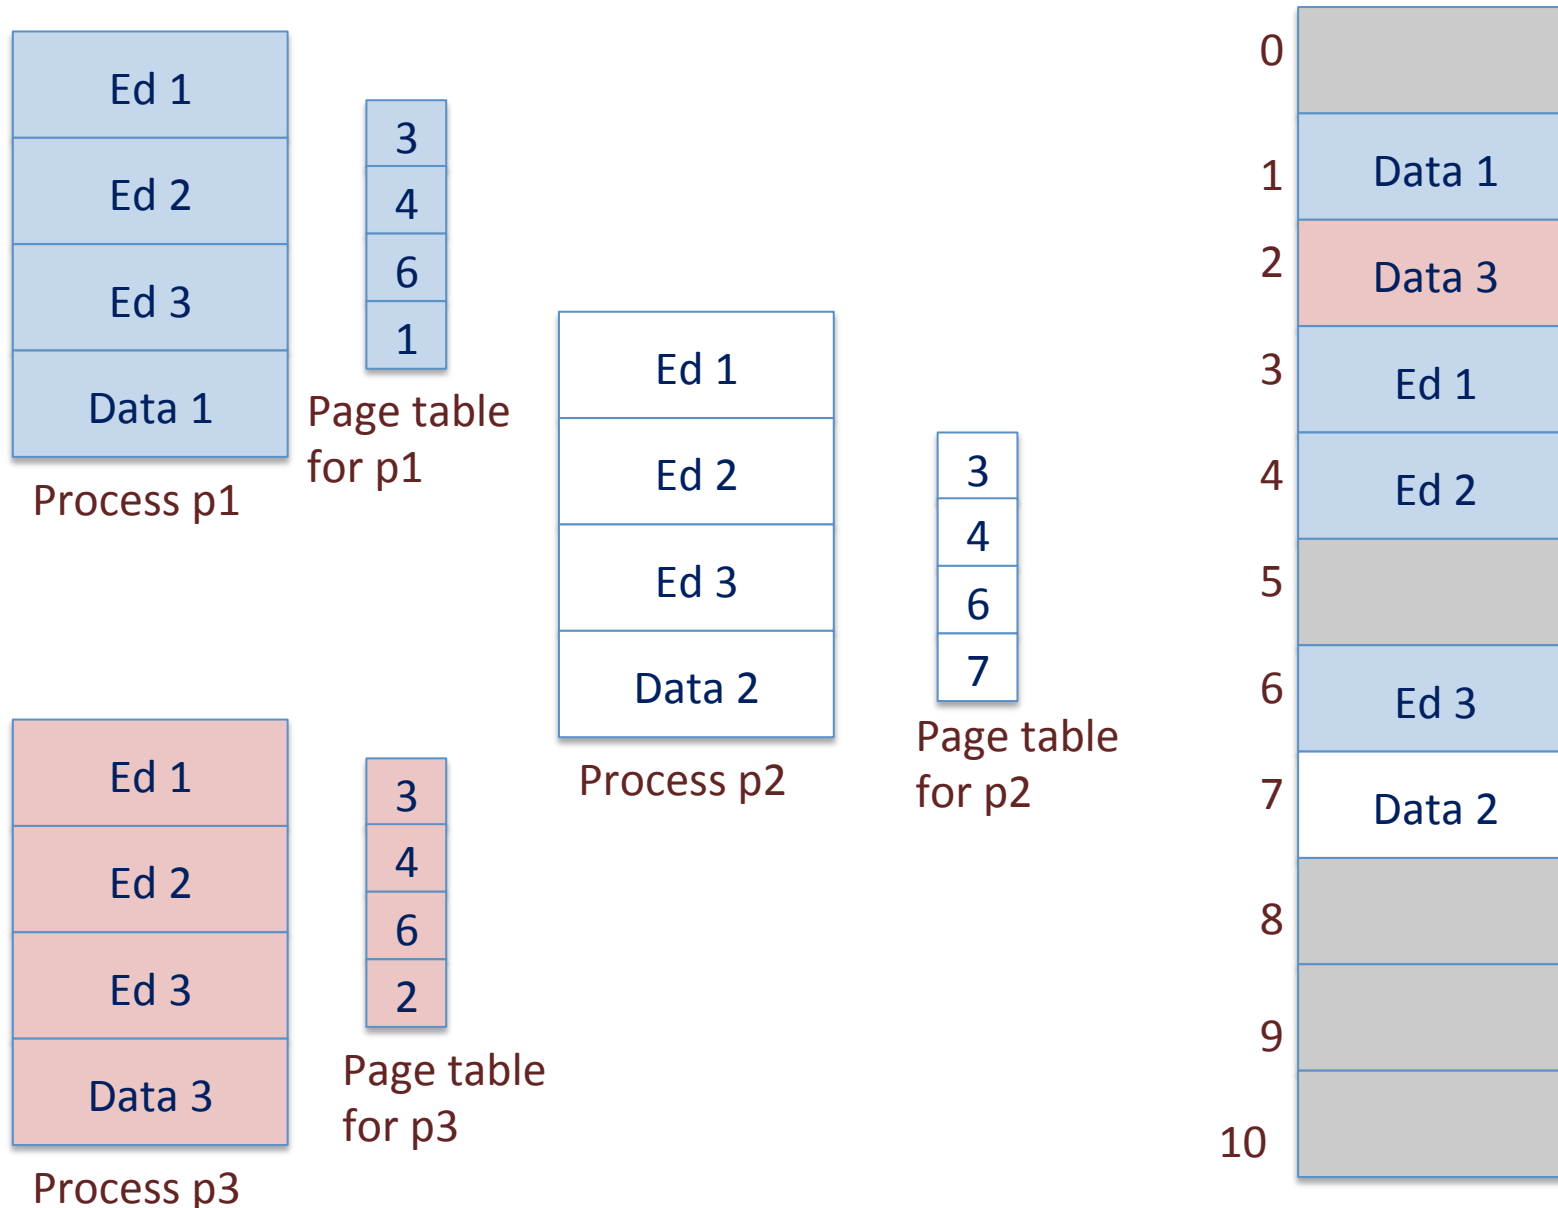
# valid (v) or invalid (i) bit in a page table



# Shared Pages

- An advantage of paging is the possibility of sharing common code
- Shared code
  - One copy of read-only (reentrant) code shared among processes (i.e., text editors, compilers, window systems).
  - Shared code must appear in same location in the logical address space of all processes.
- Private code and data
  - Each process keeps a separate copy of the code and data.
  - The pages for the private code and data can appear anywhere in the logical address space.

# Shared Pages Example





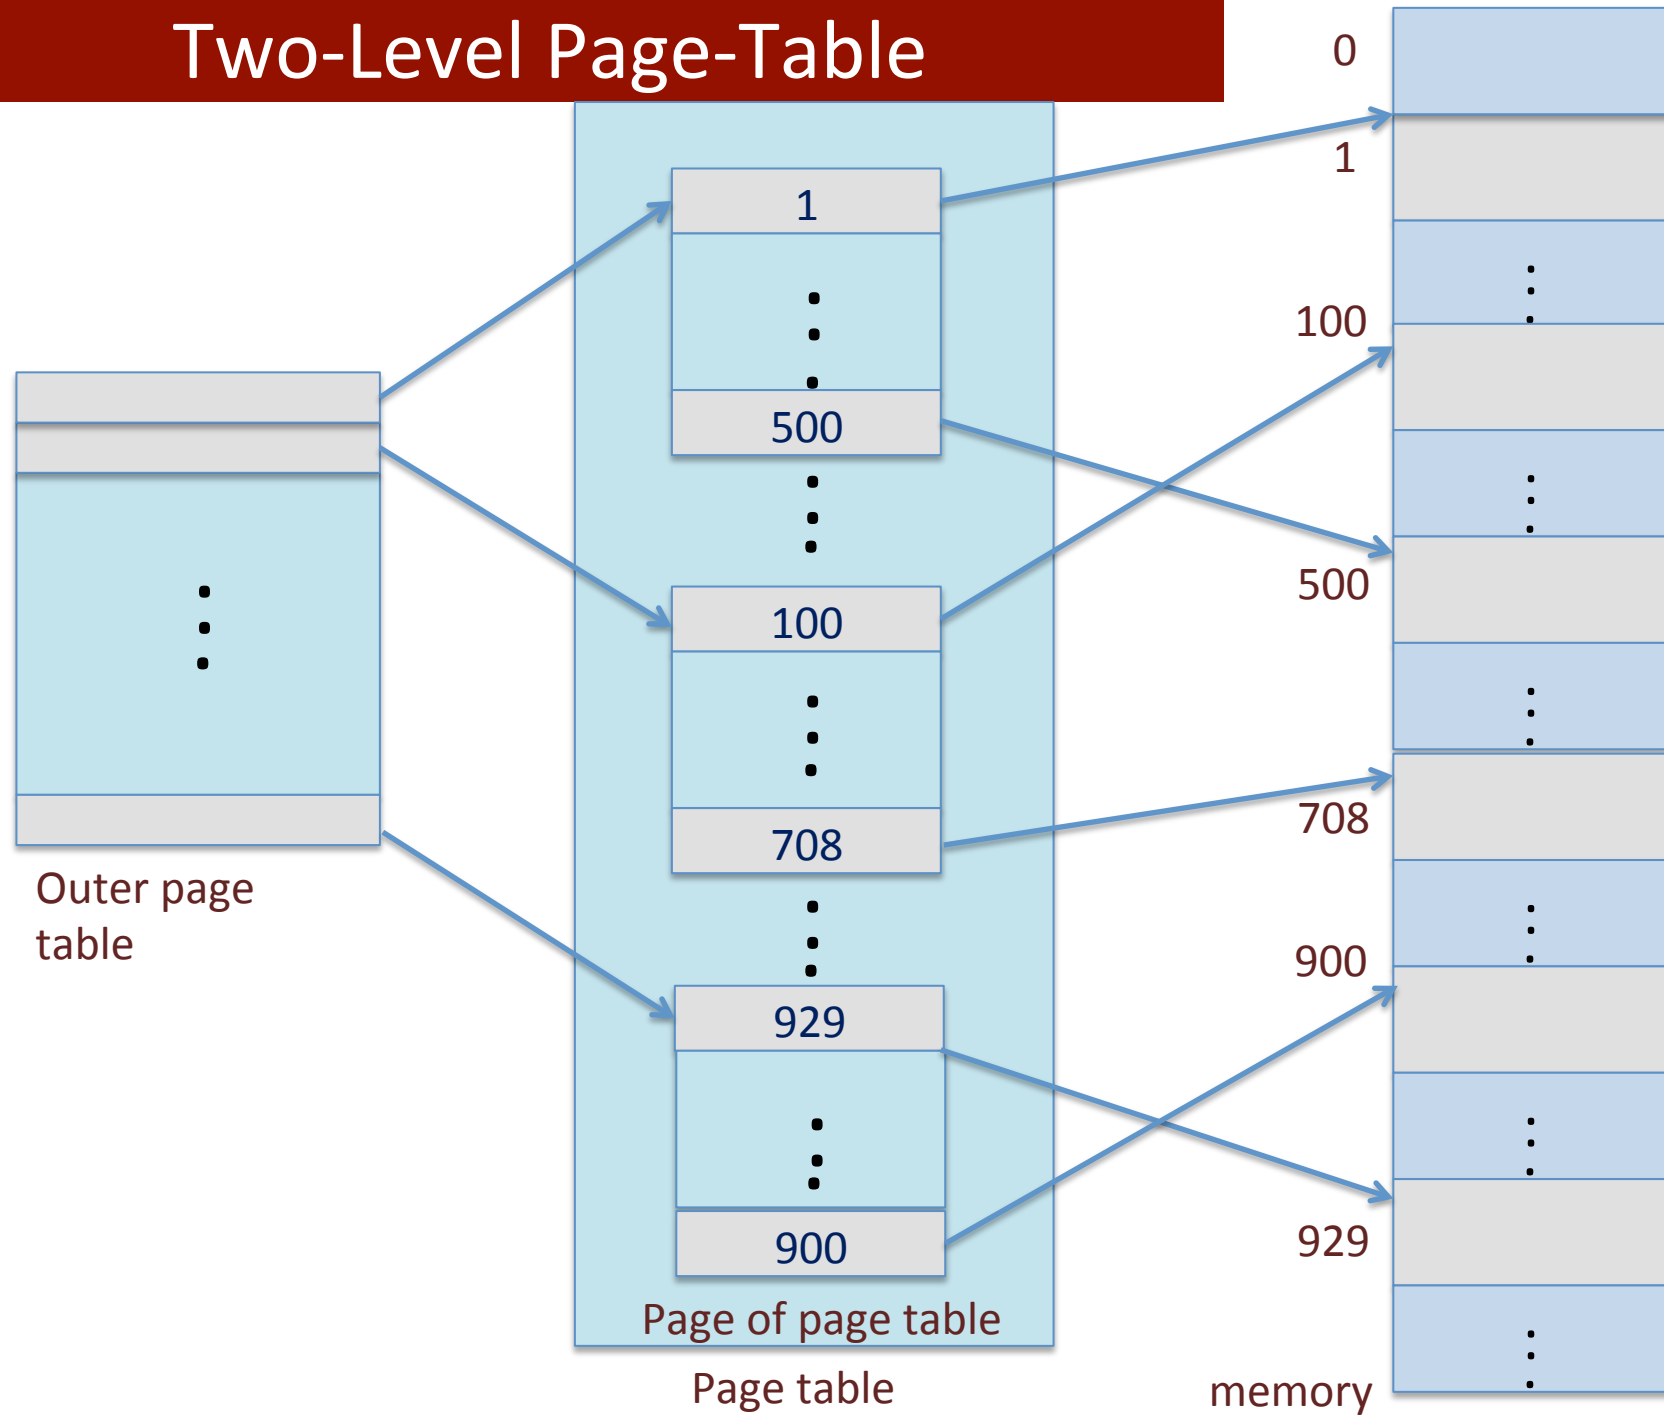
# Page Table Structure

- Modern computer systems support a large logical-address space
- Do not want to allocate the page table contiguously in main memory
- Hierarchical Page Tables
- Hashed Page Tables
- Inverted Page Tables

# Hierarchical Page Tables

- Break up the logical address space into multiple page tables.
- A simple technique is a **two-level page table**.
- Page table itself is also paged

# Two-Level Page-Table



# Two-Level Paging Example

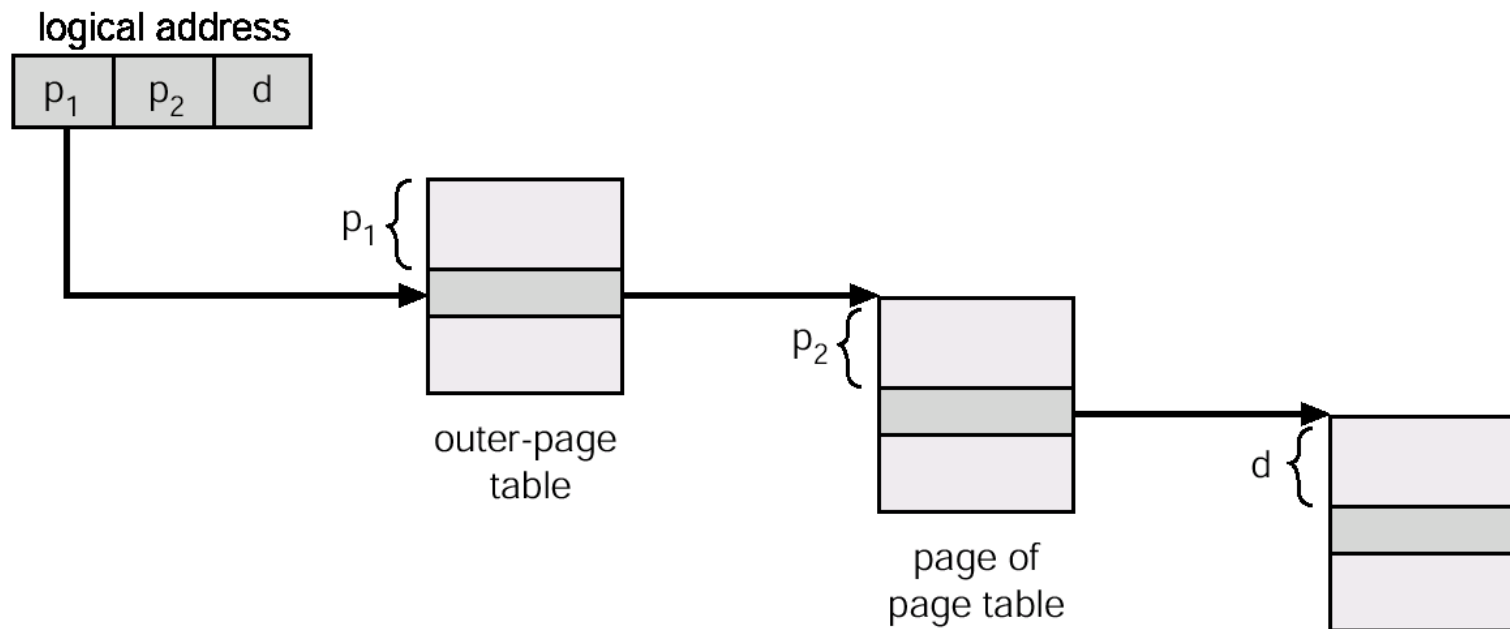
- A logical address (on 32-bit machine with 4K page size) is divided into:
  - a **page number** consisting of 20 bits.
  - a **page offset** consisting of 12 bits.
- Since the page table is paged, the **page number** is further divided into:
  - a 10-bit page number.
  - a 10-bit page offset.
- Thus, a logical address is as follows:

page number		page offset
$p_1$	$p_2$	$d$
10	10	12

where  $p_1$  is an index into the outer page table, and  $p_2$  is the displacement within the page of the outer page table.

# Address-Translation Scheme

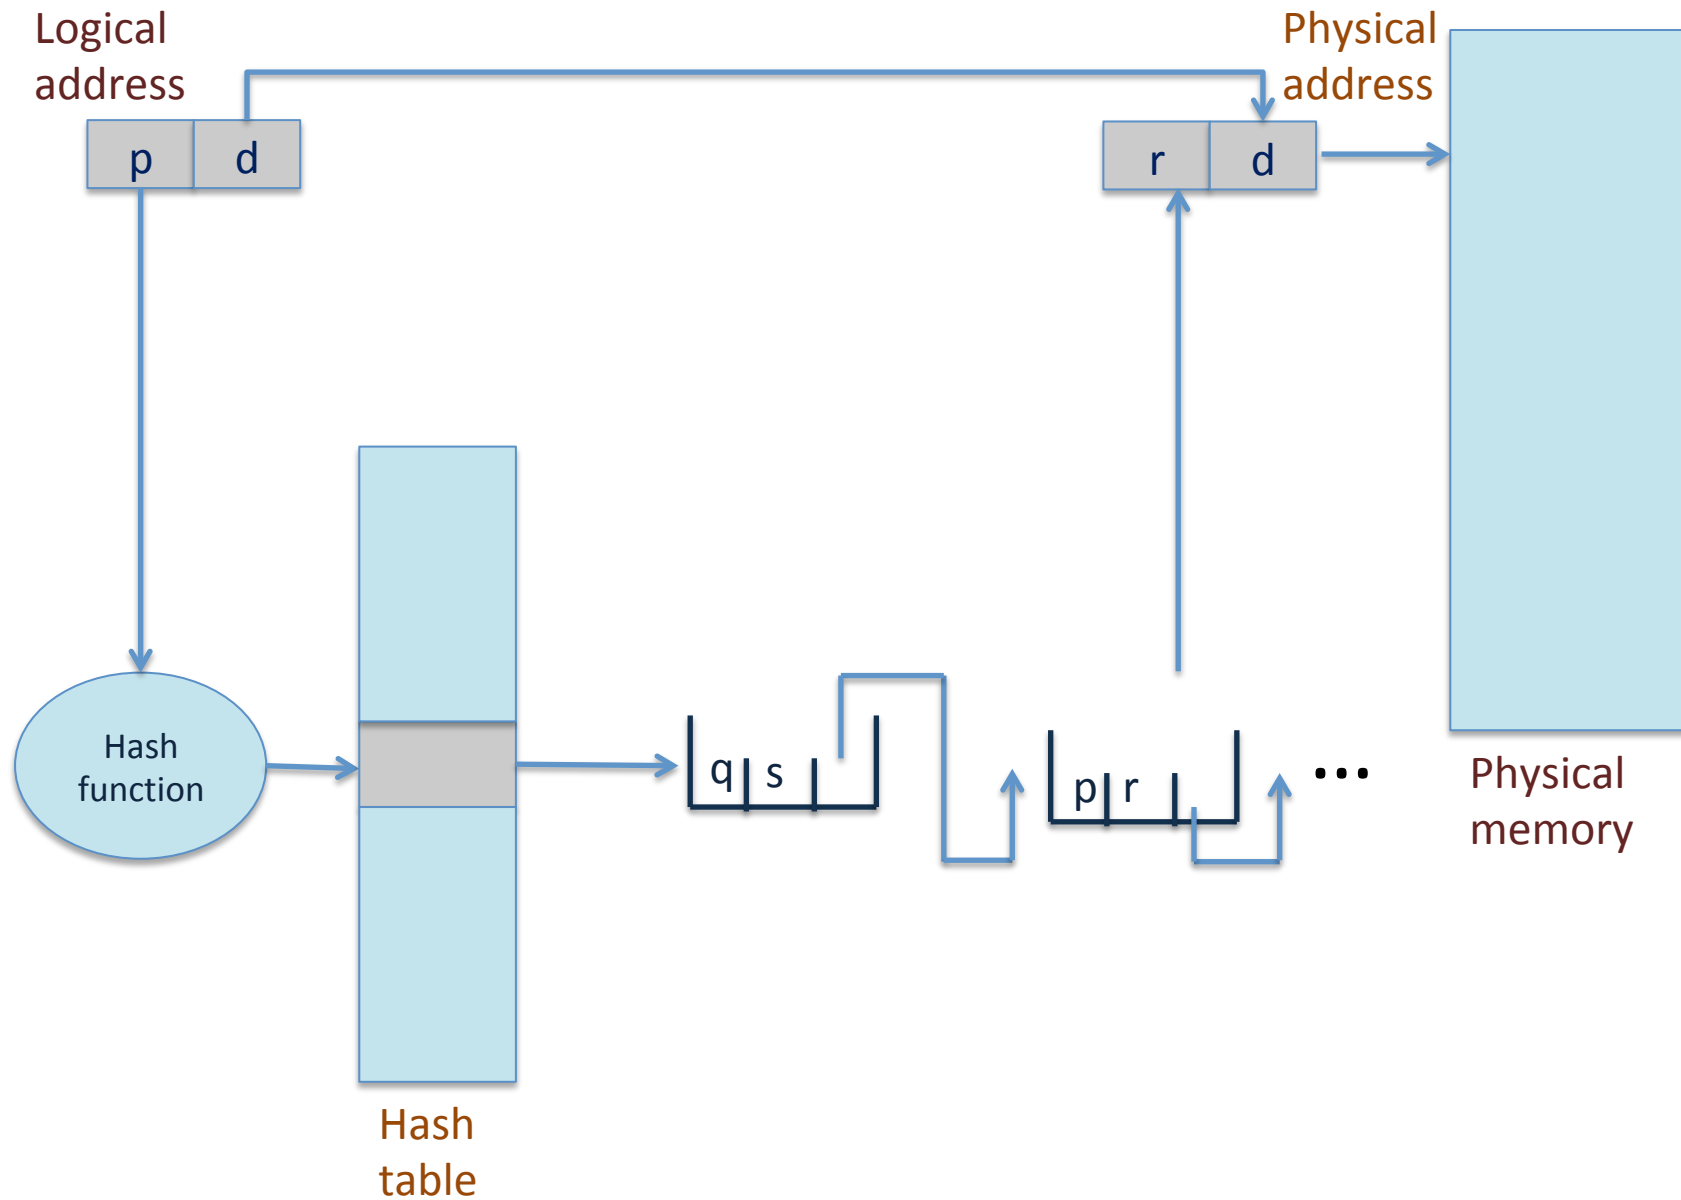
Address-translation for a two-level 32-bit paging architecture



# Hashed Page Tables

- Common in address spaces  $> 32$  bits.
- The virtual (logical) page number is hashed into a page table. This hashed page table contains a chain of elements hashing to the same location.
- Virtual page numbers are compared in this chain searching for a match. If a match is found, the corresponding physical frame is extracted.

# Hashed Page Table

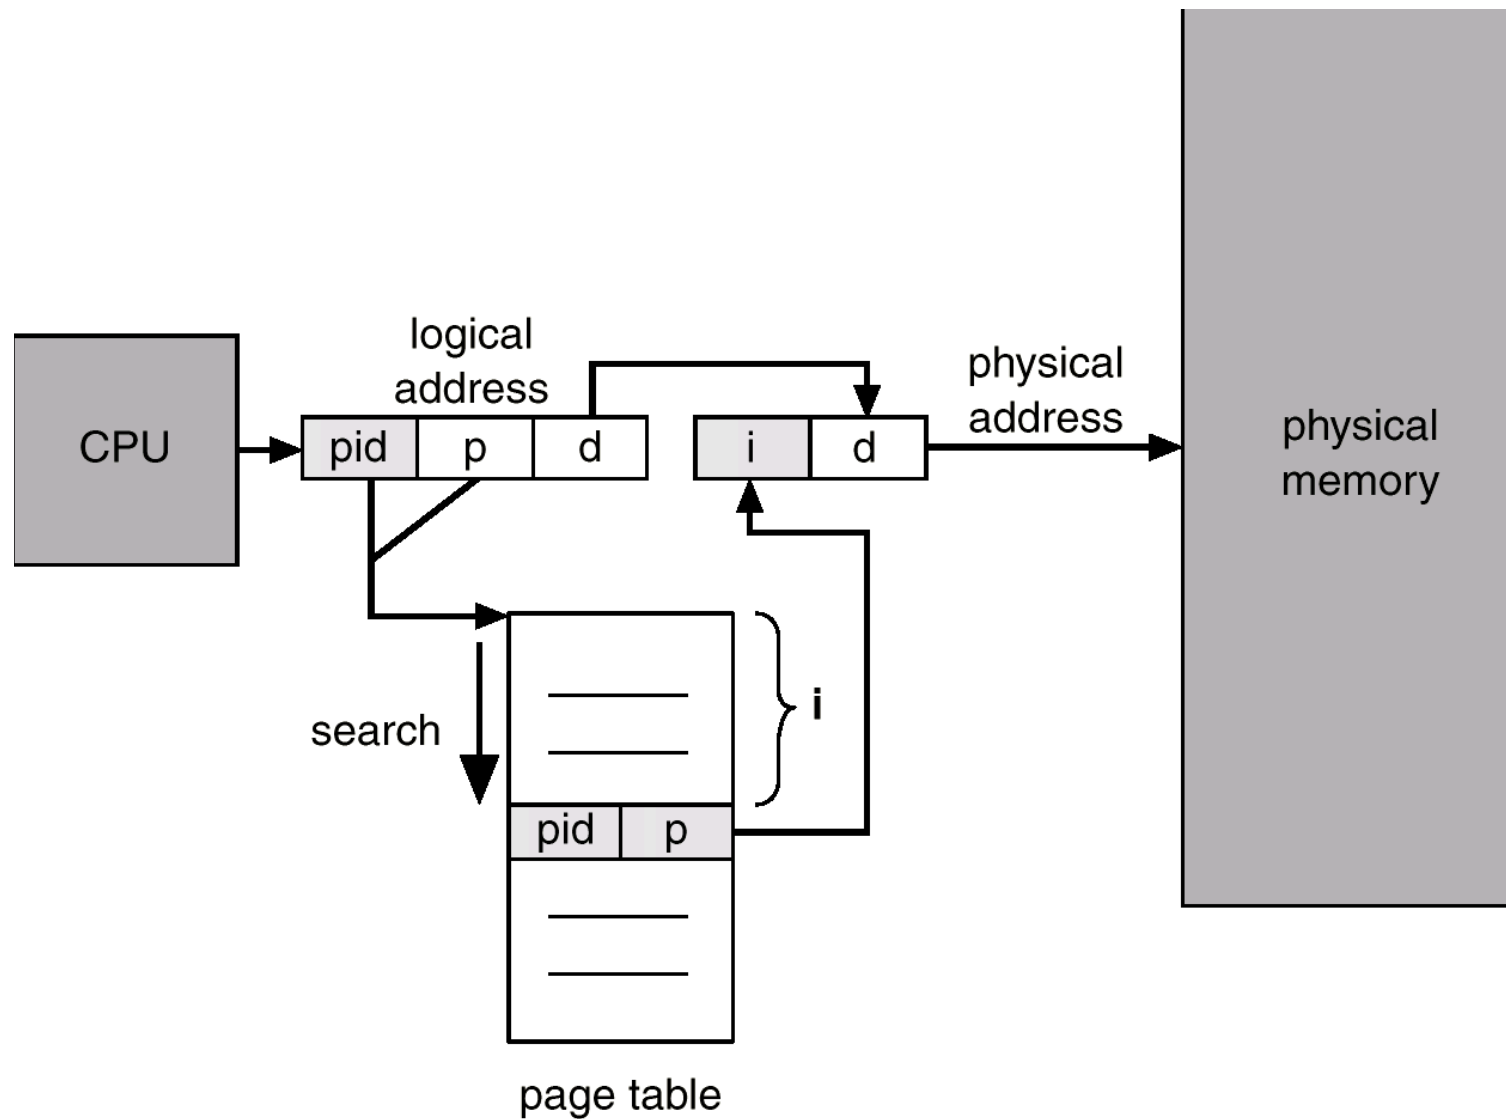


# Inverted Page Table



- Only one page table in the system, with one entry for each frame of physical memory.
- Entry consists of the virtual address of the page, with information about the process that owns that page.
- Decreases memory needed to store each page table, but increases time needed to search the table when a page reference occurs.
- Use hash table to limit the search to one — or at most a few — page-table entries.



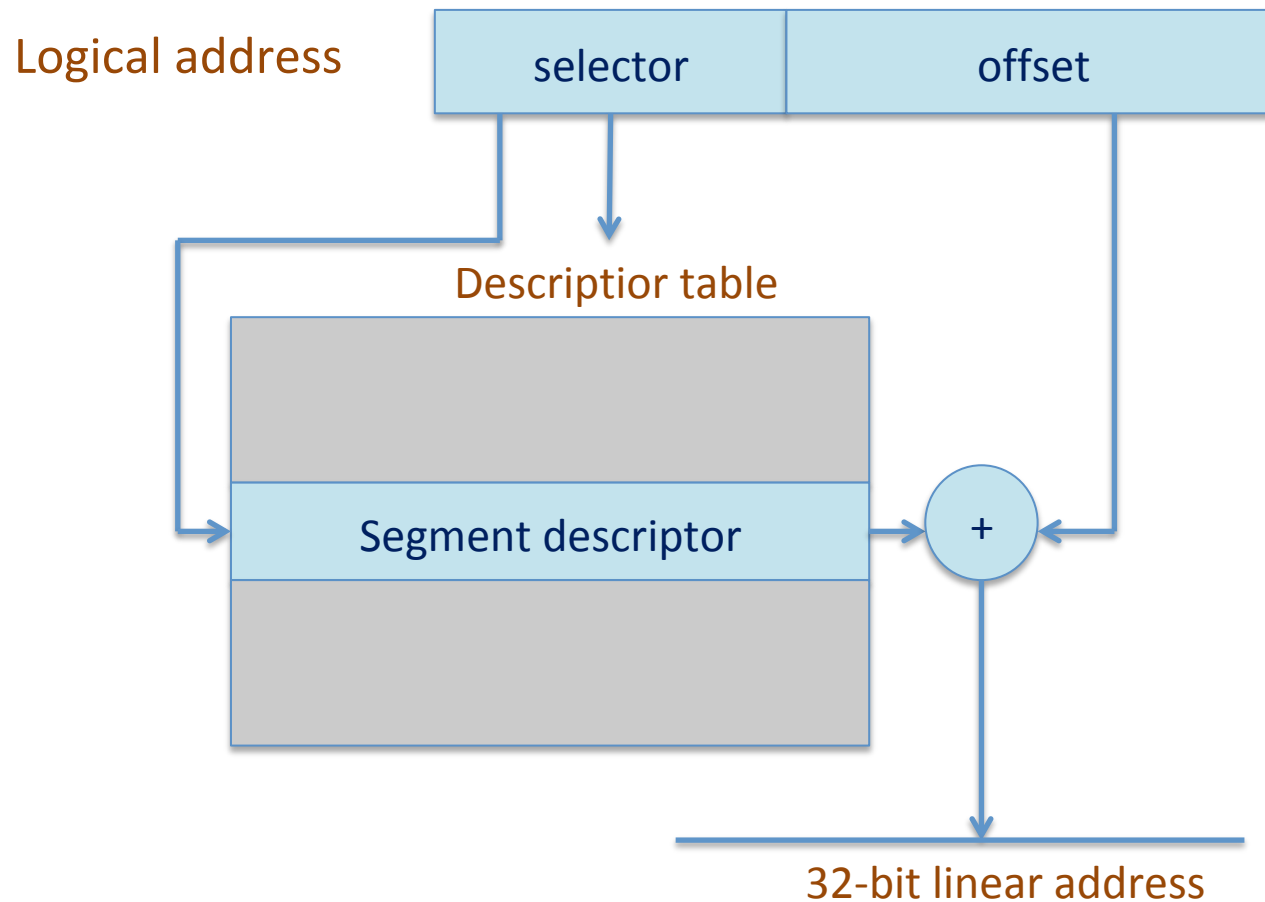
# Inverted Page Table Architecture



# Logical to Physical Address Translation in Pentium

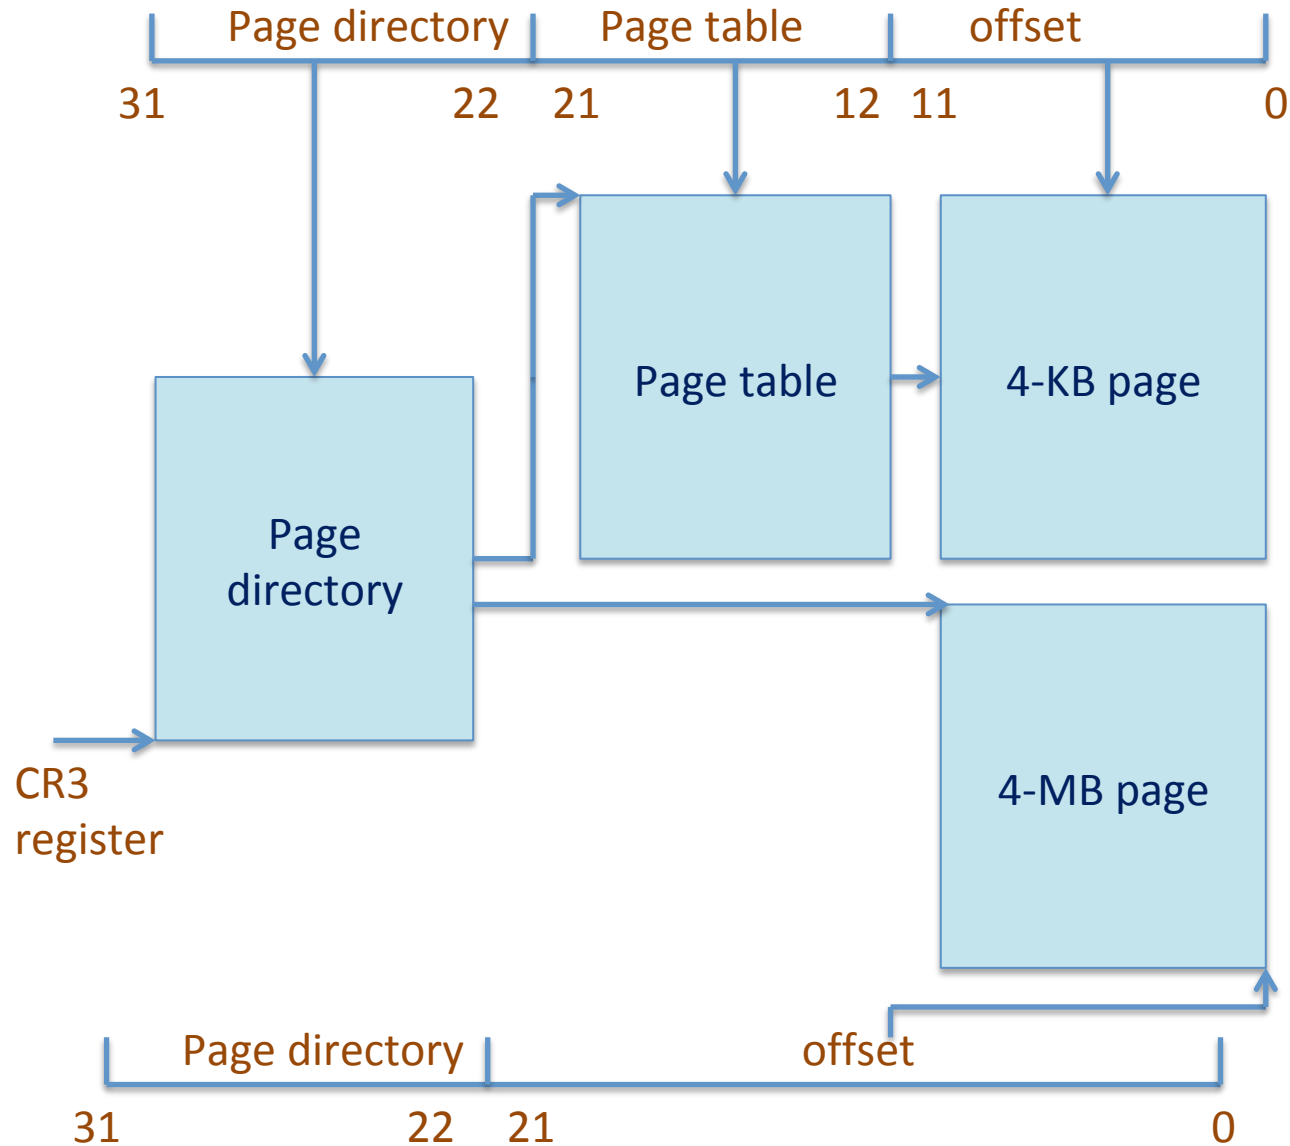


# Intel Pentium Segmentation



# Pentium Paging Architecture

(logical address)

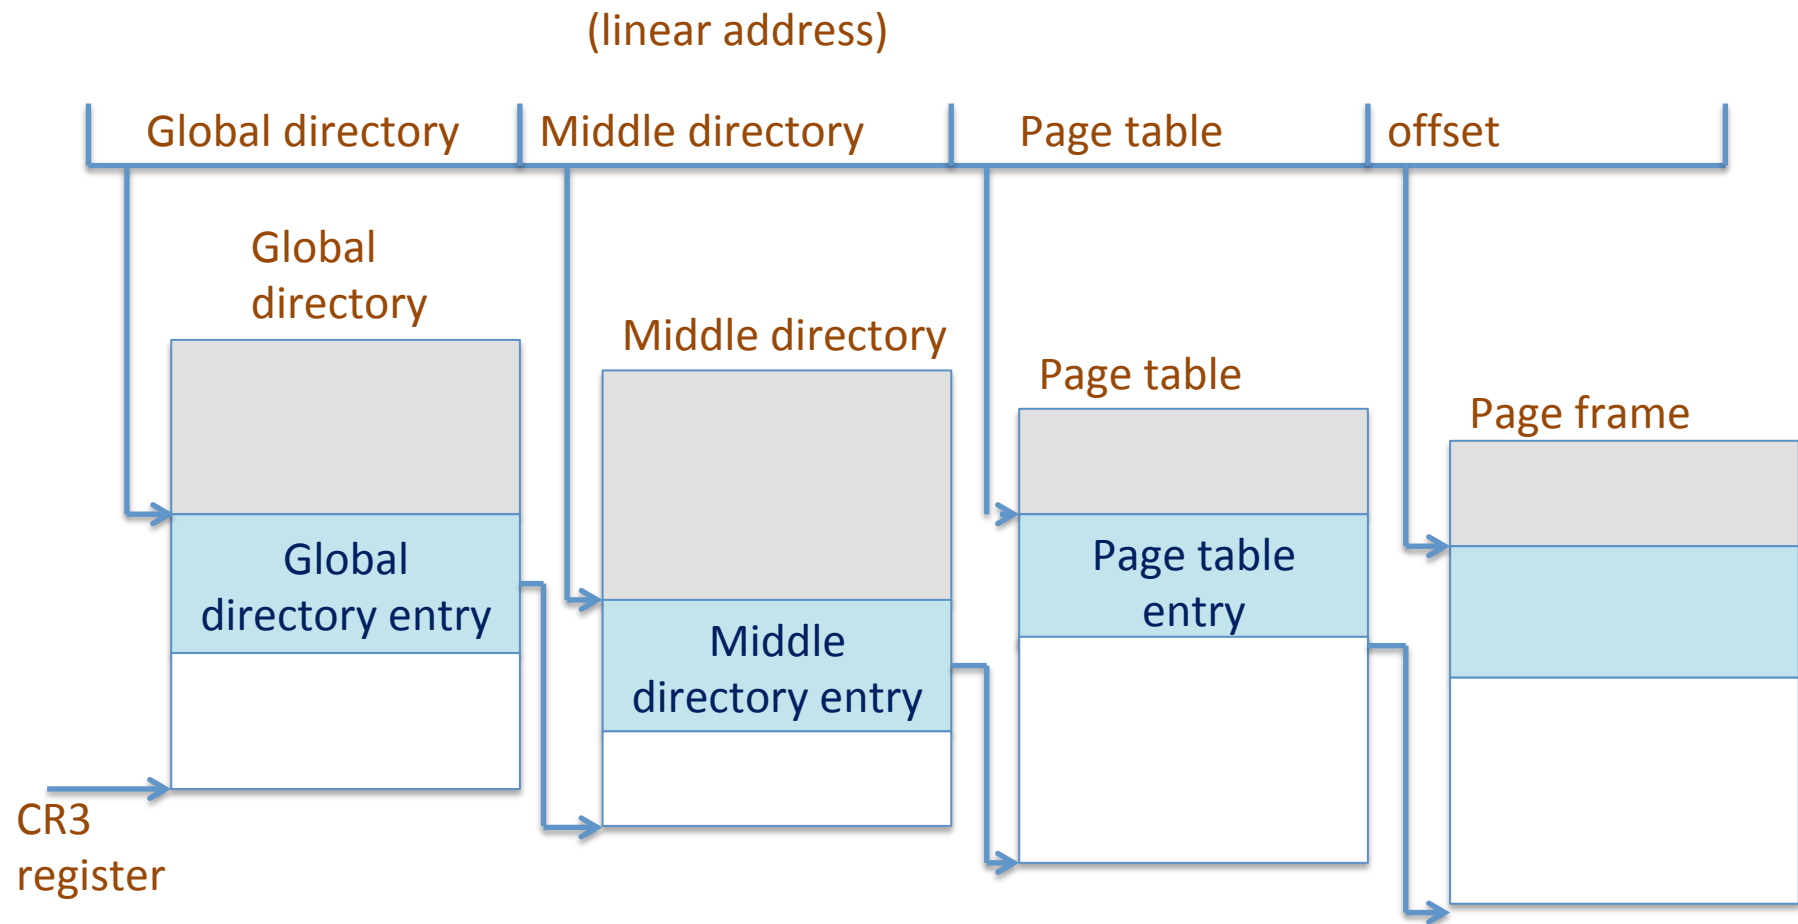


# Linear Address in Linux

Broken into four parts:

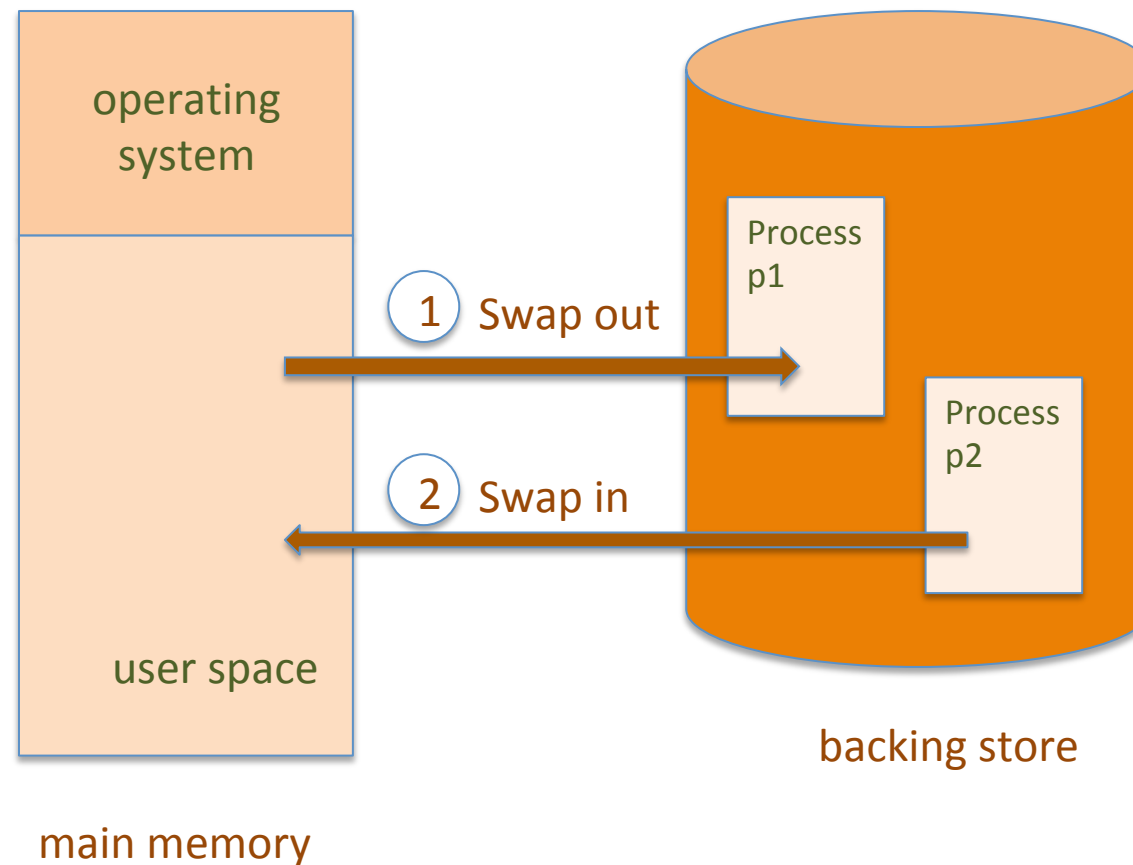
global directory	middle directory	page table	offset
---------------------	---------------------	---------------	--------

# Three-level Paging in Linux



# Swapping

- A process can be **swapped** temporarily out of memory to a backing store, and then brought back into memory for continued execution.



# Acknowledgments

- These slides are adapted from
  - Öznur Özkasap (Koç University)
  - Operating System and Concepts (9<sup>th</sup> edition) Wiley