

CPU Scheduling Algorithms

Didem Unat

Lecture 6

COMP304 - Operating Systems (OS)

Terminology

- **CPU utilization** – keep the CPU as busy as possible
- **Throughput** – # of processes that complete their execution per time unit
- **Turnaround time** – amount of time to execute a particular process
- **Waiting time** – amount of time a process has been waiting in the ready queue
- **Response time** – amount of time it takes from when a request was submitted until the first response is produced, **not** output (for time-sharing environment)

1. First-Come, First Served (FCFS)

<u>Process</u>	<u>Burst Time</u>
P ₁	24
P ₂	3
P ₃	3

- Suppose that the processes arrive in the order: P₁ , P₂ , P₃
The Gantt Chart for the schedule is:



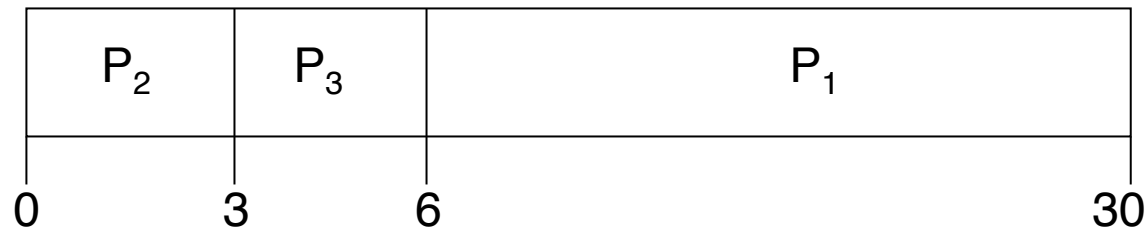
- Waiting times for: P₁ = 0; P₂ = 24; P₃ = 27
- Average waiting time: $(0 + 24 + 27)/3 = 17$

FCFS Scheduling (Cont.)

Suppose that the processes arrive in the order

P_2, P_3, P_1

- The Gantt chart for the schedule is:



- Waiting time for $P_1 = 6$; $P_2 = 0$; $P_3 = 3$
- Average waiting time: $(6 + 0 + 3)/3 = 3$
- Much better than previous case.
- Convoy effect:** short process behind long process


2. Shortest Job First (SJF) Scheduling

- Associate with each process the length of its next CPU burst. Use these lengths to schedule the process with the shortest time.
- Two schemes:
 - **Non-preemptive** – once CPU given to the process it cannot be preempted until completes its CPU burst.
 - **preemptive** – if a new process arrives with CPU burst length less than remaining time of current executing process, preempt. This is known as the **Shortest-Remaining-Time-First (SRTF)**.
- SJF is **optimal** – gives minimum average waiting time for a given set of processes.

Non-preemptive (SJF) Scheduling

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P ₁	0.0	7
P ₂	2.0	4
P ₃	4.0	1
P ₄	5.0	4

- SJF (non-preemptive)




- Average waiting time ?
= $(0 + 6 + 3 + 7)/4 = 4$

Example of Preemptive SJF

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P ₁	0.0	7
P ₂	2.0	4
P ₃	4.0	1
P ₄	5.0	4

- SJF (preemptive)



- Average waiting time ?
 $= (9 + 1 + 0 + 2)/4 = 3$


3. Priority Scheduling

- A **priority** number (integer) is associated with each process
- The CPU is allocated to the process with the highest priority (smallest integer \equiv highest priority).
- Can be preemptive and nonpreemptive
- FCFS and SJF are special cases of priority scheduling
 - Why?
 - In FCFS priority is based on the arrival time
 - SJF is a priority scheduling where priority is the predicted next CPU burst time.

Priority Scheduling

Pi	$\tau(p_i)$	Priority
0	350	5
1	125	2
2	475	3
3	250	1
4	75	4

•Nonpreemptive example



$$T_{\text{TRnd}}(p_0) = \tau(p_0) + \tau(p_4) + \tau(p_2) + \tau(p_1) + \tau(p_3) = 350 + 75 + 475 + 125 + 250 = 1275$$

$$R(p_0) = 925$$

$$R(p_1) = 250$$

$$R(p_2) = 375$$

$$T_{\text{TRnd}}(p_1) = \tau(p_1) + \tau(p_3) = 125 + 250 = 375$$

$$T_{\text{TRnd}}(p_2) = \tau(p_2) + \tau(p_1) + \tau(p_3) = 475 + 125 + 250 = 850$$

$$T_{\text{TRnd}}(p_3) = \tau(p_3) = 250$$

$$R(p_3) = 0$$


$$T_{\text{TRnd}}(p_4) = \tau(p_4) + \tau(p_2) + \tau(p_1) + \tau(p_3) = 75 + 475 + 125 + 250 = 925$$

$$R(p_4) = 850$$

$$\text{Average response time } R_{\text{avg}} = (925 + 250 + 375 + 0 + 850) / 5 = 2400 / 5 = 480$$

Problem with Priority Scheduling

- Problem \equiv **Starvation** – low priority processes may never execute.
- Solution \equiv **Aging** – as time progresses increase the priority of the process.



IBM 7094 operator's console

- Rumor has it that,
 - when they shut down the IBM 7094 at MIT in 1973, they found a low-priority process that had been submitted in 1967 and had not yet been run.

Priority

- What if a high-priority process needs to access the data that is currently being held by a low- priority process?
- The high-priority process is blocked by the low-priority process. This is **priority inversion**.
- This can be solved with **priority-inheritance protocol**.
 - The low priority process accessing the data inherits the high priority until it is done with the resource.
 - When the low-priority process finishes, its priority reverts back to the original.


4. Round Robin (RR)

- Each process gets a small unit of CPU time (**time quantum**), usually 10-100 milliseconds.
- After this time has elapsed, the process is **preempted** and added to the end of the ready queue.
- If there are n processes in the ready queue and the time quantum is q , then each process gets $1/n$ of the CPU time in chunks of at most q time units at once. No process waits more than $(n-1)q$ time units.

Example of RR with Time Quantum = 20

<u>Process</u>	<u>Burst Time</u>
P ₁	53
P ₂	17
P ₃	68
P ₄	24


- The Gantt chart is:



- Typically, **higher average turnaround** than SJF, but **better response**.

Round Robin (Time Quantum=50)


i	$\tau(p_i)$
0	350
1	125
2	475
3	250
4	75



$$R(p_0) = 0$$

Round Robin (Time Quantum=50)

i	$\tau(p_i)$
0	350
1	125
2	475
3	250
4	75




$$R(p_0) = 0$$

$$R(p_1) = 50$$

Round Robin (Time Quantum=50)

i	$\tau(p_i)$
0	350
1	125
2	475
3	250
4	75




$$R(p_0) = 0$$

$$R(p_1) = 50$$

$$R(p_2) = 100$$

Round Robin (Time Quantum=50)

i	$\tau(p_i)$
0	350
1	125
2	475
3	250
4	75



$$R(p_0) = 0$$

$$R(p_1) = 50$$


$$R(p_2) = 100$$

$$R(p_3) = 150$$

$$R(p_4) = 200$$

Round Robin (Time Quantum=50)

i	$\tau(p_i)$
0	350
1	125
2	475
3	250
4	75



$$T_{\text{TRnd}}(p_1) = 550$$

$$T_{\text{TRnd}}(p_4) = 475$$

$$R(p_0) = 0$$

$$R(p_1) = 50$$


$$R(p_2) = 100$$

$$R(p_3) = 150$$

$$R(p_4) = 200$$

Round Robin (Time Quantum=50)

i	$\tau(p_i)$
0	350
1	125
2	475
3	250
4	75



$$T_{\text{TRnd}}(p_1) = 550$$

$$T_{\text{TRnd}}(p_3) = 950$$

$$T_{\text{TRnd}}(p_4) = 475$$

$$R(p_0) = 0$$

$$R(p_1) = 50$$


$$R(p_2) = 100$$

$$R(p_3) = 150$$

$$R(p_4) = 200$$

Round Robin (Time Quantum=50)

i	$\tau(p_i)$
0	350
1	125
2	475
3	250
4	75



$$T_{\text{TRnd}}(p_0) = 1100$$

$$T_{\text{TRnd}}(p_1) = 550$$

$$T_{\text{TRnd}}(p_3) = 950$$

$$T_{\text{TRnd}}(p_4) = 475$$

$$R(p_0) = 0$$

$$R(p_1) = 50$$


$$R(p_2) = 100$$

$$R(p_3) = 150$$

$$R(p_4) = 200$$

Round Robin (Time Quantum=50)

i	$\tau(p_i)$
0	350
1	125
2	475
3	250
4	75



$$T_{\text{TRnd}}(p_0) = 1100$$

$$T_{\text{TRnd}}(p_1) = 550$$

$$T_{\text{TRnd}}(p_2) = 1275$$

$$T_{\text{TRnd}}(p_3) = 950$$

$$T_{\text{TRnd}}(p_4) = 475$$

$$R(p_0) = 0$$

$$R(p_1) = 50$$

$$R(p_2) = 100$$


$$R(p_3) = 150$$

$$R(p_4) = 200$$

Round Robin (Time Quantum=50)

- Equitable
- Most widely-used

i	$\tau(p_i)$
0	350
1	125
2	475
3	250
4	75



$$T_{\text{TRnd}}(p_0) = 1100$$

$$R(p_0) = 0$$

$$T_{\text{TRnd}}(p_1) = 550$$

$$R(p_1) = 50$$

$$T_{\text{TRnd}}(p_2) = 1275$$

$$R(p_2) = 100$$

$$T_{\text{TRnd}}(p_3) = 950$$

$$R(p_3) = 150$$

$$T_{\text{TRnd}}(p_4) = 475$$

$$R(p_4) = 200$$


$$T_{\text{TRnd-avg}} = (1100+550+1275+950+475)/5 = 4350/5 = 870$$

$$\text{Average response (wait) time } R_{\text{avg}} = (0+50+100+150+200)/5 = 500/5 = 100$$

RR (TQ=50, Overhead =10)

•Context Switch Overhead must be considered

i	$\tau(p_i)$
0	350
1	125
2	475
3	250
4	75



$$T_{\text{TRnd}}(p_0) = 1320$$

$$R(p_0) = 0$$

$$T_{\text{TRnd}}(p_1) = 660$$

$$R(p_1) = 60$$

$$T_{\text{TRnd}}(p_2) = 1535$$

$$R(p_2) = 120$$

$$T_{\text{TRnd}}(p_3) = 1140$$

$$R(p_3) = 180$$


$$T_{\text{TRnd}}(p_4) = 565$$

$$R(p_4) = 240$$

$$T_{\text{TRnd-avg}} = (1320+660+1535+1140+565)/5 = 5220/5 = 1044$$

$$\text{Average response (wait) time } R_{\text{avg}} = (0+60+120+180+240)/5 = 600/5 = 120$$

Time Quantum and Context Switch Time



- q large \Rightarrow FIFO
- q small $\Rightarrow q$ must be large with respect to context switch, otherwise overhead is too high.

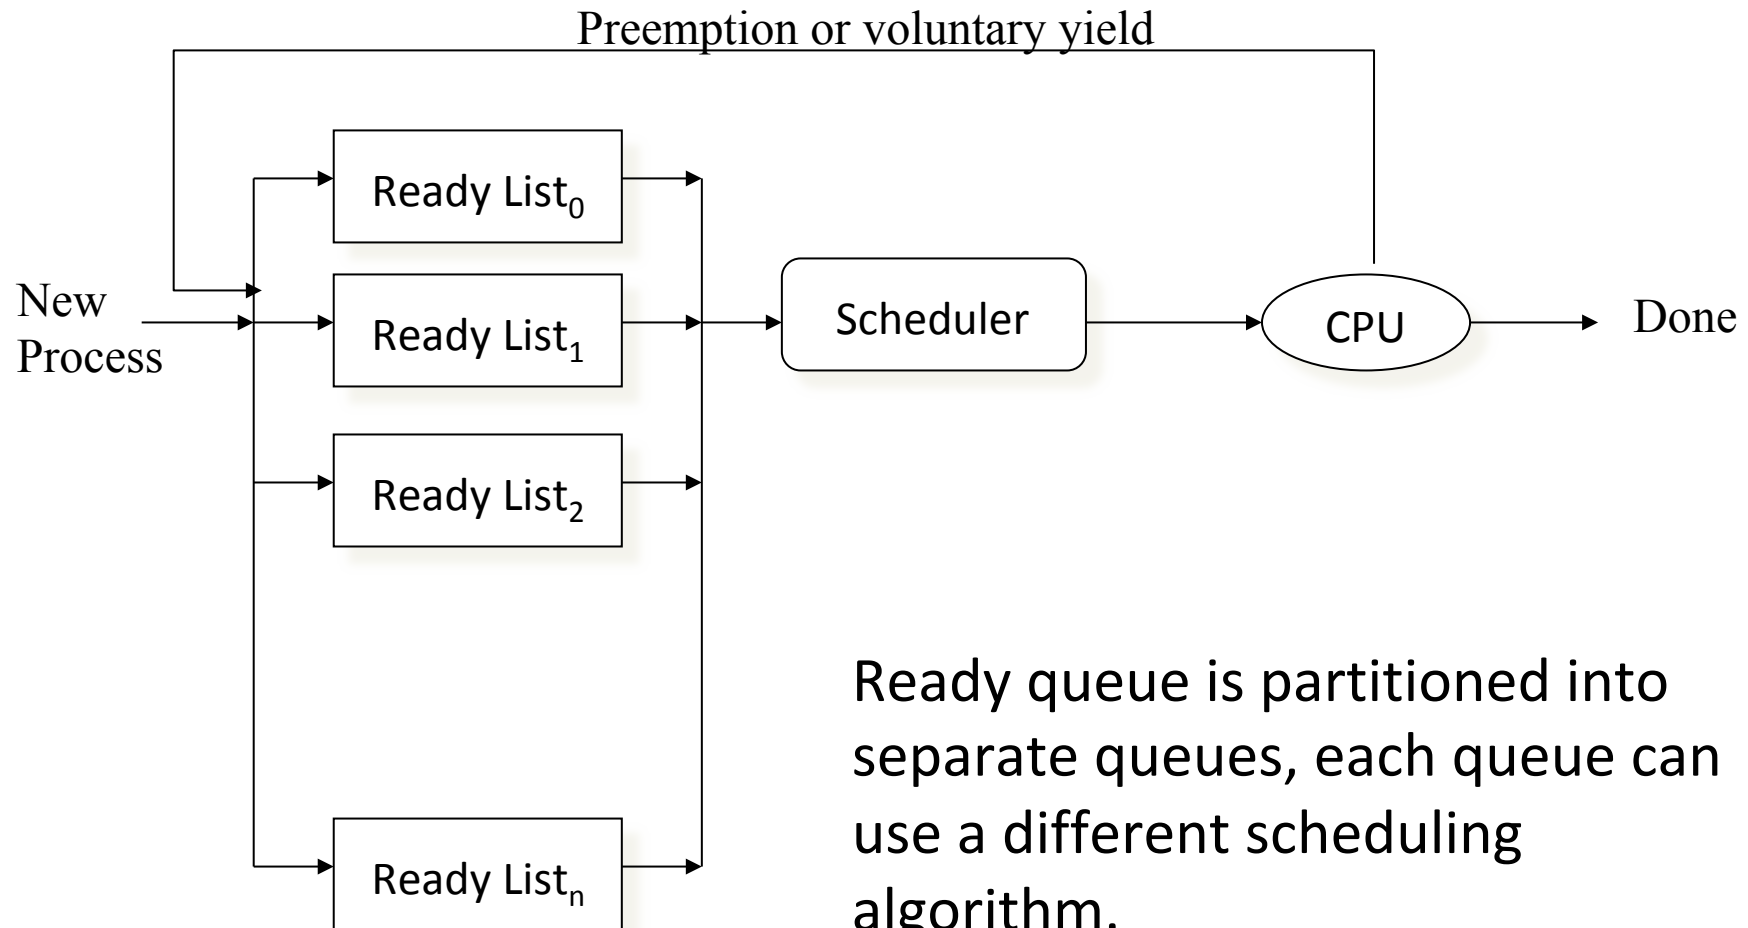
RR algorithm

- Consider a variant of the RR scheduling algorithm where the entries in the ready queue are pointers to the PCBs.
- What would be the effect of putting two pointers to the same process in the ready queue?
- How would you modify the basic RR algorithm to achieve the same effect without the duplicate pointers?

5. Multilevel Queue


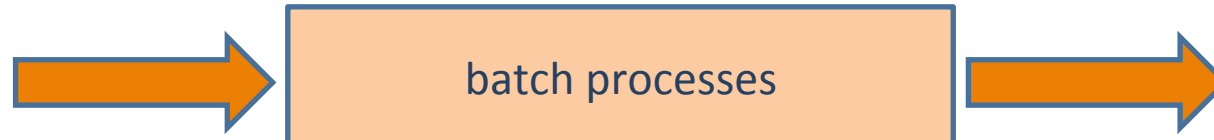

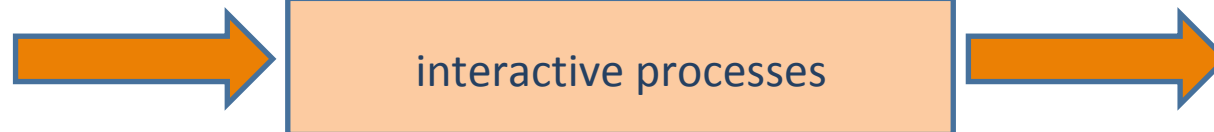

- Ready queue is partitioned into separate queues:
foreground (interactive)
background (batch)
- Each queue has its own scheduling algorithm,
foreground – RR
background – FCFS
- Scheduling must be done between the queues.
 - Fixed priority scheduling – (i.e., serve all from foreground then from background). Possibility of **starvation**.
 - Time slice – each queue gets a certain amount of CPU time which it can schedule amongst its processes; i.e., 80% to foreground in RR, 20% to background in FCFS

Multi-Level Queues



Example- Multi-level Queue

Highest priority




Lowest priority

6. Multilevel Feedback Queue

- Multilevel queue with feedback scheduling is similar to multilevel queue; however, it allows processes to move between queues.
- **Aging** can be implemented this way.
- Multilevel-feedback-queue scheduler defined by the following parameters:
 - number of queues
 - scheduling algorithms for each queue
 - method used to determine when to upgrade a process
 - method used to determine when to demote a process
 - method used to determine which queue a process will enter when that process needs service

Example of Multilevel Feedback Queue


- Three queues:
 - Q_0 – time quantum 8 milliseconds
 - Q_1 – time quantum 16 milliseconds
 - Q_2 – FCFS non-preemption
- **Example Scheduling**
 - A new job enters queue Q_0 . When it gains CPU, job receives 8 milliseconds. If it does not finish in 8 milliseconds, job is moved to queue Q_1 .
 - At Q_1 job is again served FCFS and receives 16 additional milliseconds.
 - If it still does not complete, it is preempted and moved to queue Q_2 .



Multiple-Processor Scheduling

- CPU scheduling is more complex when multiple CPUs are available
 - **Asymmetric multiprocessing** – only one processor accesses the system data structures, alleviating the need for data sharing
 - **Symmetric multiprocessing (SMP)** – each processor is self-scheduling, all processes in common ready queue, or each has its own private queue of ready processes
- **Processor affinity** – process has affinity for processor on which it is currently running
 - **soft affinity** (can be changed at a later time)
 - **hard affinity** (process doesn't move to another processor)
 - Variations including **processor sets**

Multicore Scheduling



- In SMP, each CPU has its own private ready queue for the processes waiting to be run on that CPU

Process Migration

- If SMP, need to keep all CPUs loaded for efficiency
- **Load balancing** attempts to keep workload evenly distributed
- **Push migration** – periodically checks load on each processor, and if found any of them overloaded, pushes task from overloaded CPU to other CPUs
- **Pull migration** – idle processors pull waiting tasks from busy processor

Starvation

- Which of the following scheduling algorithms could result in starvation?
 - First-come, first serve
 - Shortest job first
 - Round robin
 - Priority

Reading

- Read Chapter 5 (Textbook)
- Read Chapter 4 (Linux Kernel Development)
- Acknowledgments
 - These slides are adapted from
 - Öznur Özkasap (Koç University)
 - Operating System and Concepts (9th edition) Wiley