

# Synchronization-III

## Semaphores

Didem Unat

Lecture 10

COMP304 - Operating Systems (OS)

# Mutex Locks

```
do {  
    acquire lock  
    critical section  
    release lock  
    remainder section  
} while (true);
```

```
acquire() {  
    while (!available)  
        ; /* busy wait */  
    available = false;;  
}  
  
release() {  
    available = true;  
}
```

- Software interface for locks
- Calls to **acquire()** and **release()** must be atomic
  - implemented via hardware atomic instructions
- But this solution requires **busy waiting**
  - And also called a **spinlock**

# Semaphores

- We want to be able to write more complex constructs
  - need a language to do so. We define semaphores which we assume are atomic operations.
- Semaphores are more general synchronization tools
  - Operating System Primitive
  - Two standard atomic operations modify semaphore variable *S*: *wait()* and *signal()*

**WAIT ( S ):**

```
while ( S <= 0 );  
S = S - 1;
```

**SIGNAL ( S ):**

```
S = S + 1;
```

- As given here, these are not atomic as written in "macro code". We define these operations, however, to be atomic (Protected by a hardware lock.)

# Critical Section for n Processes

- Shared semaphore:

**semaphore mutex = 1;** //initial value

- Process  $P_i$ :

```
do {  
    wait(mutex)  
    critical section  
    signal(mutex)  
    remainder section  
} while (true);
```

# Shared Account Balance Example

```
semaphore mutex = 1;
```

```
proc_0() {  
    . . .  
    /* Enter the CS */  
    wait(mutex);  
    balance += amount;  
    signal(mutex);  
    . . .  
}
```

```
proc_1() {  
    . . .  
    /* Enter the CS */  
    wait(mutex);  
    balance -= amount;  
    signal(mutex);  
    . . .  
}
```

//CS stands for critical section

# Semaphore as a general synchronization tool

- Provides **mutual exclusion**

Semaphore S = 1; // initialized to 1 or initialized to # of resources

wait (S);

Critical Section

signal (S);

- **Counting** semaphore – integer value can range over an unrestricted domain
  - For example: resources in the hardware: semaphore is initialized to number of printers
- **Binary** semaphore – integer value can range only between 0 and 1; can be simpler to implement
  - This is the same as **mutex locks**

# Semaphore as a general synchronization tool

- Semaphores can be used to force synchronization ( precedence ) if the **preceeder** does a signal at the end, and the **follower** does wait at beginning.

For example, here we want P1 to execute before P2.

- Execute B in  $P_j$  only after A executed in  $P_i$
- Use semaphore flag initialized to **0**

Code:



# Classical Problems of Synchronization

- Bounded-Buffer Problem
- Readers and Writers Problem
- Dining Philosophers Problem

These classical problems are used for testing newly proposed synchronization methods.




# Bounded-Buffer Problem

- Buffer size: buffer can hold **n** items
- Binary semaphores  
**semaphore mutex;**
- Counting semaphores  
**semaphore full, empty;**

Initially:

**full = 0, empty = n, mutex = 1**



This is the same producer / consumer problem as before. But now we'll do it with signals and waits. Remember: a **wait decreases** its argument and a **signal increases** its argument.

# Bounded-Buffer Problem

- Empty and Full are counting semaphores

```
producer:
do {
    /* produce an item in nextp */
    wait (empty);          /* Do action      */

    /* add nextp to buffer */

    signal (empty);
} while(TRUE);
```

```
consumer:
do {
    wait (full);

    /* remove an item from buffer to nextc */

    signal (full);

    /* consume an item in nextc */
} while(TRUE);
```

Signals are wrong!  
Producer (consumer)  
needs to wake up the  
consumer (producer)!

# Bounded-Buffer Problem

- **Mutex** is binary semaphore, Empty and Full are counting semaphores

```
producer:
do {
    /* produce an item in nextp */
    wait (empty);          /* Do action */

    /* add nextp to buffer */

    signal (full);

} while(TRUE);
```

Does this work for multiple producers and consumers?

Only works for one producer and consumer. Need the mutex to prevent multiple producers writing into the same buffer.

```
consumer:
do {
    wait (full);

    /* remove an item from buffer to nextc */

    signal (empty);

    /* consume an item in nextc */
} while(TRUE);
```

# Bounded-Buffer Problem

- **Mutex** is binary semaphore, Empty and Full are counting semaphores

```
producer:
do {
    /* produce an item in nextp */
    wait (empty);          /* Do action      */
    wait (mutex);          /* Buffer guard*/

    /* add nextp to buffer */

    signal (mutex);
    signal (full);

} while(TRUE);
```

```
consumer:
do {
    wait (full);
    wait (mutex);

    /* remove an item from buffer to nextc */

    signal (mutex);
    signal (empty);

    /* consume an item in nextc */
} while(TRUE);
```

We need a mutex in addition to empty and full semaphores because we are operating on the same buffer

# Semaphores

- **Spinlocks** (mutexes) are useful in a system since no context switch is required
- A disadvantage of mutex solutions so far:
  - they all require **busy waiting**.
- To overcome busy waiting → **blocking a process**

# No busy waiting (blocking) Semaphores

- With each semaphore there is an associated waiting queue:
  - Keeps list of processes waiting on the semaphore
- Two operations:
  - **Block** – place the process invoking the operation on the appropriate waiting queue if semaphore == false
  - **Wakeup** – Wakes up one of the blocked processes upon getting a signal and places the process to ready queue

# Blocking Semaphores

```
typedef struct {  
    int    value;  
    struct process *list; /* list of processes waiting on S */  
} SEMAPHORE;
```

```
SEMAPHORE s;  
wait(s) {  
    s.value = s.value - 1;  
    if ( s.value < 0 ) {  
        add this process to s.list;  
        block ();  
    }  
}
```

```
SEMAPHORE s;  
signal(s) {  
    s.value = s.value + 1;  
    if ( s.value <= 0 ) {  
        remove a process P from s.list;  
        wakeup(P);  
    }  
}
```

**Block** – place the process invoking the operation on the appropriate waiting queue if semaphore is not available

**Wakeup** – Wakes up one of the blocked processes upon getting a signal and places the process to ready queue

# Deadlock and Starvation

- **Deadlock** – two or more processes are waiting indefinitely for an event that can be caused by only one of the waiting processes.
- Let S and Q be two semaphores initialized to 1

$P_0$   
wait(S);  
wait(Q);  
⋮  
signal(S);  
signal(Q)

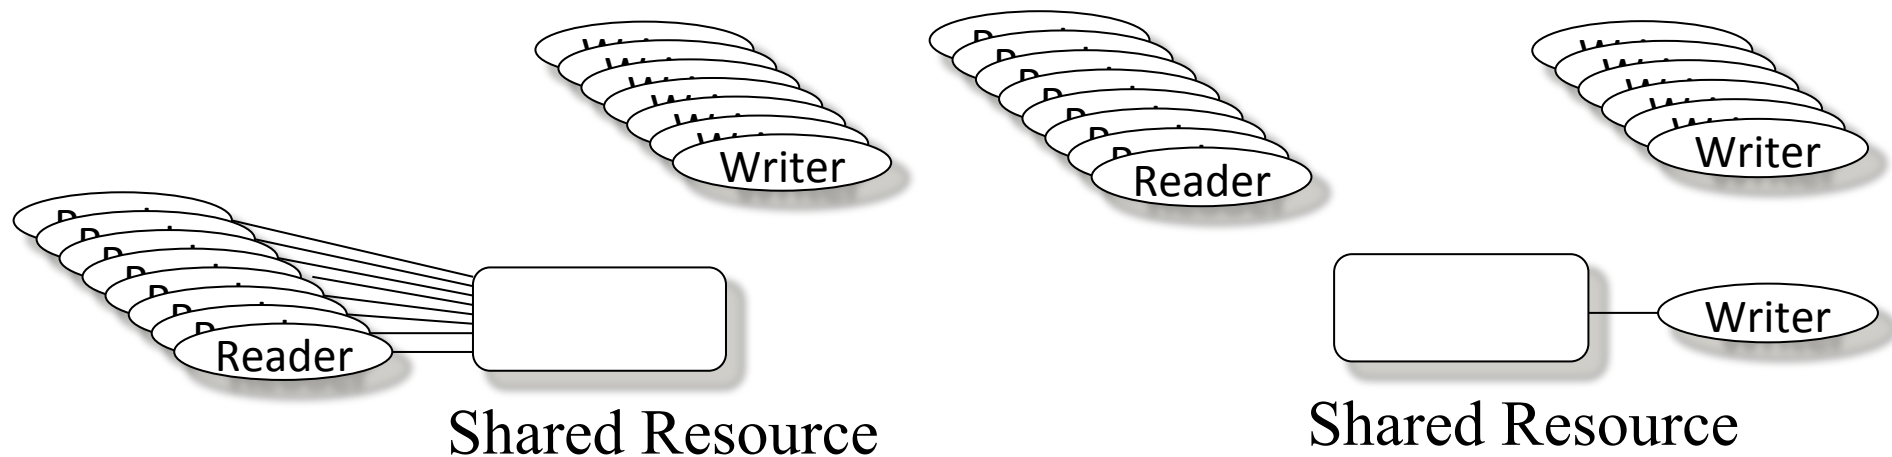
$P_1$   
wait(Q);  
wait(S);  
⋮  
signal(Q);  
signal(S);

- **Starvation** – indefinite blocking. A process may never be removed from the semaphore queue in which it is suspended.
- May occur if we add and remove processes from the list in LIFO order or based on priority.



# Readers and Writers Problem

- A data set is shared among a number of concurrent processes
  - **Readers** – only read the data set; they do **not** perform any updates
  - **Writers** – can both read and write.
- Problem
  - Allow multiple readers to read at the same time.
  - Only one single writer can access the shared data at a time.



# Readers/Writers Problem

## Locks:

- are **shared** (for the readers) and **exclusive** (for the writer).

Two possible (contradictory) guidelines can be used:

- No reader is kept waiting unless a writer holds the lock (the readers have precedence).
  - First readers problem
- If a writer is waiting for access, no new reader gains access (writer has precedence).
  - Second readers problem

Can starvation occur on either of these rules?

# First-Readers Problem

Favoring readers over writers

```
semaphore rdwrt = 1;
semaphore cntmutex = 1;
int readcount = 0;
```

**Writer:**

```
do {
    wait( rdwrt );
    /*    writing is performed    */
    signal( rdwrt );

} while(TRUE);
```

**Reader:**

```
do {
    wait( cntmutex );
    readcount = readcount + 1;
    if readcount == 1 then
        wait( rdwrt );
        signal( cntmutex );

        /*    reading is performed    */

    wait( cntmutex );
    readcount = readcount - 1;
    signal( cntmutex );
    if readcount == 0 then
        signal( rdwrt );
} while(TRUE);
```

A reader might be  
reading while a writer is  
writing at the same  
time or two writers can  
perform their writes!

How can that happen?

# First-Readers Problem

**Correct Implementation,** Favoring readers over writers

```
semaphore rdwrt = 1;
semaphore cntmutex = 1;
int readcount = 0;
```

**Writer:**

```
do {
    wait( rdwrt );
    /*    writing is performed    */
    signal( rdwrt );

} while(TRUE);
```

**Reader:**

```
do {
    wait( cntmutex );
    readcount = readcount + 1;
    if readcount == 1 then
        wait( rdwrt );
        signal( cntmutex );

        /*    reading is performed    */

    wait( cntmutex );
    readcount = readcount - 1;
    if readcount == 0 then
        signal( rdwrt );
        signal( cntmutex );

} while(TRUE);
```

# Quiz

- Many events on campus take place in SGK. Assume that the show room can only allow  $N$  many people. As soon as  $N$  audience are in the room, the staff will not accept another incoming audience until an existing audience leaves the event.
- Explain how semaphores can be used by the staff to limit the number of people in the SGK event room.
- Show your pseudo-code.

# Solution

- Semaphore is initialized to the number of allowable people in SGKM. When a reservation is accepted, the `acquire()` method is called; when someone leaves the event, the `release()` method is called. If the system reaches the number of allowable people, subsequent calls to `acquire()` will block until an existing person leaves the event and the `release` method is invoked.

# Reading

- Read Chapter 6
- Acknowledgments
  - These slides are adapted from
    - Öznur Özkasap (Koç University)
    - Operating System and Concepts (9<sup>th</sup> edition) Wiley
    - Jerry Breecher