

Synchronization-IV

Monitors

Didem Unat

Lecture 11

COMP304 - Operating Systems (OS)

Semaphores

- We want to be able to write more complex constructs so need a language to do so. We define semaphores which we assume are atomic operations.
- Semaphores are more general synchronization tools
 - Operating System Primitive
 - Two standard atomic operations modify semaphore variable *S*: *wait()* and *signal()*

WAIT (S):

```
while ( S <= 0 );  
S = S - 1;
```

SIGNAL (S):

```
S = S + 1;
```

- As given here, these are not atomic as written in "macro code". We define these operations, however, to be atomic (Protected by a hardware lock.)

Semaphore as a general synchronization tool

- Provides **mutual exclusion**

Semaphore S = 1; // initialized to 1 or initialized to # of resources

wait (S);

Critical Section

signal (S);

- **Counting** semaphore – integer value can range over an unrestricted domain
 - For example: resources in the hardware: semaphore is initialized to number of resources
- **Binary** semaphore – integer value can range only between 0 and 1; can be simpler to implement
 - This is the same as **mutex locks**

Semaphore as a general synchronization tool

- Semaphores can be used to force synchronization (precedence) if the **preceeder** does a signal at the end, and the **follower** does wait at beginning.

For example, here we want P1 to execute before P2.

- Execute B in P_j only after A executed in P_i
- Use semaphore flag initialized to **0**

Code:



Blocking Semaphores

```
typedef struct {  
    int    value;  
    struct process *list; /* list of processes waiting on S */  
} SEMAPHORE;
```

```
SEMAPHORE s;  
wait(s) {  
    s.value = s.value - 1;  
    if ( s.value < 0 ) {  
        add this process to s.list;  
        block ();  
    }  
}
```

```
SEMAPHORE s;  
signal(s) {  
    s.value = s.value + 1;  
    if ( s.value <= 0 ) {  
        remove a process P from s.L;  
        wakeup(P);  
    }  
}
```

Block – place the process invoking the operation on the appropriate waiting queue if semaphore is not available

Wakeup – Wakes up one of the blocked processes upon getting a signal and places the process to ready queue

Semaphores vs Locks

- Processes that are blocked at the level of program logic are placed on queues, rather than busy-waiting
- Busy-waiting may be used for the mutual exclusion
 - But these are very short critical sections
- Unlike locks, counting semaphores can take an integer value representing total number of resources

Problems with Semaphores (and Locks)

- Semaphores are shared global variables
 - Can be accessed from anywhere
- Used for both critical sections (mutual exclusion) and for coordination (scheduling or ordering execution)
- Incorrect use of semaphore operations
 - Call signal first and later on call wait
 - `signal(mutex) wait(mutex)`
 - Call wait after another wait
 - `wait(mutex) wait(mutex)`
 - Omitting of wait or signal
- Thus, they are prone to bugs
- To deal with such issues,
 - Introduce a high-level synchronization construct - **monitors**

Monitors

- A monitor is a **programming language construct** that supports controlled access to shared data
 - First developed in Concurrent Pascal
 - It resembles an object-oriented approach for synchronization
- *A monitor encapsulates*
 - **shared data** structures
 - **procedures** that operate on the shared data (protects shared data)
 - **synchronization** between concurrent processes that invoke those procedures

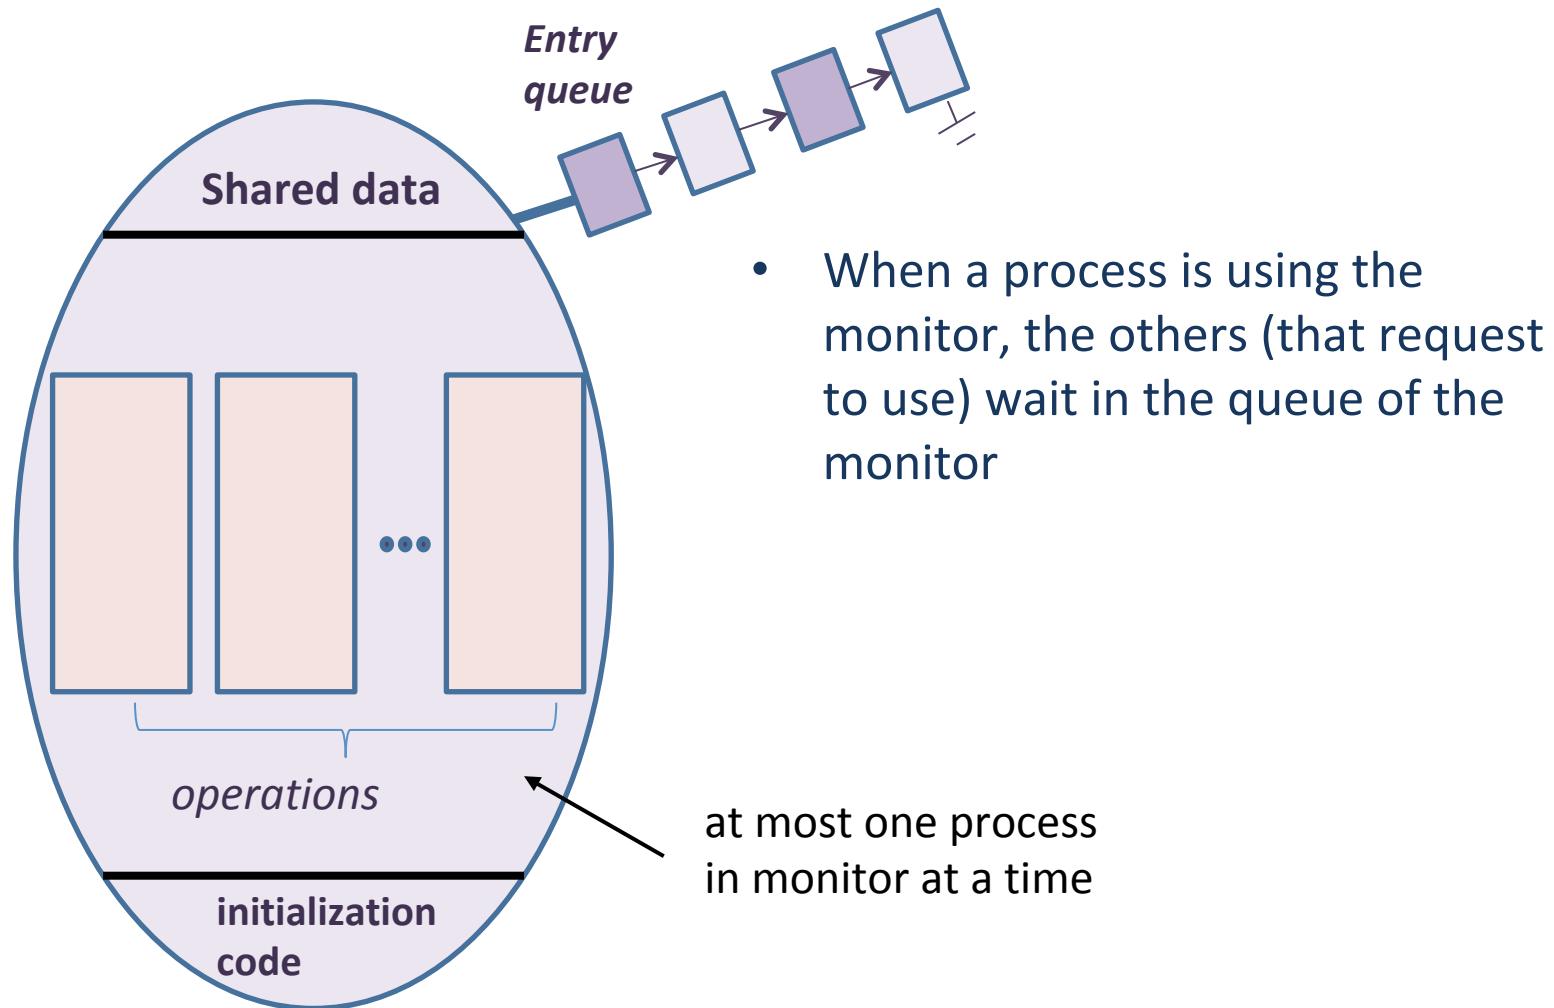
```
monitor monitor-name
{
    // shared variable declarations
    procedure P1 (...) {.....}

    procedure Pn (...) {.....}

    initialization(...) { ... }
}
```

Monitor construct ensures that only one process at a time can be **active** within the monitor

Schematic View of a Monitor



Example: Shared Balance


```
monitor sharedBalance {  
  
    double balance;  
  
    void deposit(double amount) {  
        balance += amount;  
    }  
    void withdraw(double amount) {  
        balance -= amount;  
    }  
    . . .  
}
```

- Balance is a **shared** variable
- Deposit and withdraw are **procedures**
- Processes do not directly read/write into the shared variables but access them via these procedures


Producer-Consumer Problem

- Buffer size: buffer can hold **n** items

```
monitor ProducerConsumer {  
  
    item buffer[N];  
  
    void produce(item x) {  
        //add an item to the buffer  
    }  
    void consume(item *x) {  
        //remove an item from the buffer  
    }  
    . . .  
}
```

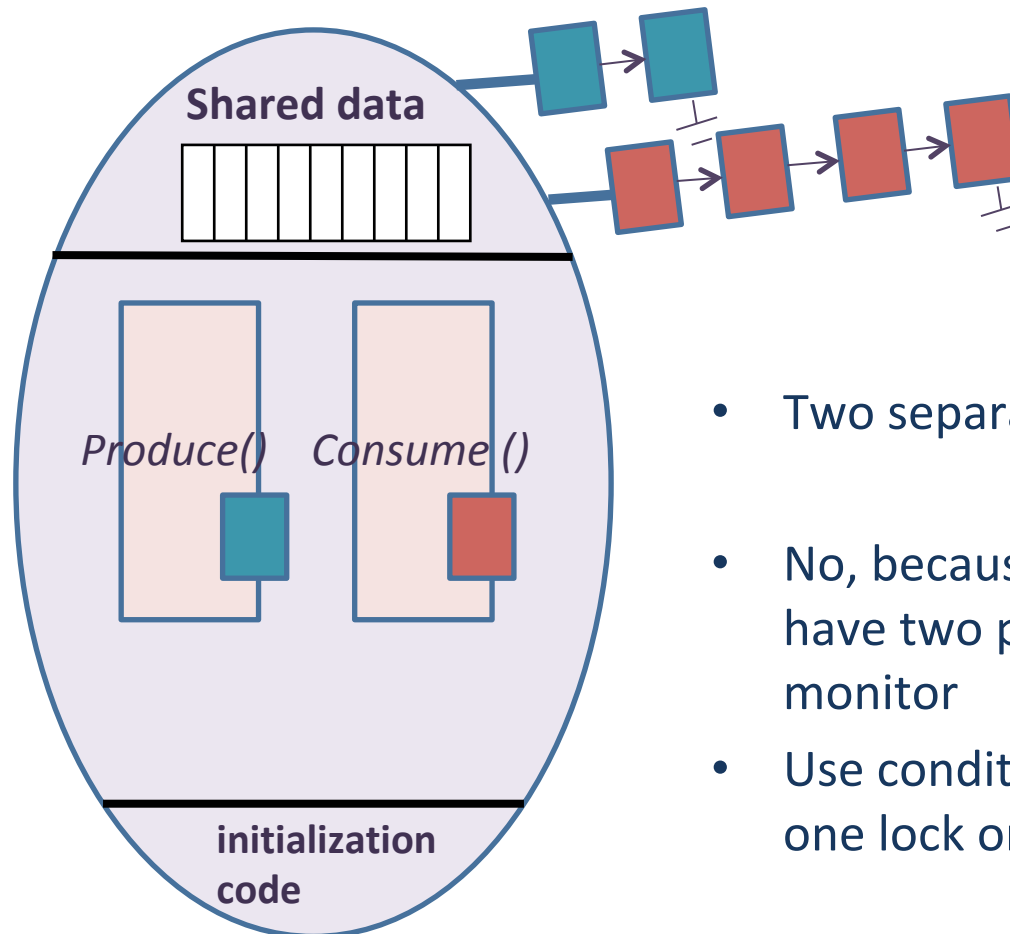


Example: Producer-Consumer Problem



- Buffer is empty
- A consumer (red) is in the queue
- Now what?

Example: Producer-Consumer Problem



- Two separate queues?
- No, because then we would have two processes in the monitor
- Use condition variables sharing one lock on the queue

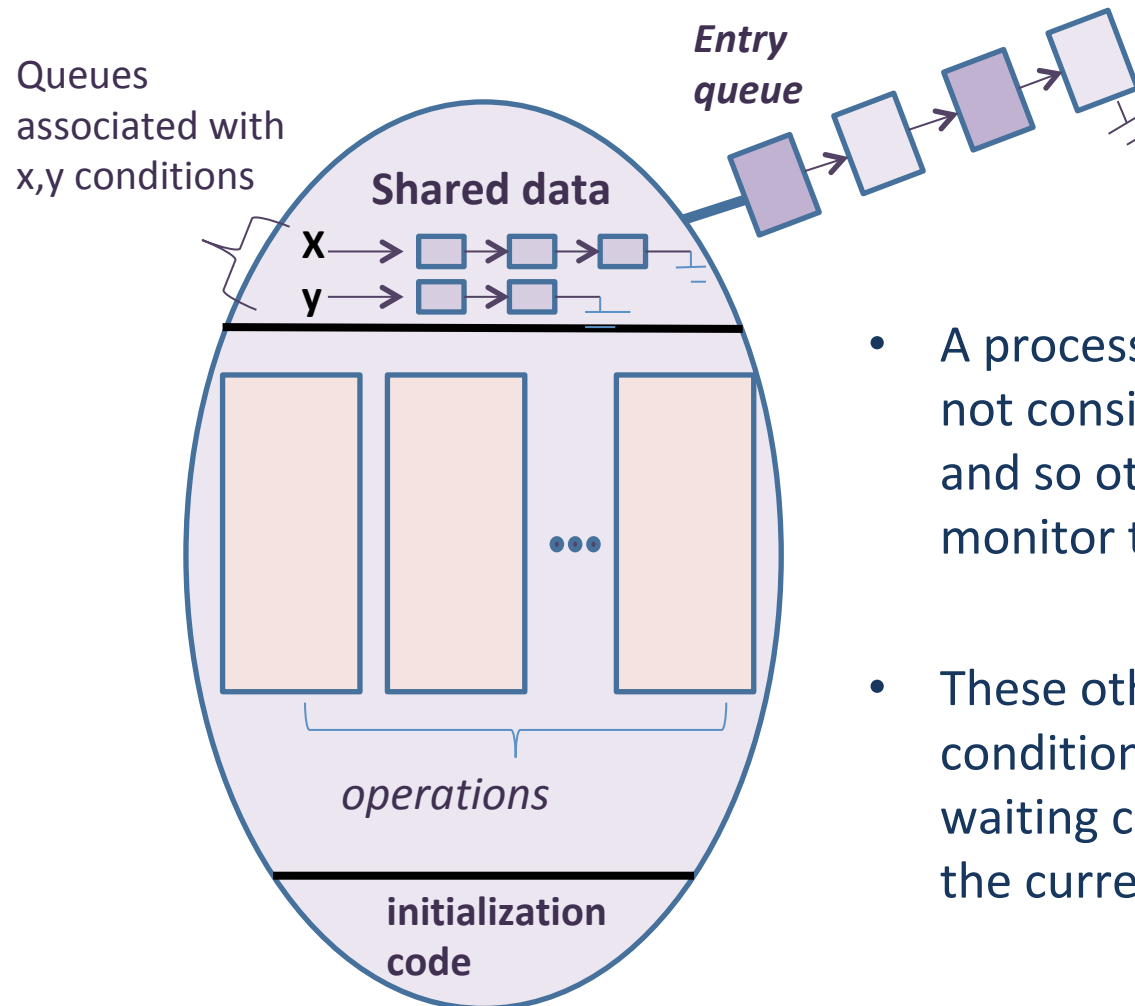
Conditional Variables

- Conceptually a condition variable is a queue of processes, associated with a monitor on which a process may wait for some condition to become true
- Sometimes called a rendezvous point
 - To allow a process to wait within the monitor, a **condition variable** must be declared, as

condition c;

- Three operations on condition variables 'c'
 - **wait(c)**
 - The calling process is suspended until another process invokes it
 - **signal(c)**
 - wake up at most one waiting process
 - if no waiting processes, signal has no effect
 - this is different than semaphores: no history!
 - **broadcast(c)**
 - wake up all waiting processes

Monitor with Condition Variables



- A process waiting for a condition is not considered to occupy the monitor, and so other processes may enter the monitor to change the monitor's state
- These other processes may signal the condition variable to indicate that waiting condition has become true in the current state

Producer and Consumer with Monitors

```
Monitor ProducerConsumer {
    item buffer[N];
    condition not_full, not_empty;

    produce(item x) {
        if (array "buffer" is full) // determined by a count
            wait(not_full);
        insert "x" in array "buffer"
        signal(not_empty);
    }

    consume(item *x) {
        if (array "buffer" is empty) //determined maybe by a count
            wait(not_empty);
        *x = get item from array "buffer"
        signal(not_full);
    }
}
```


Producer and Consumer with Monitors

```
Monitor ProducerConsumer {  
    item buffer[N];  
    condition not_full, not_empty;
```

```
.....  
produce(item x) {  
    .....  
    if (array "buffer" is full) // determined by a count  
        wait(not_full);  
    insert "x" in array "buffer"  
    signal(not_empty);  
    .....  
}
```

EnterMonitor


ExitMonitor

```
consume(item *x) {  
    if (array "buffer" is empty) //determined maybe by a count  
        wait(not_empty);  
    *x = get item from array "buffer"  
    signal(not_full);  
}  
}
```

- Monitor inserts the lock operations before at the entry of produce/consume and releases the lock at the exit of produce/consume
- That's why we don't need a separate mutex lock for multiple producers (or consumers)

Dining Philosophers Problem

- 5 philosophers with 5 chopsticks sit around a circular table.
 - They each want to eat at random times
 - Must pick up the 2 chopsticks to eat
 - Pick one chopstick at a time
- While a philosopher is thinking, she drops the chopsticks on the table




Dining Philosophers Problem

- Shared data
semaphore chopstick[5]; Initially all values are 1

Philosopher i :

```
while (true) {  
    wait(chopstick[i])  
    wait(chopstick[(i+1) % 5])  
  
    // eat  
  
    signal(chopstick[i]);  
    signal(chopstick[(i+1) % 5]);  
  
    // think  
}
```




Deadlock
may occur!

Dining Philosophers Problem

- Shared data
semaphore chopstick[5]; Initially all values are 1

Philosopher i :


```
while (true) {  
    wait(chopstick[i])  
    wait(chopstick[(i+1) % 5])  
  
    // eat  
  
    signal(chopstick[i]);  
    signal(chopstick[(i+1) % 5]);  
  
    // think  
}
```



Starvation
may occur!

Several Solutions

- Allow pickup only if both chopsticks are available
- Odd # philosophers always pick up left chopstick first, even # philosophers always pick up right chopstick first



A deadlock-free solution does not necessarily eliminate the possibility of starvation.

Monitor Solution to Dining Philosophers

- This implements a deadlock-free solutions with a restriction that a philosopher may pick up chopsticks only if both of them are available
- Is this solution starvation free?

```
monitor DP
{
    enum { THINKING, HUNGRY, EATING} state [5] ;
    condition self [5];

    void pickup (int i) {
        state[i] = HUNGRY;
        test(i);
        if (state[i] != EATING)
            self[i].wait;
    }

    void putdown (int i) {
        state[i] = THINKING;
        // test left and right neighbors
        test((i + 4) % 5);
        test((i + 1) % 5);
    }
}
```

Monitor Solution to Dining Philosophers (cont.)

```
void test (int i) {
    if ( (state[(i + 4) % 5] != EATING) &&
        (state[i] == HUNGRY) &&
        (state[(i + 1) % 5] != EATING) ) {

        state[i] = EATING ;
        self[i].signal () ;
    }
}

initialization_code() {
    for (int i = 0; i < 5; i++)
        state[i] = THINKING;
}
```

- If my neighbors are not in eating state and I am hungry, then I can eat!
- All philosophers' states are set to Thinking initially

Solution to DP Problem (cont.)

- Each philosopher i invokes the operations `pickup()` and `putdown()` in the following sequence:

`dp.pickup (i)`

EAT

`dp.putdown (i)`

THINK

Language Support

- Java, C# and several other languages have support for Monitors
 - Here is a Java example

```
// To make a method synchronized, simply add the synchronized  
// keyword to its declaration:
```

```
public class SynchronizedCounter {  
    private int c = 0;  
    public synchronized void increment() {  
        c++;  
    }  
    public synchronized void decrement() {  
        c--;  
    }  
    public synchronized int value() {  
        return c;  
    }  
}
```

Summary of Monitors

- A **monitor** is a synchronization construct that allows processes
 - To have both mutual exclusion and
 - The ability to wait (block) for a certain condition to become true.
- Monitors also have a mechanism for signaling other processes that their condition has been met.
- A monitor consists of a mutex object and **condition variables**, procedures to access them

Summary of Synchronization

- Mutual exclusion is required to ensure no two concurrent processes are in their critical section at the same time
 - To prevent race conditions (corruption of shared data)
 - First identified and solved by Dijkstra in 1965
- Hardware solutions
 - Test-and-Test
 - Compare-and-Swap
- Software solutions (only provide higher-level abstractions to their hardware solutions)
 - Locks (busy-wait)
 - Semaphores (blocking, counting semaphores)
 - Monitors (high level language constructs)

OS Support for Synchronization

- Windows
 - Uses **spinlocks** on multiprocessor systems.
 - Also provides **dispatcher objects** which may act as mutexes and semaphores.
 - Dispatcher objects may also provide **events**. An event acts much like a **condition variable**.
- Linux
 - Disables interrupts to implement short critical sections
 - Provides semaphores and spinlocks
 - Pthreads: provides mutex locks and condition variables

Reading

- Read Chapter 6
- Wikipedia Article on Monitors
 - http://en.wikipedia.org/wiki/Monitor_%28synchronization%29
- Acknowledgments
 - These slides are adapted from
 - Öznur Özkasap (Koç University)
 - Operating System and Concepts (9th edition) Wiley
 - Jerry Breecher (Worcester Polytechnic Institute)
 - Mark Zbikowski and Gary Kimura (Univ. of Washington)